

CS162  
Operating Systems and  
Systems Programming  
Lecture 9

Monitors (Continued)  
Scheduling  
Core Concepts and Classic Policies

Professor Natacha Crooks  
<https://cs162.org/>

# Recall: Monitors are better!

---

Use *locks* for mutual exclusion and *condition variables* for scheduling constraints

**Monitor:** a **lock** and zero or more **condition variables** for managing concurrent access to shared data

- A monitor is a paradigm for concurrent programming
- Some languages like Java provide this natively
  - Most others use actual locks and condition variables

# Recall: Wait & Signal Pattern

---

```
...  
acquire(&buf_lock)  
...  
cond_signal(&buf_CV);  
...  
release(&buf_lock);
```

```
acquire(&buf_lock);  
...  
while (isEmpty(&queue)) {  
    cond_wait(&buf_CV, &buf_lock)  
};  
...  
lock.Release();
```

# Hoare Semantics

---

## Thread A

```
...  
acquire(&buf_lock)  
...  
cond_signal(&buf_CV);  
...  
release(&buf_lock);
```

## Thread B

```
acquire(&buf_lock);  
...  
if (isEmpty(&queue)) {  
    cond_wait(&buf_CV, &buf_lock)  
};  
...  
lock.Release();
```

- 1. When call signal, handover buf\_lock to thread B.**
- 2. Thread B gets immediately scheduled (nothing can run in between).**
- 3. Thread B eventually releases lock.**

# Mesa Semantics

---

## Thread A

```
...  
acquire(&buf_lock)  
...  
cond_signal(&buf_CV);  
...  
release(&buf_lock);
```

## Thread B

```
acquire(&buf_lock);  
...  
while (isEmpty(&queue)) {  
    cond_wait(&buf_CV, &buf_lock)  
};  
...  
lock.Release();
```

1. When call signal, keep lock. Place Thread B on READY queue (no special priority)
2. Thread A eventually releases buf\_lock.
3. Thread B eventually gets scheduled and acquires buf\_lock. Thread C may have run in between.
4. Thread B eventually releases buf\_lock.

# Basic Structure of *Mesa* Monitor Program

---

Monitors represent the synchronization logic of the program

- Wait if necessary
- Signal when change something so any waiting threads can proceed

```
lock
while (need to wait) {
    condvar.wait();
}
unlock
```

do something so no need to wait

```
lock


condvar.signal();

unlock
```

} Check and/or update  
state variables  
Wait if necessary

} Check and/or update  
state variables

# Readers/Writers Problem



Motivation: Consider a shared database

- Two classes of users:
  - » Readers – never modify database
  - » Writers – read and modify database
- Is using a single lock on the whole database sufficient?
  - » Like to have many readers at the same time
  - » Only one writer at a time

# Basic Readers/Writers Solution

---

## Correctness Constraints:

- Readers can access database when no writers
- Writers can access database when no readers or writers
- Only one thread manipulates state variables at a time

## Basic structure of a solution:

### –Reader()

```
Wait until no writers
Access data base
Check out - wake up a waiting writer
```

### –Writer()

```
Wait until no active readers or writers
Access database
Check out - wake up waiting readers or writer
```



# Basic Readers/Writers Solution

---

State variables (Protected by a lock called “lock”):

- » int AR: Number of active readers; initially = 0
- » int WR: Number of waiting readers; initially = 0
- » int AW: Number of active writers; initially = 0
- » int WW: Number of waiting writers; initially = 0
- » Condition okToRead = NIL
- » Condition okToWrite = NIL

# Code for a Reader

---

```
Reader() {
    // First check self into system
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);
    // Perform actual read-only access
    AccessDatabase(ReadOnly);
    // Now, check out of system
    acquire(&lock);
    AR--; // No longer active
    if (AR == 0 && WW > 0) // No other active readers
        cond_signal(&okToWrite); // Wake up one writer
    release(&lock);
}
```

# Code for a Writer

---

```
Writer() {
    // First check self into system
    acquire(&lock);
    while ((AW + AR) > 0) { // Is it safe to write?
        WW++; // No. Active users exist
        cond_wait(&okToWrite, &lock); // Sleep on
cond var
        WW--; // No longer waiting
    }
    AW++; // Now we are active!
    release(&lock);
    // Perform actual read/write access
    AccessDatabase(ReadWrite);
    // Now, check out of system
    acquire(&lock);
    AW--; // No longer active
    if (WW > 0) { // Give priority to writers
        cond_signal(&okToWrite); // Wake up one writer
    } else if (WR > 0) { // Otherwise, wake reader
        cond_broadcast(&okToRead); // Wake all
readers
    }
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

Use an example to simulate the solution

Consider the following sequence of operators:

– R1, R2, W1, R3

Initially:  $AR = 0$ ,  $WR = 0$ ,  $AW = 0$ ,  $WW = 0$

# Simulation of Readers/Writers Solution

---

R1 comes along (no waiting threads)

$AR = 0, WR = 0, AW = 0, WW = 0$

```
Reader() {  
    acquire(&lock)  
    while ((AW + WW) > 0) { // Is it safe to read?  
        WR++; // No. Writers exist  
        cond_wait(&okToRead, &lock); // Sleep on cond  
var  
        WR--; // No longer waiting  
    }  
    AR++; // Now we are active!  
    release(&lock);  
  
    AccessDBase(ReadOnly);  
  
    acquire(&lock);  
    AR--;  
    if (AR == 0 && WW > 0)  
        cond_signal(&okToWrite);  
    release(&lock);  
}
```

# Simulation of Readers/Writers Solution

---

R1 comes along (no waiting threads)

$AR = 0, WR = 0, AW = 0, WW = 0$

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);

    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

R1 comes along (no waiting threads)

AR = 1, WR = 0, AW = 0, WW = 0

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);

    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

R1 comes along (no waiting threads)

AR = 1, WR = 0, AW = 0, WW = 0

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);

    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&lock);
}
```



# Simulation of Readers/Writers Solution

---

R1 accessing dbase (no other threads)

AR = 1, WR = 0, AW = 0, WW = 0

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);
    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&lōck);
}
```

# Simulation of Readers/Writers Solution

---

R2 comes along (R1 accessing dbase)

$AR = 1, WR = 0, AW = 0, WW = 0$

```
Reader() {  
    acquire(&lock);  
    while ((AW + WW) > 0) { // Is it safe to read?  
        WR++; // No. Writers exist  
        cond_wait(&okToRead, &lock); // Sleep on cond  
var  
        WR--; // No longer waiting  
    }  
    AR++; // Now we are active!  
    release(&lock);  
  
    AccessDBase(ReadOnly);  
  
    acquire(&lock);  
    AR--;  
    if (AR == 0 && WW > 0)  
        cond_signal(&okToWrite);  
    release(&löck);  
}
```

# Simulation of Readers/Writers Solution

---

R2 comes along (R1 accessing dbase)

AR = 1, WR = 0, AW = 0, WW = 0

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);

    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

R2 comes along (R1 accessing dbase)

AR = 2, WR = 0, AW = 0, WW = 0

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);

    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

R2 comes along (R1 accessing dbase)

AR = 2, WR = 0, AW = 0, WW = 0

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);

    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&lōck);
}
```

# Simulation of Readers/Writers Solution

---

R1 and R2 accessing dbase

AR = 2, WR = 0, AW = 0, WW = 0

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);
    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
```

Assume readers take a while to access database  
Situation: Locks released, only AR is non-zero

# Simulation of Readers/Writers Solution

---

W1 comes along (R1 and R2 are still accessing dbase)

$AR = 2, WR = 0, AW = 0, WW = 0$

```
Writer() {
    acquire(&lock);
    while ((AW + AR) > 0) { // Is it safe to
write?
        WW++; // No. Active users exist
        cond_wait(&okToWrite, &lock); // Sleep on
cond var
        WW--; // No longer waiting
    }
    AW++;
    release(&lock);

    AccessDBase(ReadWrite);

    acquire(&lock);
    AW--;
    if (WW > 0) {
        cond_signal(&okToWrite);
    } else if (WR > 0) {
        cond_broadcast(&okToRead);
    }
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

W1 comes along (R1 and R2 are still accessing dbase)

$AR = 2, WR = 0, AW = 0, WW = 0$

```
Writer() {
    acquire(&lock);
    while ((AW + AR) > 0) { // Is it safe to
write?
        WW++; // No. Active users exist
        cond_wait(&okToWrite, &lock); // Sleep on
cond var
        WW--; // No longer waiting
    }
    AW++;
    release(&lock);

    AccessDBase(ReadWrite);

    acquire(&lock);
    AW--;
    if (WW > 0) {
        cond_signal(&okToWrite);
    } else if (WR > 0) {
        cond_broadcast(&okToRead);
    }
    release(&lock);
}
```



# Simulation of Readers/Writers Solution

---

- W1 comes along (R1 and R2 are still accessing dbase)
- $AR = 2$ ,  $WR = 0$ ,  $AW = 0$ ,  $WW = 1$

```
Writer() {
    acquire(&lock);
    while ((AW + AR) > 0) { // Is it safe to
write?
        WW++; // No. Active users exist
        cond_wait(&okToWrite, &lock); // Sleep on
cond var
        WW--; // No longer waiting
    }
    AW++;
    release(&lock);

    AccessDBase(ReadWrite);

    acquire(&lock);
    AW--;
    if (WW > 0) {
        cond_signal(&okToWrite);
    } else if (WR > 0) {
        cond_broadcast(&okToRead);
    }
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

R3 comes along (R1 and R2 accessing dbase, W1 waiting)

$AR = 2, WR = 0, AW = 0, WW = 1$

```
Reader() {  
    acquire(&lock);  
    while ((AW + WW) > 0) { // Is it safe to read?  
        WR++; // No. Writers exist  
        cond_wait(&okToRead, &lock); // Sleep on cond  
var  
        WR--; // No longer waiting  
    }  
    AR++; // Now we are active!  
    release(&lock);  
  
    AccessDBase(ReadOnly);  
  
    acquire(&lock);  
    AR--;  
    if (AR == 0 && WW > 0)  
        cond_signal(&okToWrite);  
    release(&lock);  
}
```

# Simulation of Readers/Writers Solution

---

R3 comes along (R1 and R2 accessing dbase, W1 waiting)

$AR = 2$ ,  $WR = 0$ ,  $AW = 0$ ,  $WW = 1$

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);

    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&löck);
}
```

# Simulation of Readers/Writers Solution

---

R3 comes along (R1 and R2 accessing dbase, W1 waiting)

AR = 2, WR = 1, AW = 0, WW = 1

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No! Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    lock.release();

    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&löck);
}
```

# Simulation of Readers/Writers Solution

---

R3 comes along (R1, R2 accessing dbase, W1 waiting)

AR = 2, WR = 1, AW = 0, WW = 1

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);

    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond signal(&okToWrite);
    release(&löck);
}
```

# Simulation of Readers/Writers Solution

---

R1 and R2 accessing dbase, W1 and R3 waiting

$AR = 2$ ,  $WR = 1$ ,  $AW = 0$ ,  $WW = 1$

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);

    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
```

Status:

- R1 and R2 still reading
- W1 and R3 waiting on okToWrite and okToRead, respectively

# Simulation of Readers/Writers Solution

---

R2 finishes (R1 accessing dbase, W1 and R3 waiting)

$AR = 2$ ,  $WR = 1$ ,  $AW = 0$ ,  $WW = 1$

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);

    AccessDBase(ReadOnly);
    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

R2 finishes (R1 accessing dbase, W1 and R3 waiting)

AR = 1, WR = 1, AW = 0, WW = 1

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);

    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&lock);
}
```



# Simulation of Readers/Writers Solution

---

R2 finishes (R1 accessing dbase, W1 and R3 waiting)

AR = 1, WR = 1, AW = 0, WW = 1

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);

    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

R2 finishes (R1 accessing dbase, W1 and R3 waiting)

$AR = 1, WR = 1, AW = 0, WW = 1$

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);

    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

R1 finishes (W1 and R3 waiting)

AR = 1, WR = 1, AW = 0, WW = 1

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);

    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

R1 finishes (W1, R3 waiting)

AR = 0, WR = 1, AW = 0, WW = 1

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);

    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

R1 finishes (W1, R3 waiting)

AR = 0, WR = 1, AW = 0, WW = 1

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);

    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

R1 signals a writer (W1 and R3 waiting)

$AR = 0$ ,  $WR = 1$ ,  $AW = 0$ ,  $WW = 1$

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);

    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

W1 gets signal (R3 still waiting)

AR = 0, WR = 1, AW = 0, WW = 1

```
Writer() {
    acquire(&lock);
    while ((AW + AR) > 0) { // Is it safe to
write?
        WW++; // No. Active users exist
        cond_wait(&okToWrite, &lock); // Sleep on
cond var
        WW--; // No longer waiting
    }
    AW++;
    release(&lock);

    AccessDBase(ReadWrite);

    acquire(&lock);
    AW--;
    if (WW > 0) {
        cond_signal(&okToWrite);
    } else if (WR > 0) {
        cond_broadcast(&okToRead);
    }
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

W1 gets signal (R3 still waiting)

AR = 0, WR = 1, AW = 0, WW = 0

```
Writer() {
    acquire(&lock);
    while ((AW + AR) > 0) { // Is it safe to
write?
        WW++; // No Active users exist
        cond_wait(&okToWrite, &lock); // Sleep on
cond var
        WW--; // No longer waiting
    }
    AW++;
    release(&lock);

    AccessDBase(ReadWrite);

    acquire(&lock);
    AW--;
    if (WW > 0) {
        cond_signal(&okToWrite);
    } else if (WR > 0) {
        cond_broadcast(&okToRead);
    }
    release(&lock);
}
```



# Simulation of Readers/Writers Solution

---

W1 gets signal (R3 still waiting)

AR = 0, WR = 1, AW = 1, WW = 0

```
Writer() {
    acquire(&lock);
    while ((AW + AR) > 0) { // Is it safe to
write?
        WW++; // No. Active users exist
        cond_wait(&okToWrite, &lock); // Sleep on
cond var
        WW--; // No longer waiting
    }
    AW++;
    release(&lock);

    AccessDBase(ReadWrite);

    acquire(&lock);
    AW--;
    if (WW > 0) {
        cond_signal(&okToWrite);
    } else if (WR > 0) {
        cond_broadcast(&okToRead);
    }
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

W1 accessing dbase (R3 still waiting)

AR = 0, WR = 1, AW = 1, WW = 0

```
Writer() {
    acquire(&lock);
    while ((AW + AR) > 0) {    // Is it safe to
write?
        WW++;    // No. Active users exist
        cond_wait(&okToWrite, &lock);    // Sleep on
cond var
        WW--;    // No longer waiting
    }
    AW++;
    release(&lock);
    AccessDBase(ReadWrite);

    acquire(&lock);
    AW--;
    if (WW > 0) {
        cond_signal(&okToWrite);
    } else if (WR > 0) {
        cond_broadcast(&okToRead);
    }
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

W1 finishes (R3 still waiting)

AR = 0, WR = 1, AW = 1, WW = 0

```
Writer() {
    acquire(&lock);
    while ((AW + AR) > 0) { // Is it safe to
write?
        WW++; // No. Active users exist
        cond_wait(&okToWrite, &lock); // Sleep on
cond var
        WW--; // No longer waiting
    }
    AW++;
    release(&lock);

    AccessDBase(ReadWrite);
    acquire(&lock);
    AW--;
    if (WW > 0) {
        cond_signal(&okToWrite);
    } else if (WR > 0) {
        cond_broadcast(&okToRead);
    }
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

W1 finishes (R3 still waiting)

AR = 0, WR = 1, AW = 0, WW = 0

```
Writer() {
    acquire(&lock);
    while ((AW + AR) > 0) { // Is it safe to
write?
        WW++; // No. Active users exist
        cond_wait(&okToWrite, &lock); // Sleep on
cond var
        WW--; // No longer waiting
    }
    AW++;
    release(&lock);

    AccessDBase(ReadWrite);

    acquire(&lock);
    AW--;
    if (WW > 0) {
        cond_signal(&okToWrite);
    } else if (WR > 0) {
        cond_broadcast(&okToRead);
    }
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

W1 finishes (R3 still waiting)

AR = 0, WR = 1, AW = 0, WW = 0

```
Writer() {
    acquire(&lock);
    while ((AW + AR) > 0) { // Is it safe to
write?
        WW++; // No. Active users exist
        cond_wait(&okToWrite, &lock); // Sleep on
cond var
        WW--; // No longer waiting
    }
    AW++;
    release(&lock);

    AccessDBase(ReadWrite);

    acquire(&lock);
    AW--;
    if (WW > 0) {
        cond_signal(&okToWrite);
    } else if (WR > 0) {
        cond_broadcast(&okToRead);
    }
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

W1 signaling readers (R3 still waiting)

AR = 0, WR = 1, AW = 0, WW = 0

```
Writer() {
    acquire(&lock);
    while ((AW + AR) > 0) { // Is it safe to
write?
        WW++; // No. Active users exist
        cond_wait(&okToWrite, &lock); // Sleep on
cond var
        WW--; // No longer waiting
    }
    AW++;
    release(&lock);

    AccessDBase(ReadWrite);

    acquire(&lock);
    AW--;
    if (WW > 0) {
        cond_signal(&okToWrite);
    } else if (WR > 0) {
        cond_broadcast(&okToRead);
    }
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

R3 gets signal (no waiting threads)

AR = 0, WR = 1, AW = 0, WW = 0

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);

    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

R3 gets signal (no waiting threads)

AR = 0, WR = 0, AW = 0, WW = 0

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);

    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&lock);
}
```



# Simulation of Readers/Writers Solution

---

R3 accessing dbase (no waiting threads)

AR = 1, WR = 0, AW = 0, WW = 0

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);
    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

R3 finishes (no waiting threads)

AR = 1, WR = 0, AW = 0, WW = 0

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);

    AccessDBase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&lock);
}
```

# Simulation of Readers/Writers Solution

---

R3 finishes (no waiting threads)

AR = 0, WR = 0, AW = 0, WW = 0

```
Reader() {
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);

    AccessDbase(ReadOnly);

    acquire(&lock);
    AR--;
    if (AR == 0 && WW > 0)
        cond_signal(&okToWrite);
    release(&lock);
}
```

# Questions

---

Can readers starve? Consider Reader() entry code:

```
while ((AW + WW) > 0) { // Is it safe to read?
    WR++; // No. Writers exist
    cond_wait(&okToRead, &lock); // Sleep on
cond var
    WR--; // No longer waiting
}
AR++; // Now we are active!
```

What if we erase the condition check in Reader exit?

```
AR--; // No longer active
if (AR == 0 && WW > 0) // No other active
readers
    cond_signal(&okToWrite); // Wake up one
writer
```

# Questions

---

Further, what if we turn the `signal()` into `broadcast()`

```
AR--; // No longer active  
cond_broadcast(&okToWrite); // Wake up sleepers
```

Finally, what if we use only one condition variable (call it “`okContinue`”) instead of two separate ones?

- Both readers and writers sleep on this variable
- Must use `broadcast()` instead of `signal()`

# Code for a Reader

---

```
Reader() {
    // First check self into system
    acquire(&lock);
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        cond_wait(&okToRead, &lock); // Sleep on cond
var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    release(&lock);
    // Perform actual read-only access
    AccessDatabase(ReadOnly);
    // Now, check out of system
    acquire(&lock);
    AR--; // No longer active
    if (AR == 0 && WW > 0) // No other active readers
        cond_signal(&okToWrite); // Wake up one writer
    release(&lock);
}
```

# Code for a Writer

---

```
Writer() {
    // First check self into system
    acquire(&lock);
    while ((AW + AR) > 0) { // Is it safe to write?
        WW++; // No. Active users exist
        cond_wait(&okToWrite, &lock); // Sleep on
cond var
        WW--; // No longer waiting
    }
    AW++; // Now we are active!
    release(&lock);
    // Perform actual read/write access
    AccessDatabase(ReadWrite);
    // Now, check out of system
    acquire(&lock);
    AW--; // No longer active
    if (WW > 0) { // Give priority to writers
        cond_signal(&okToWrite); // Wake up one writer
    } else if (WR > 0) { // Otherwise, wake reader
        cond_broadcast(&okToRead); // Wake all
readers
    }
    release(&lock);
}
```

# Mesa Monitor Conclusion

---

Monitors represent the synchronization logic of the program

- Wait if necessary
- Signal when change something so any waiting threads can proceed

```
lock
while (need to wait) {
    condvar.wait();
}
unlock
```

do something so no need to wait

```
lock

condvar.signal();

unlock
```

} Check and/or update  
state variables  
Wait if necessary

} Check and/or update  
state variables



# C-Language Support for Synchronization

---

C language: Pretty straightforward synchronization

Just make sure you know *all* the code paths out of a critical section

```
int Rtn() {  
    acquire(&lock);  
    ...  
    if (exception) {  
        release(&lock);  
        return errReturnCode;  
    }  
    ...  
    release(&lock);  
    return OK;  
}
```

# Concurrency and Synchronization in C

---

Harder with more locks

```
void Rtn() {  
    lock1.acquire();  
    ...  
    if (error) {  
        lock1.release();  
        return;  
    }  
    ...  
    lock2.acquire();  
    ...  
    if (error) {  
        lock2.release();  
        lock1.release();  
        return;  
    }  
    ...  
    lock2.release();  
    lock1.release();  
}
```

# C++ Language Support for Synchronization

---

Languages with exceptions like C++

- Languages that support exceptions are problematic (easy to make a non-local exit without releasing lock)

```
void Rtn() {  
    lock.acquire();  
    ...  
    DoFoo();  
    ...  
    lock.release();  
}  
void DoFoo() {  
    ...  
    if (exception) throw errException;  
    ...  
}
```

- Notice that an exception in DoFoo() will exit without releasing the lock!

# C++ Language Support for Synchronization (con't)

---

Must catch all exceptions in critical sections

- Catch exceptions, release lock, and re-throw exception:

```
void Rtn() {  
    lock.acquire();  
    try {  
        ...  
        DoFoo();  
        ...  
    } catch (...) {           // catch exception  
        lock.release();      // release lock  
        throw;               // re-throw the  
exception  
    }  
    lock.release();  
}  
void DoFoo() {  
    ...  
    if (exception) throw errException;  
    ...  
}
```

## Much better: C++ Lock Guards

---

```
#include <mutex>
int global_i = 0;
std::mutex global_mutex;

void safe_increment() {
    std::lock_guard<std::mutex> lock(global_mutex);
    ...
    global_i++;
    // Mutex released when 'lock' goes out of scope
}
```

# Python with Keyword

---

More versatile than we show here (can be used to close files, database connections, etc.)

```
lock = threading.Lock()
...
with lock: # Automatically calls acquire()
    some_var += 1
...
# release() called however we leave block
```

# Java synchronized Keyword

---

Every Java object has an associated lock:

- Lock is acquired on entry and released on exit from a **synchronized** method
- Lock is properly released if exception occurs inside a **synchronized** method
  - Mutex execution of synchronized methods (beware deadlock)

```
class Account {  
    private int balance;  
  
    // object constructor  
    public Account (int initialBalance) {  
        balance = initialBalance;  
    }  
    public synchronized int getBalance() {  
        return balance;  
    }  
    public synchronized void deposit(int amount) {  
        balance += amount;  
    }  
}
```

# Java Support for Monitors

---

Along with a lock, every object has a single condition variable associated with it

To wait inside a synchronized method:

- void wait();
- void wait(long timeout);

To signal while in a synchronized method:

- void notify();
- void notifyAll();



# Where are we going with synchronization?

---

Programs	Shared Programs
Higher-level API	Locks   Semaphores   Monitors   Send/Receive
Hardware	Load/Store   Disable Ints   Test&Set   Compare&Swap

Implement various higher-level synchronization primitives using atomic operations

# Topic Breakdown

---

Virtualizing the CPU

Process Abstraction and API

Threads and Concurrency

Scheduling

Virtualizing Memory

Virtual Memory

Paging

Persistence

IO devices

File Systems

Distributed Systems

Challenges with distribution

Data Processing & Storage

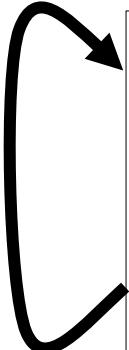
# Goals for Today

---

- What is scheduling?
- What makes a good scheduling policy?
- What are existing schedulers and how do they perform?

# The Scheduling Loop!

---

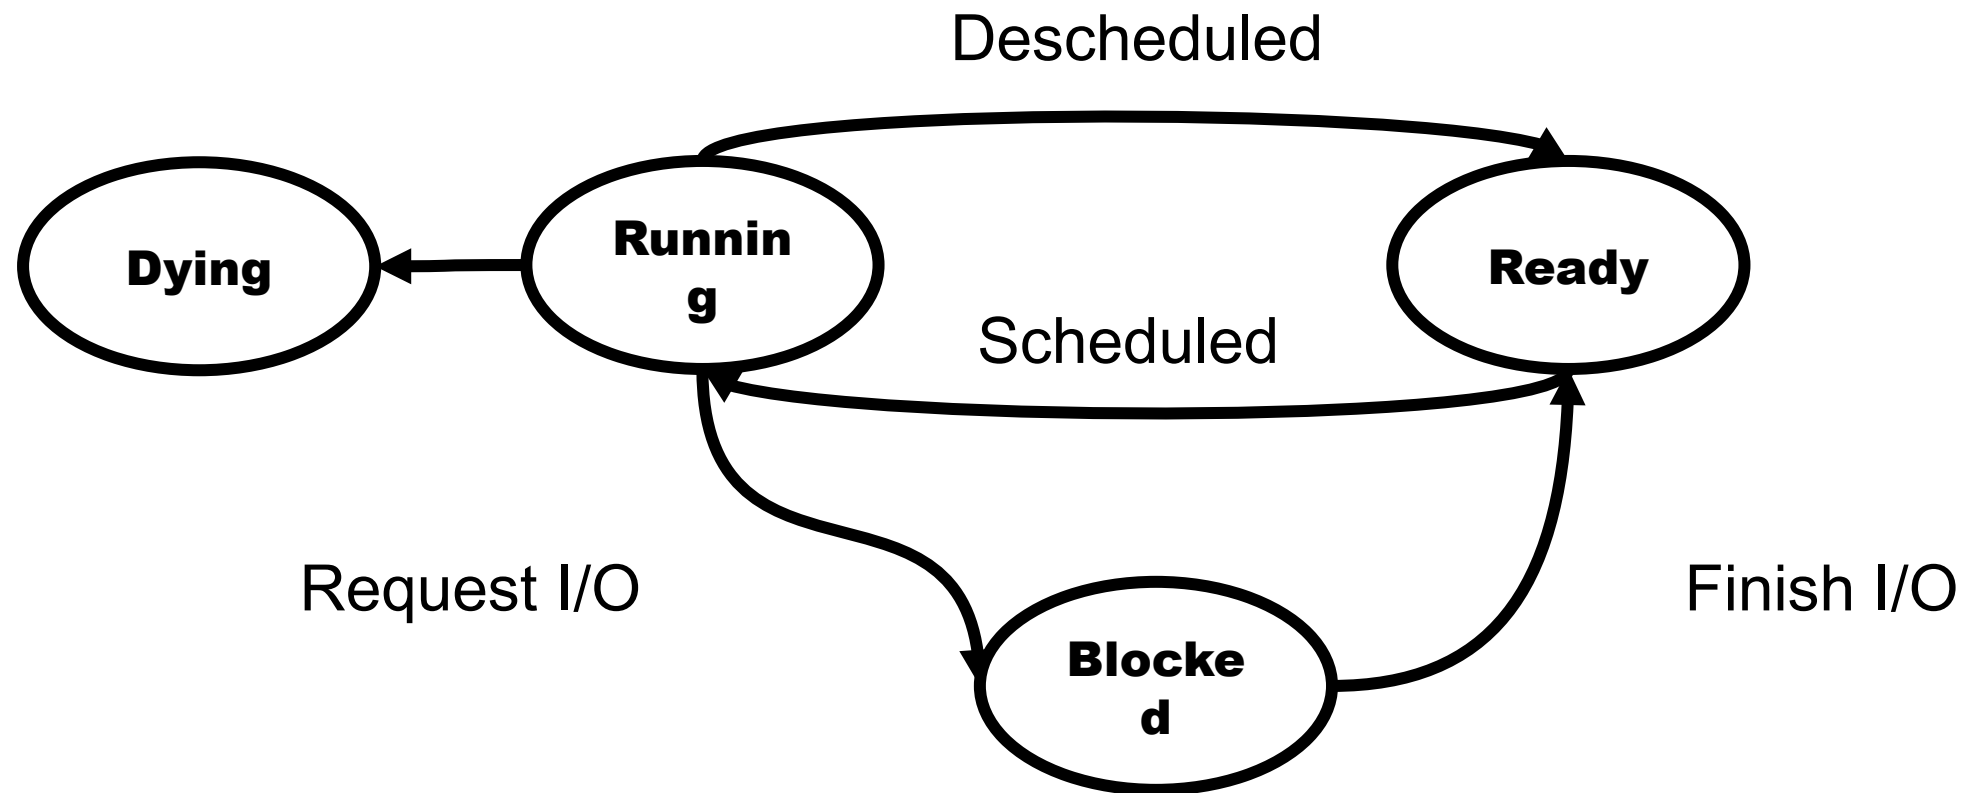


```
if (readyThreads(TCBs) ) {  
    nextTCB =  
    selectThread(TCBs);  
    run(nextTCB);  
} else {  
    run_idle_thread();  
}
```


1. Which task to run next?
2. How frequently does this loop run?
3. What happens if `run` never returns?

# Recall: Thread Life Cycle

---



# Recall: What triggers a scheduling decision?



# What makes a good scheduling policy?

---

## **A hopeless Queue.**

**The Queue For the UK Queen**

**6 miles (10 KM) long.**

**Visible from Space.**

## **A bad but more realistic queue.**

**The DMV**

# What makes a good scheduling policy?

**What does the DMV  
care about?**



**What do individual  
users care about?**





# Important Performance Metrics

---

Response time (or latency).

User-perceived time to do some task

Throughput.

The rate at which tasks are completed

Scheduling overhead.

The time to switch from one task to another.

Predictability.

Variance in response times for repeated requests.

# Important Performance Metrics

---

## Fairness

Equality in the performance perceived by one task

## Starvation

The lack of progress for one task, due to resources being allocated to different tasks

# Sample Scheduling Policies

---

Assume DMV job A takes 1 second, job B takes 2 days

**Policy Idea:** Only ever schedule users with Job A

What is the metric we are optimizing?

A) Throughput B) Latency C) Predictability D) Low-Overhead

Can the schedule lead to starvation?

A) Yes B) No

Is the schedule fair?

A) Yes B) No

# Sample Scheduling Policies

---

Assume DMV consists only of jobs of type A.

**Policy Idea:** Schedule jobs randomly

What is the metric we are optimizing?

A) Throughput B) Latency C) Predictability D) Low-Overhead

Can the schedule lead to starvation?

A) Yes B) No

Is the schedule fair?

A) Yes B) No

# Sample Scheduling Policies

---

Assume DMV consists only of 100 different types of jobs. Some jobs need Clerk A, some Clerks A&B, others Clerk C.

**Policy Idea** Every time schedule a job, compute all possible orderings of jobs, pick one that finishes quickest

What is the metric we are optimizing?

A) Throughput B) Latency C) Predictability D) Low-Overhead

Can the schedule lead to starvation?

A) Yes B) No

Is the schedule fair?

A) Yes B) No

# Scheduling Policy Goals/Criteria

---

Minimise Response  
Time

Maximise Throughput

While remaining fair and starvation-free

# Useful metrics

---

## Waiting time for P

Total Time spent waiting for *CPU*

## Average waiting time

Average of all processes' wait time

## Response Time for P

Time to when process gets first scheduled

## Completion time

Waiting time + Run time

## Average completion time

Average of all processes' completion time

# Assumptions

---

Threads are independent!

One thread = One User

Unrealistic but simplify the problem so it can be solved

Only look at **work-conserving** scheduler  
=> Never leave processor idle if work to do



# Workload Assumptions

---

A workload is a set of tasks for some system to perform, including how long tasks last and when they arrive

## Compute-Bound

Tasks that primarily perform  
compute

Fully utilise CPU

## IO Bound

Mostly wait for IO, limited  
compute

Often in the  
Blocked state

# First-Come, First-Served (FCFS)

---

Run tasks in order of arrival.

Run task until completion (or blocks on IO).  
No preemption

This is the DMV model.

Also called FIFO

# First-Come, First-Served (FCFS)

---


Process

Burst Time

$P_1$       3

$P_2$       3

$P_3$       24



What is the average completion time?

( )

What is the average waiting time?

( )

# First-Come, First-Served (FCFS)

---


Process

Burst Time

$P_3$  24

$P_2$  3

$P_1$  3



What is the average completion time?

( )

What is the average waiting time?

( )

# The Convoy Effect

---

FIFO/FCFS very sensitive to arrival order

*Convoy effect*

Short process stuck behind long process

Lots of small tasks build up behind long tasks

FIFO is **non-preemptible**

# The Convoy Effect

---

FIFO/FCFS very sensitive to arrival order

*Convoy effect*

Short process stuck behind long process

Lots of small tasks build up behind long tasks

FIFO is *non-preemptible*



# The Convoy Effect

---


FIFO/FCFS very sensitive to arrival order

*Convoy effect*

Short process stuck behind long process

Lots of small tasks build up behind long tasks

FIFO is *non-preemptible*



# The Convoy Effect

---


FIFO/FCFS very sensitive to arrival order

*Convoy effect*

Short process stuck behind long process

Lots of small tasks build up behind long tasks

FIFO is *non-preemptible*





# The Convoy Effect

---


FIFO/FCFS very sensitive to arrival order

*Convoy effect*

Short process stuck behind long process

Lots of small tasks build up behind long tasks

FIFO is *non-preemptible*



# The Convoy Effect

---


FIFO/FCFS very sensitive to arrival order

*Convoy effect*

Short process stuck behind long process

Lots of small tasks build up behind long tasks

FIFO is *non-preemptible*



# The Convoy Effect

---


FIFO/FCFS very sensitive to arrival order

*Convoy effect*

Short process stuck behind long process

Lots of small tasks build up behind long tasks

FIFO is *non-preemptible*



# The Convoy Effect

---


FIFO/FCFS very sensitive to arrival order

*Convoy effect*

Short process stuck behind long process

Lots of small tasks build up behind long tasks

FIFO is *non-preemptible*



# The Convoy Effect

---

FIFO/FCFS very sensitive to arrival order

*Convoy effect*

Short process stuck behind long process

Lots of small tasks build up behind long tasks

FIFO is *non-preemptible*

**Can FIFO lead to  
starvation?**



# FCFS/FIFO Summary

---

## The good

Simple  
Low Overhead  
No Starvation

## The bad

**Sensitive to arrival  
order (poor  
predictability)**

## The ugly

**Convoy Effect.  
Bad for Interactive  
Tasks**

# Shortest Job First

---


How can we minimise average completion time?

By scheduling jobs in order of  
estimated completion time

This is the “10 items or less” line at Safeway

# Shortest Job First

---



What is the average completion time? ()

Can prove that SJF generates optimal average completion time if all jobs arrive at the same time




# Are we done?

---

Can SJF lead to starvation?

Yes

Any scheduling policy that always favours a  
**fixed property** for scheduling leads to  
starvation




# Are we done?

---

Can SJF lead to starvation?

Yes

Any scheduling policy that always favours a  
**fixed property** for scheduling leads to  
starvation




# Are we done?

---

Can SJF lead to starvation?

Yes

Any scheduling policy that always favours a  
**fixed property** for scheduling leads to  
starvation




# Are we done?

---

Is SFJ subject to the convoy effect?

Yes

Any **non-preemptible** scheduling policy suffers from convoy effect




# Are we done?

---

Is SFJ subject to the convoy effect?

Yes

Any **non-preemptible** scheduling policy suffers from convoy effect



# SJF Summary

---

## The good

Optimal Average Completion  
Time when jobs arrive  
simultaneously

## The bad

Sensitive to arrival order  
(poor predictability)

## The ugly

Can lead to starvation!

Requires knowing duration of job

# Shortest Time to Completion First (STCF)

---

Introduce the notion of **preemption**

A running task can be de-scheduled before completion.

STCF

Schedule the task with the **least amount of time left**

# Shortest Time to Completion First (STCF)

---

## STCF

Schedule the task with the **least amount of time left**


<u>Process</u>	<u>Burst Time (left)</u>	<u>Arrival Time</u>
$P_1$	3	10
$P_2$	6	1
$P_3$	24	0
$P_4$	16	20



# Shortest Time to Completion First (STCF)

---

<u>Process</u>	<u>Burst Time (left)</u>	<u>Arrival Time</u>
$P_1$	3	10
$P_2$	6	1
$P_3$	24	0
$P_4$	16	18




0                      1

# Shortest Time to Completion First (STCF)

---

<u>Process</u>	<u>Burst Time (left)</u>	<u>Arrival Time</u>
$P_1$	3	10
$P_2$	6	1
$P_3$	23	0
$P_4$	16	18



# Shortest Time to Completion First (STCF)

---


<u>Process</u>	<u>Burst Time (left)</u>	<u>Arrival Time</u>
$P_1$	3	10
$P_2$	0	1
$P_3$	23	0
$P_4$	16	20



# Shortest Time to Completion First (STCF)

---

<u>Process</u>	<u>Burst Time (left)</u>	<u>Arrival Time</u>
$P_1$	3	10
$P_2$	0	1
$P_3$	20	0
$P_4$	16	18



# Shortest Time to Completion First (STCF)

---


<u>Process</u>	<u>Burst Time (left)</u>	<u>Arrival Time</u>
----------------	--------------------------	---------------------

$P_1$	0	10
-------	---	----

$P_2$	0	1
-------	---	---

$P_3$	15	0
-------	----	---


$P_4$	16	18
-------	----	----



# Shortest Time to Completion First (STCF)

---


<u>Process</u>	<u>Burst Time (left)</u>	<u>Arrival Time</u>
$P_1$	0	10
$P_2$	0	1
$P_3$	0	0
$P_4$	15	18



# Shortest Time to Completion First (STCF)

---

<u>Process</u>		<u>Burst Time (left)</u>	<u>Arrival Time</u>
$P_1$	0		10
$P_2$	0		1
$P_3$	0		0
$P_4$	15		18



# Are we done?

---

Can STCF lead to starvation?

Yes

Any scheduling policy that always favours a  
**fixed property** for scheduling leads starvation

No change!



# Are we done?

---

Is STCF subject to the convoy effect?

No!

STCF is a preemptible policy

# STCF Summary

---

## The good

Optimal Average Completion  
Time Always

## The bad

## The ugly

Can lead to starvation!

Requires knowing duration of job

# Taking a step back

---

Property	FCFS	SJF	STCF
Optimise Average Completion Time		✓	✓
Prevent Starvation			
Prevent Convoy Effect	✓		
Psychic Skills Not Needed			✓



Can we design a non-psychic, starvation-free scheduler with good response time?

# Round-Robin Scheduling

---

RR runs a job for a **time slice**  
(a scheduling quantum)

Once time slice over,  
Switch to next job in ready queue.  
=> Called time-slicing

## RR with Time Quantum = 20


---

<u>Process</u>	<u>Burst Time</u>
$P_1$	53
$P_2$	8
$P_3$	68
$P_4$	24

# RR with Time Quantum = 20

---


<u>Process</u>	<u>Burst Time</u>
$P_1$	53 => 33
$P_2$	8
$P_3$	68
$P_4$	24



# RR with Time Quantum = 20

---


<u>Process</u>	<u>Burst Time</u>
$P_1$	33
$P_2$	8 $\Rightarrow$ 0
$P_3$	68
$P_4$	24



# RR with Time Quantum = 20

---

<u>Process</u>	<u>Burst Time</u>
$P_1$	33
$P_2$	0
$P_3$	68 => 48
$P_4$	24






# RR with Time Quantum = 20

---


<u>Process</u>	<u>Burst Time</u>
$P_1$	33
$P_2$	0
$P_3$	48
$P_4$	24 => 4



# RR with Time Quantum = 20

---


<u>Process</u>	<u>Burst Time</u>
$P_1$	33 => 13
$P_2$	0
$P_3$	48
$P_4$	4



# RR with Time Quantum = 20

---


<u>Process</u>	<u>Burst Time</u>
$P_1$	13
$P_2$	0
$P_3$	48 => 28
$P_4$	4



# RR with Time Quantum = 20

---

<u>Process</u>	<u>Burst Time</u>
$P_1$	13
$P_2$	0
$P_3$	28
$P_4$	4 $\Rightarrow$ 0



# RR with Time Quantum = 20

---

Waiting time


- $P_1 = 0 + (68-20) + (112-88) = 72$
- $P_2 = (20-0) = 20$
- $P_3 = (28-0) + (88-48) + (125-108) + 0 = 85$
- $P_4 = (48-0) + (108-68) = 88$

Average waiting time

()


Average completion time

()




# Decrease Completion Time

- $T_1$ : Burst Length 10     $T_3$ : Burst Length 10
- $T_2$ : Burst Length 5



Average Completion Time =  $(10 + 15 + 25)/3 = 16.7$



Average Completion Time =  $(20 + 10 + 25)/3 = 18.3$

# Switching is not free!

---

Small scheduling quantas lead to  
frequent context switches

- Mode switch overhead
- Trash cache-state

$q$  must be large with respect to context switch,  
otherwise overhead is too high

# Are we done?

---

Can RR lead to starvation?

No

No process waits more than  $(n-1)q$  time units



# Are we done?

---

Can RR suffer from convoy effect?

No

Only run a time-slice at a time

# RR Summary

---

## The good

Bounded response time

## The bad

Completion time can be high  
(stretches out long jobs)

## The ugly

Overhead of context switching

# Taking a step back

---

Property	FCFS	SJF	STCF
Optimise Average Completion Time		✓	✓
Prevent Starvation			
Prevent Convoy Effect	✓		
Psychic Skills Not Needed			✓



# Taking a step back

---

Property	FCFS	SJF	STCF	RR
Optimise Average Completion Time		✓	✓	
Optimise Average Response Time				✓
Prevent Starvation	✓			✓
Prevent Convoy Effect			✓	✓
Psychic Skills Not Needed	✓			✓

# FCFS and Round Robin Showdown

---


Assuming zero-cost context-switching time,  
is RR always better than FCFS?

10 jobs, each take 100s of CPU time  
RR scheduler quantum of 1s  
All jobs start at the same time

Job #	FIFO	RR
1	100	991
2	200	992
...	...	...
9	900	999
10	1000	1000

# Earlier Example with Different Time Quantum

Best FCFS:



Quantum	P1	P2	P3	P4	Average
Best FCFS	85	8	16	32	69.5
Q=1	137	30	153	81	100.5
Q=5	135	28	153	82	99.5
Q=8	133	16	153	80	99,5
Q=10	135	18	153	92	104.5
Q=20	125	28	153	112	104.5
Worst FCFS	121	153	68	145	121.75