

# CS162

## Operating Systems and Systems Programming

### Lecture 10

## Scheduling

### Core Concepts and Classic Policies

Professor Natacha Crooks

<https://cs162.org/>

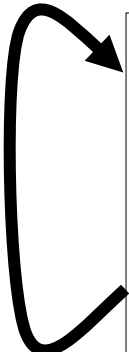
# Goals for Today

---

- What is scheduling?
- What makes a good scheduling policy?
- What are existing schedulers and how do they perform?

# The Scheduling Loop!

---

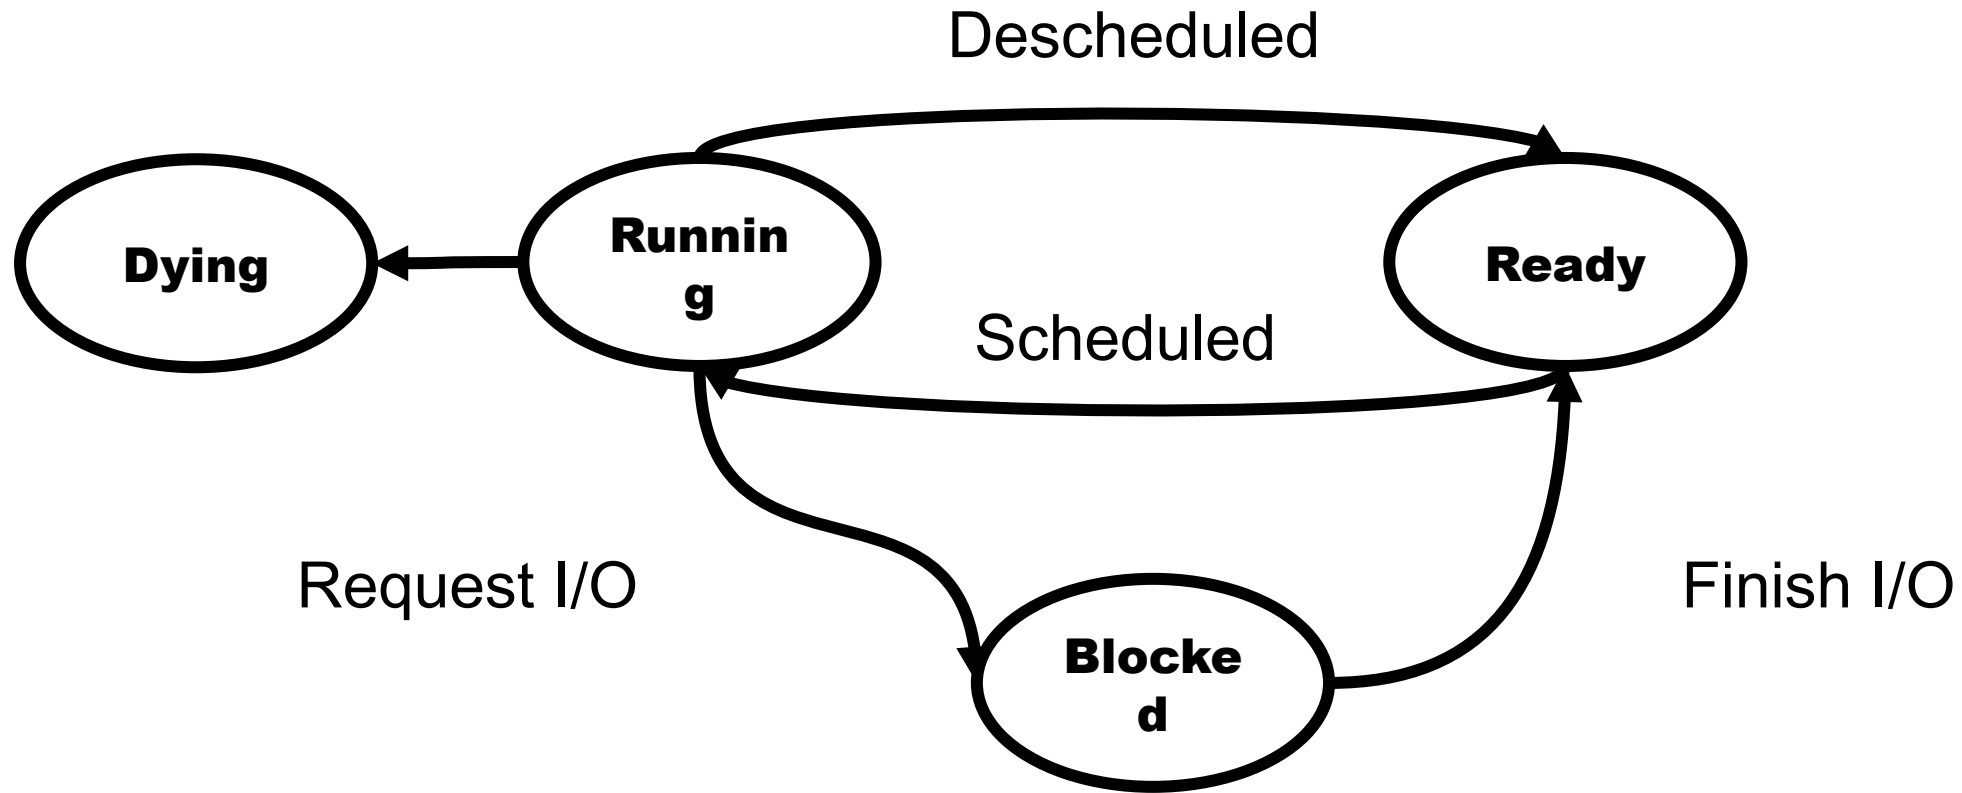


```
if (readyThreads(TCBs) ) {  
    nextTCB =  
    selectThread(TCBs);  
    run(nextTCB);  
} else {  
    run_idle_thread();  
}
```

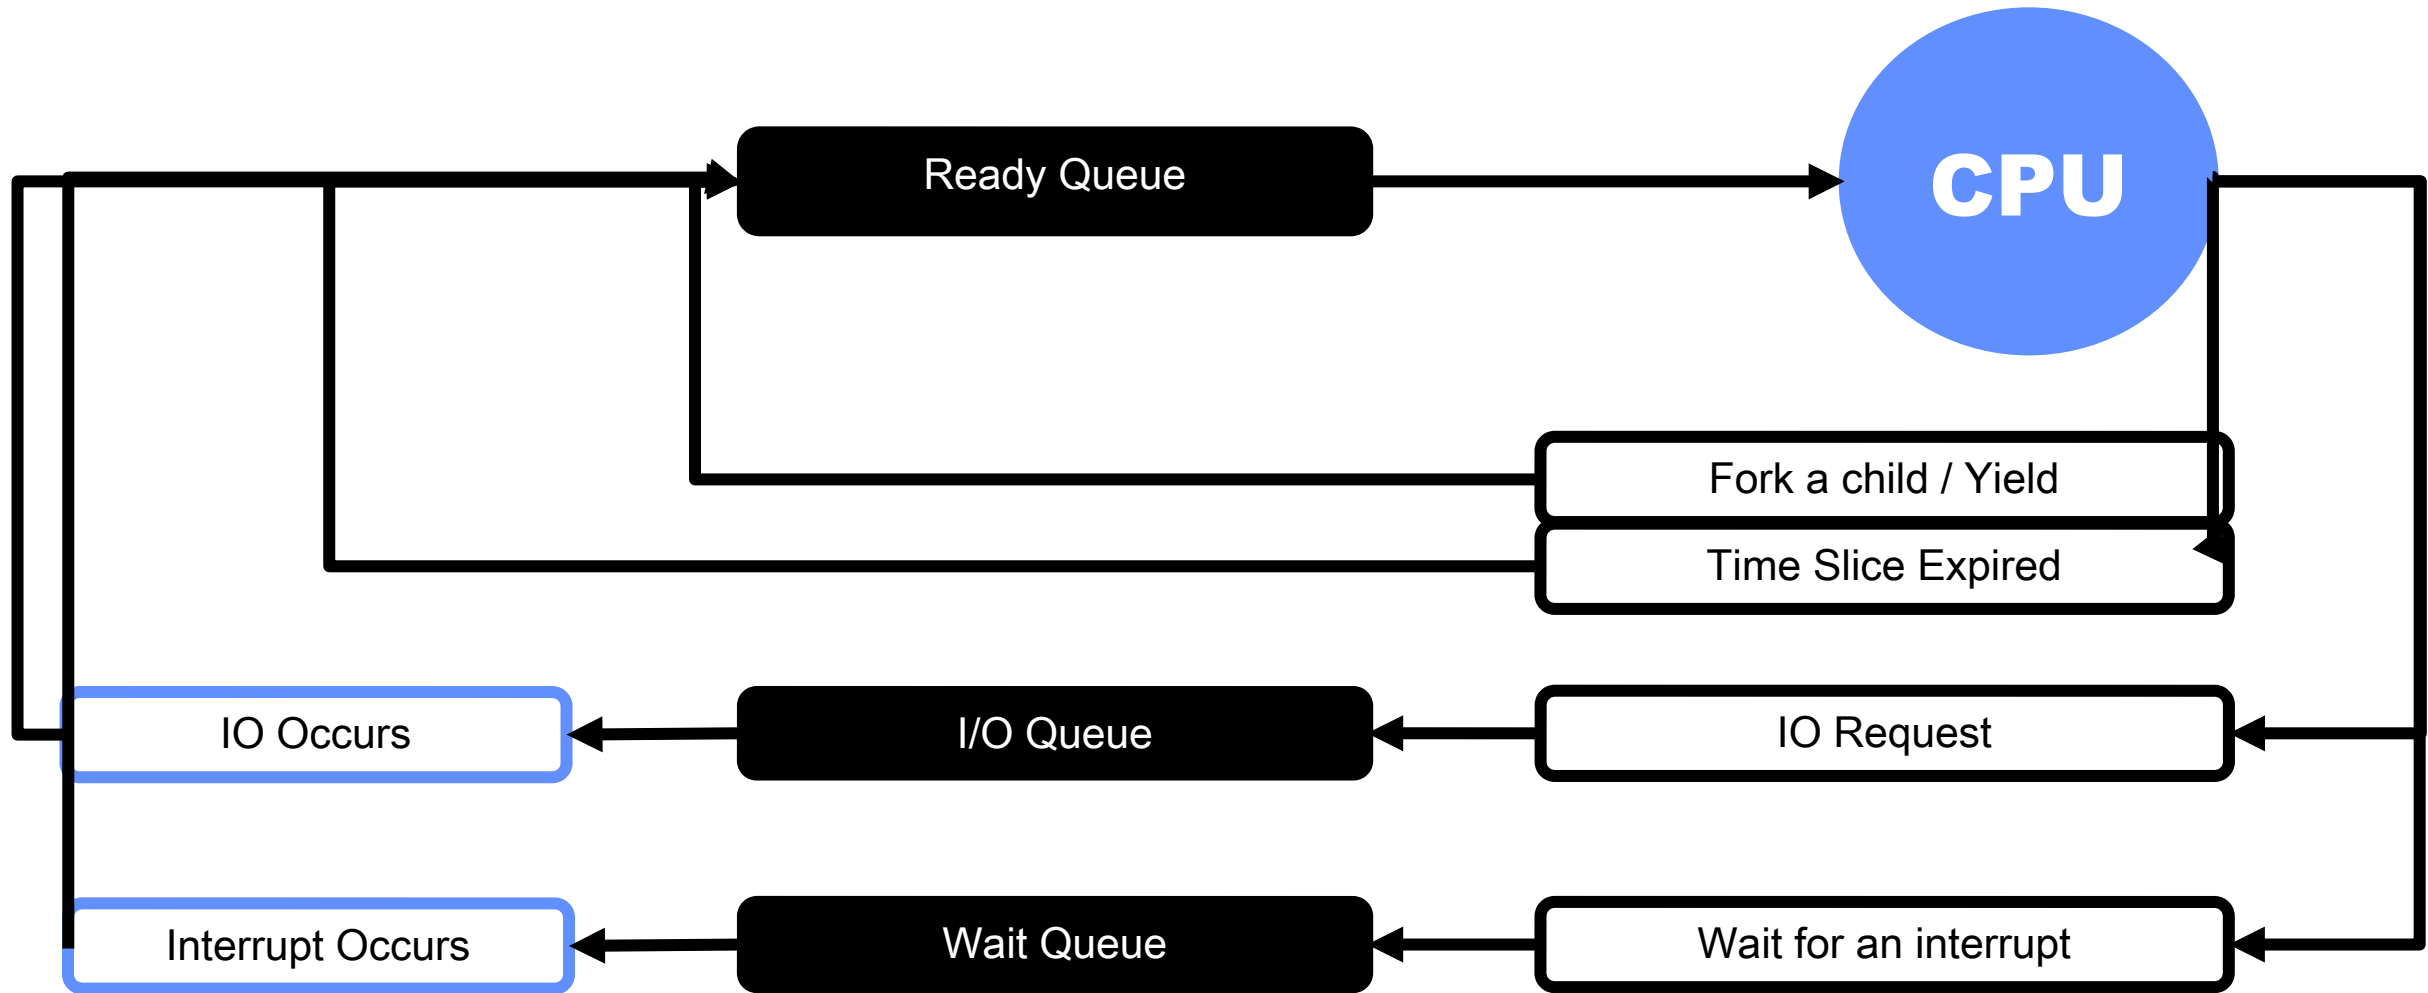
1. Which task to run next?
2. How frequently does this loop run?
3. What happens if `run` never returns?

# Recall: Thread Life Cycle

---



# Recall: What triggers a scheduling decision?



# What makes a good scheduling policy?

---

## **A hopeless Queue.**

**The Queue For the UK Queen**

**6 miles (10 KM) long.**

**Visible from Space.**

## **A bad but more realistic queue.**

**The DMV**

# What makes a good scheduling policy?

**What does the DMV  
care about?**



**What do individual  
users care about?**



# Important Performance Metrics

---

Response time (or latency).

User-perceived time to do some task

Throughput.

The rate at which tasks are completed

Scheduling overhead.

The time to switch from one task to another.

Predictability.

Variance in response times for repeated requests.



# Important Performance Metrics

---

## Fairness

Equality in the performance perceived by one task

## Starvation

The lack of progress for one task, due to resources being allocated to different tasks

# Sample Scheduling Policies

---

Assume DMV job A takes 1 second, job B takes 2 days

**Policy Idea:** Only ever schedule users with Job A

What is the metric we are optimizing?

A) Throughput B) Latency C) Predictability D) Low-Overhead

Can the schedule lead to starvation?

A) Yes B) No

Is the schedule fair?

A) Yes B) No

# Sample Scheduling Policies

---

Assume DMV consists only of jobs of type A.

**Policy Idea:** Schedule jobs randomly

What is the metric we are optimizing?

A) Throughput B) Latency C) Predictability D) Low-Overhead

Can the schedule lead to starvation?

A) Yes B) No

Is the schedule fair?

A) Yes B) No

# Sample Scheduling Policies

---

Assume DMV consists only of 100 different types of jobs. Some jobs need Clerk A, some Clerks A&B, others Clerk C.

**Policy Idea** Every time schedule a job, compute all possible orderings of jobs, pick one that finishes quickest

What is the metric we are optimizing?

A) Throughput B) Latency C) Predictability D) Low-Overhead

Can the schedule lead to starvation?

A) Yes B) No

Is the schedule fair?

A) Yes B) No

# Scheduling Policy Goals/Criteria

---

Minimise Latency

Maximise Throughput

While remaining fair and starvation-free

# Useful metrics

---

## Waiting time for P

Total Time spent waiting for *CPU*

## Average waiting time

Average of all processes' wait time

## Response Time for P

Time to when process gets first scheduled

## Completion time

Waiting time + Run time

## Average completion time

Average of all processes' completion time

# Assumptions

---

Threads are independent!

One thread = One User

Unrealistic but simplify the problem so it can be solved

Only look at **work-conserving** scheduler  
=> Never leave processor idle if work to do

# Workload Assumptions

---

A workload is a set of tasks for some system to perform, including how long tasks last and when they arrive

## Compute-Bound

Tasks that primarily perform  
compute

Fully utilise CPU

## IO Bound

Mostly wait for IO, limited  
compute

Often in the  
Blocked state



# First-Come, First-Served (FCFS)

---

Run tasks in order of arrival.

Run task until completion (or blocks on IO).  
No preemption

This is the DMV model.

Also called FIFO

# First-Come, First-Served (FCFS)

---


Process

Burst Time

$P_1$       3

$P_2$       3

$P_3$       24



What is the average completion time?

( )

What is the average waiting time?

( )

# First-Come, First-Served (FCFS)

---


Process

Burst Time

$P_3$  24

$P_2$  3

$P_1$  3



What is the average completion time?

( )

What is the average waiting time?

( )

# The Convoy Effect

---

FIFO/FCFS very sensitive to arrival order

*Convoy effect*

Short process stuck behind long process

Lots of small tasks build up behind long tasks

FIFO is **non-preemptible**

# The Convoy Effect

---

FIFO/FCFS very sensitive to arrival order

*Convoy effect*

Short process stuck behind long process

Lots of small tasks build up behind long tasks

FIFO is *non-preemptible*



# The Convoy Effect

---


FIFO/FCFS very sensitive to arrival order

*Convoy effect*

Short process stuck behind long process

Lots of small tasks build up behind long tasks

FIFO is *non-preemptible*



# The Convoy Effect

---


FIFO/FCFS very sensitive to arrival order

*Convoy effect*

Short process stuck behind long process

Lots of small tasks build up behind long tasks

FIFO is *non-preemptible*



# The Convoy Effect

---


FIFO/FCFS very sensitive to arrival order

*Convoy effect*

Short process stuck behind long process

Lots of small tasks build up behind long tasks

FIFO is *non-preemptible*





# The Convoy Effect

---


FIFO/FCFS very sensitive to arrival order

*Convoy effect*

Short process stuck behind long process

Lots of small tasks build up behind long tasks

FIFO is *non-preemptible*



# The Convoy Effect

---


FIFO/FCFS very sensitive to arrival order

*Convoy effect*

Short process stuck behind long process

Lots of small tasks build up behind long tasks

FIFO is *non-preemptible*



# The Convoy Effect

---


FIFO/FCFS very sensitive to arrival order

*Convoy effect*

Short process stuck behind long process

Lots of small tasks build up behind long tasks

FIFO is *non-preemptible*



# The Convoy Effect

---

FIFO/FCFS very sensitive to arrival order


*Convoy effect*

Short process stuck behind long process

Lots of small tasks build up behind long tasks

FIFO is *non-preemptible*

**Can FIFO lead to  
starvation?**



# FCFS/FIFO Summary

---

## The good

Simple  
Low Overhead  
No Starvation

## The bad

**Sensitive to arrival  
order (poor  
predictability)**

## The ugly

**Convoy Effect.  
Bad for Interactive  
Tasks**

# Shortest Job First


---

How can we minimise average completion time?

By scheduling jobs in order of  
estimated completion time

This is the “10 items or less” line at Safeway

# Shortest Job First



What is the average completion time? ()

Can prove that SJF generates optimal average completion time if all jobs arrive at the same time


# Are we done?

---

Can SJF lead to starvation?

Yes

Any scheduling policy that always favours a  
**fixed property** for scheduling leads to  
starvation






# Are we done?

---

Can SJF lead to starvation?

Yes

Any scheduling policy that always favours a  
**fixed property** for scheduling leads to  
starvation




# Are we done?

---

Can SJF lead to starvation?

Yes

Any scheduling policy that always favours a  
**fixed property** for scheduling leads to  
starvation




# Are we done?

---

Is SFJ subject to the convoy effect?

Yes

Any **non-preemptible** scheduling policy suffers from convoy effect




# Are we done?

---

Is SFJ subject to the convoy effect?

Yes

Any **non-preemptible** scheduling policy suffers from convoy effect



# SJF Summary

---

## The good

Optimal Average Completion  
Time when jobs arrive  
simultaneously

## The bad

Sensitive to arrival order  
(poor predictability)

## The ugly

Can lead to starvation!

Requires knowing duration of job

# Shortest Time to Completion First (STCF)

---

Introduce the notion of **preemption**

A running task can be de-scheduled before completion.

STCF

Schedule the task with the **least amount of time left**

# Shortest Time to Completion First (STCF)

---

## STCF


Schedule the task with the **least amount of time left**

<u>Process</u>	<u>Burst Time (left)</u>	<u>Arrival Time</u>
$P_1$	3	10
$P_2$	6	1
$P_3$	24	0
$P_4$	16	20

# Shortest Time to Completion First (STCF)

---

<u>Process</u>	<u>Burst Time (left)</u>	<u>Arrival Time</u>
$P_1$	3	10
$P_2$	6	1
$P_3$	24	0
$P_4$	16	18



0                      1



# Shortest Time to Completion First (STCF)

---

<u>Process</u>	<u>Burst Time (left)</u>	<u>Arrival Time</u>
$P_1$	3	10
$P_2$	6	1
$P_3$	23	0
$P_4$	16	18



# Shortest Time to Completion First (STCF)

---

<u>Process</u>	<u>Burst Time (left)</u>	<u>Arrival Time</u>
$P_1$	3	10
$P_2$	0	1
$P_3$	23	0
$P_4$	16	18



# Shortest Time to Completion First (STCF)

---


<u>Process</u>	<u>Burst Time (left)</u>	<u>Arrival Time</u>
$P_1$	3	10
$P_2$	0	1
$P_3$	20	0
$P_4$	16	18



# Shortest Time to Completion First (STCF)

---


<u>Process</u>	<u>Burst Time (left)</u>	<u>Arrival Time</u>
$P_1$	0	10
$P_2$	0	1
$P_3$	20	0
$P_4$	16	18



# Shortest Time to Completion First (STCF)


---

<u>Process</u>	<u>Burst Time (left)</u>	<u>Arrival Time</u>
$P_1$	0	10
$P_2$	0	1
$P_3$	15	0
$P_4$	16	18



# Shortest Time to Completion First (STCF)


<u>Process</u>	<u>Burst Time (left)</u>	<u>Arrival Time</u>
$P_1$	0	10
$P_2$	0	1
$P_3$	0	0
$P_4$	15	18



# Shortest Time to Completion First (STCF)

---

<u>Process</u>	<u>Burst Time (left)</u>	<u>Arrival Time</u>
$P_1$	0	10
$P_2$	0	1
$P_3$	0	0
$P_4$	15	18



# Are we done?

---

Can STCF lead to starvation?

Yes

Any scheduling policy that always favours a  
**fixed property** for scheduling leads starvation

No change!



# Are we done?

---

Is STCF subject to the convoy effect?

No!

STCF is a preemptible policy

# STCF Summary

---

## The good

Optimal Average Completion  
Time Always

## The bad

## The ugly

Can lead to starvation!

Requires knowing duration of job

# Taking a step back

---

Property	FCFS	SJF	STCF
Optimise Average Completion Time		✓	✓
Prevent Starvation			
Prevent Convoy Effect	✓		
Psychic Skills Not Needed			✓



Can we design a non-psychic, starvation-free scheduler with good response time?

# Round-Robin Scheduling

---

RR runs a job for a **time slice**  
(a scheduling quantum)

Once time slice over,  
Switch to next job in ready queue.  
=> Called time-slicing

## RR with Time Quantum = 20


---

<u>Process</u>	<u>Burst Time</u>
$P_1$	53
$P_2$	8
$P_3$	68
$P_4$	24

# RR with Time Quantum = 20

---


<u>Process</u>	<u>Burst Time</u>
$P_1$	53 => 33
$P_2$	8
$P_3$	68
$P_4$	24



# RR with Time Quantum = 20

---


<u>Process</u>	<u>Burst Time</u>
$P_1$	33
$P_2$	8 $\Rightarrow$ 0
$P_3$	68
$P_4$	24



# RR with Time Quantum = 20

---

<u>Process</u>	<u>Burst Time</u>
$P_1$	33
$P_2$	0
$P_3$	68 => 48
$P_4$	24






# RR with Time Quantum = 20

---


<u>Process</u>	<u>Burst Time</u>
$P_1$	33
$P_2$	0
$P_3$	48
$P_4$	24 => 4



# RR with Time Quantum = 20

---


<u>Process</u>	<u>Burst Time</u>
$P_1$	33 => 13
$P_2$	0
$P_3$	48
$P_4$	4



# RR with Time Quantum = 20

---


<u>Process</u>	<u>Burst Time</u>
$P_1$	13
$P_2$	0
$P_3$	48 => 28
$P_4$	4



# RR with Time Quantum = 20

---

<u>Process</u>	<u>Burst Time</u>
$P_1$	13
$P_2$	0
$P_3$	28
$P_4$	4 $\Rightarrow$ 0



# RR with Time Quantum = 20

Waiting time


- $P_1 = 0 + (68-20) + (112-88) = 72$
- $P_2 = (20-0) = 20$
- $P_3 = (28-0) + (88-48) + (125-108) + 0 = 85$
- $P_4 = (48-0) + (108-68) = 88$

Average waiting time

()


Average completion time

()




# Decrease Completion Time

- $T_1$ : Burst Length 10     $T_3$ : Burst Length 10
- $T_2$ : Burst Length 5



Average Completion Time =  $(10 + 15 + 25)/3 = 16.7$



Average Completion Time =  $(20 + 10 + 25)/3 = 18.3$

# Switching is not free!

---

Small scheduling quantas lead to  
frequent context switches

- Mode switch overhead
- Trash cache-state

$q$  must be large with respect to context switch,  
otherwise overhead is too high

# Are we done?

---

Can RR lead to starvation?

No

No process waits more than  $(n-1)q$  time units



# Are we done?

---

Can RR suffer from convoy effect?

No

Only run a time-slice at a time

# RR Summary

---

## The good

Bounded response time

## The bad

Completion time can be high  
(stretches out long jobs)

## The ugly

Overhead of context switching

# Taking a step back

---

Property	FCFS	SJF	STCF
Optimise Average Completion Time		✓	✓
Prevent Starvation			
Prevent Convoy Effect	✓		
Psychic Skills Not Needed			✓



# Taking a step back

---

Property	FCFS	SJF	STCF	RR
Optimise Average Completion Time		✓	✓	
Optimise Average Response Time				✓
Prevent Starvation	✓			✓
Prevent Convoy Effect			✓	✓
Psychic Skills Not Needed	✓			✓

# FCFS and Round Robin Showdown

---


Assuming zero-cost context-switching time,  
is RR always better than FCFS?

10 jobs, each take 100s of CPU time  
RR scheduler quantum of 1s  
All jobs start at the same time

Job #	FIFO	RR
1	100	991
2	200	992
...	...	...
9	900	999
10	1000	1000

# Earlier Example with Different Time Quantum

Best FCFS:



Quantum	P1	P2	P3	P4	Average
Best FCFS	85	8	16	32	69.5
Q=1	137	30	153	81	100.5
Q=5	135	28	153	82	99.5
Q=8	133	16	153	80	99,5
Q=10	135	18	153	92	104.5
Q=20	125	28	153	112	104.5
Worst FCFS	121	153	68	145	121.75