

CS162  
Operating Systems and  
Systems Programming  
Lecture 15

Virtual Memory (2)

Professor Natacha Crooks

<https://cs162.org/>

# Recall: Memory Management Wishlist

---

Memory Protection

Memory Sharing

Flexible Memory Placement

Support for Sparse Addresses

Runtime Lookup Efficiency

Compact Translation Table

# Recall: Increasingly powerful mechanisms

---

**No protection. Living life on the edge**

**Can access all memory**

**Base & Bound**

**Absolute memory addressing.  
Hard to relocate**

**Base & Bound with Relocation**

**Internal fragmentation when  
address space is sparse**

**Segmentation**

**External fragmentation as  
assigning variably sized chunks**

**Paging**

# Paging

---

Divide logical address space of process into fixed sized chunks called **pages**

View physical memory as an array of fixed-sized slots called **page frames**


Each page frame can contain a single virtual-memory page

Pages should be **small** to minimise internal fragmentation (1K-16k)

# How to Implement Simple Paging?

---


Interpret virtual address as two components



# How to Implement Simple Paging?

---

Interpret virtual address as two components



# A (Simplified) Page Table

---

A page table stores  
virtual-to-physical address translations

One page table per process. Lives in memory.

Address stored in the in the Page Table Base Register

PTBR value saved/restored in PCB on context switch

# How to access a byte?

---

Extract **page number** (first p bits)

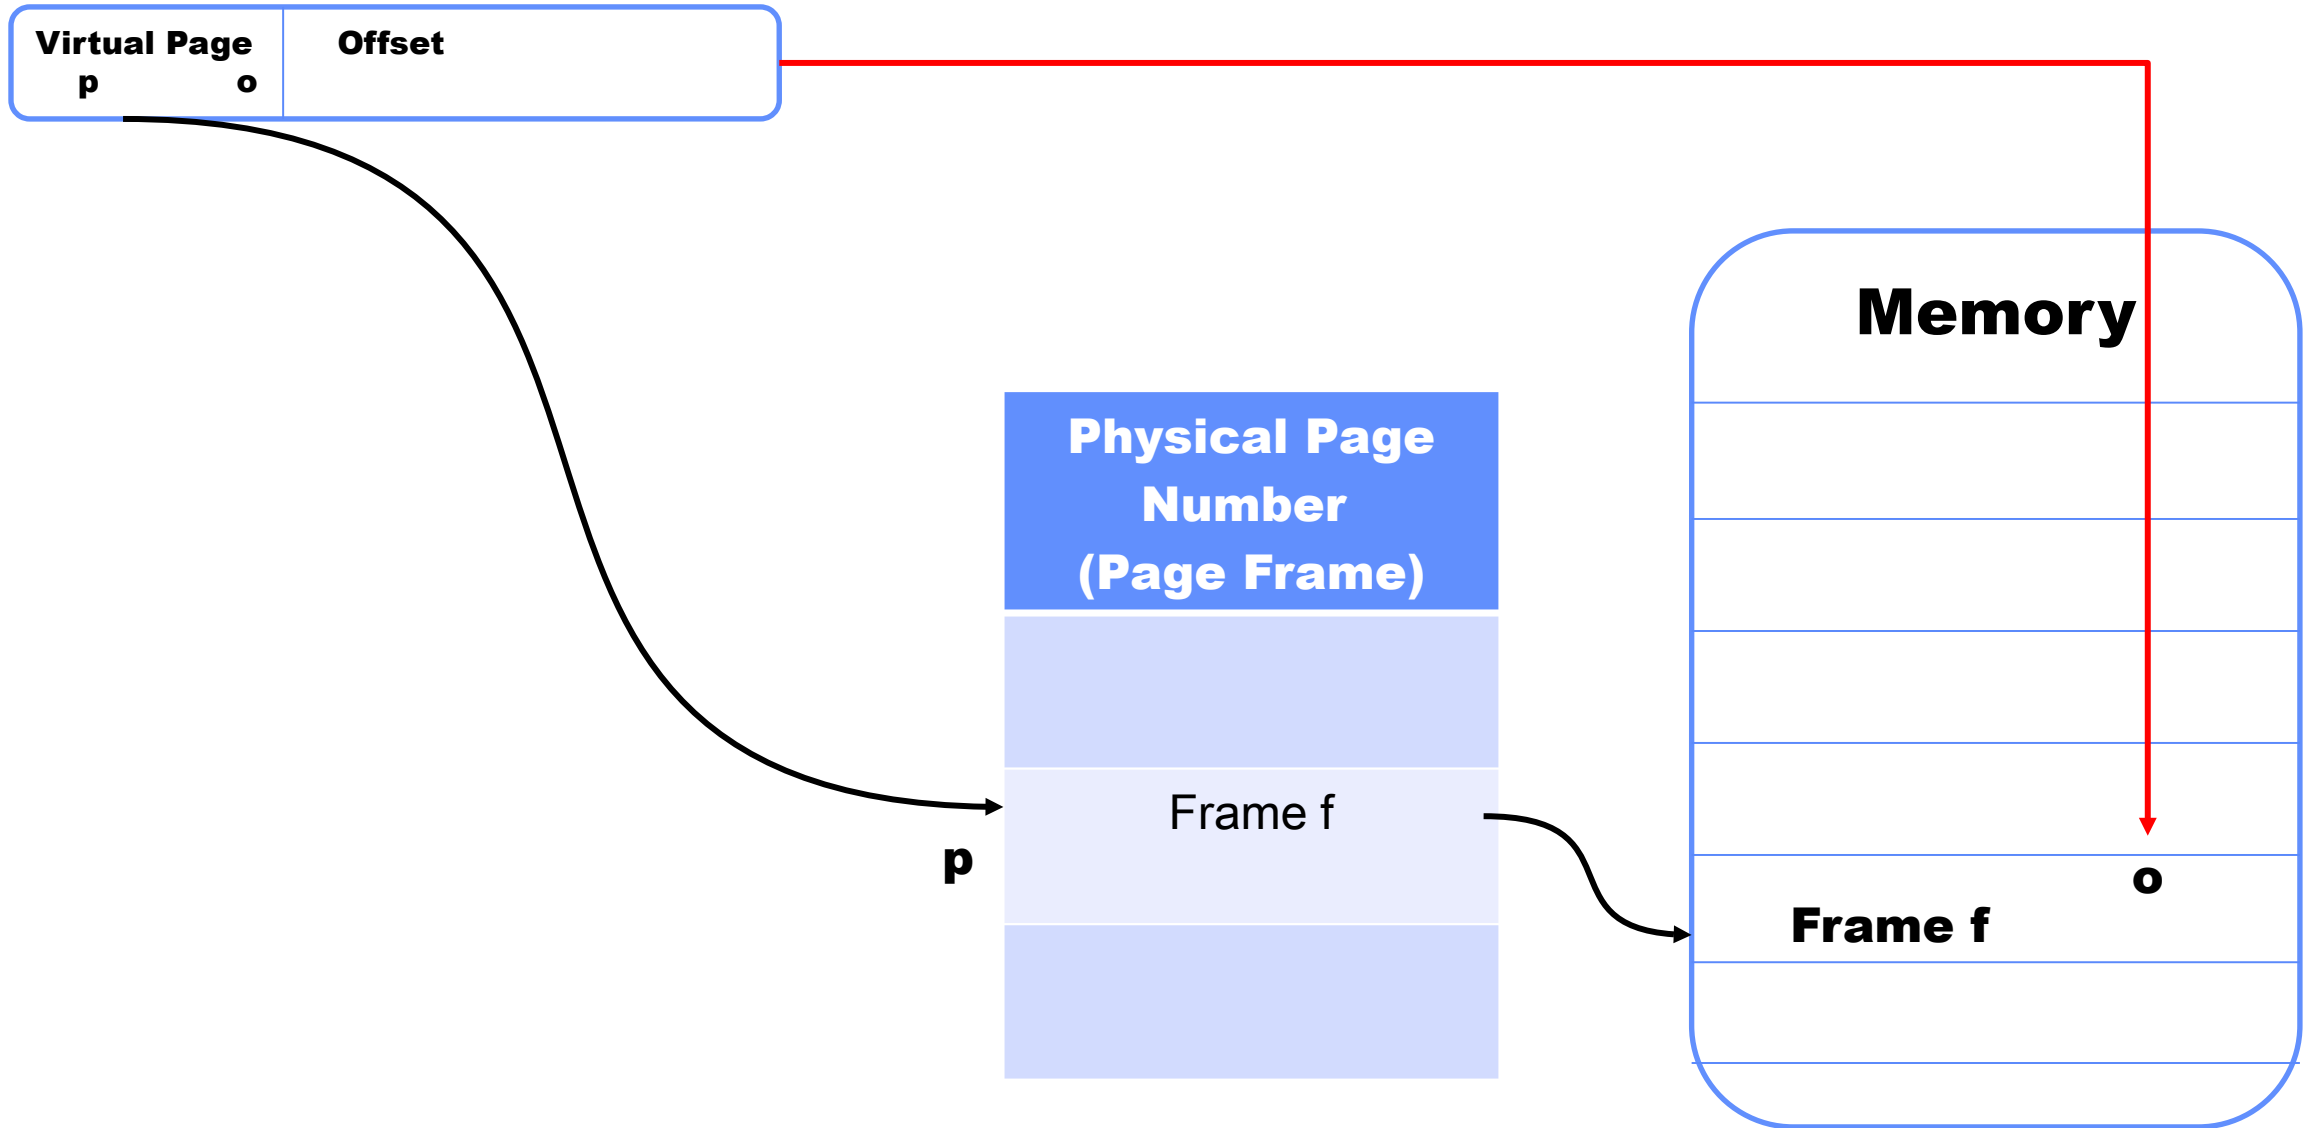
Map virtual page number into **a frame number**  
(also called physical page number) using a **page table**

Extract **offset** (last o bits)

Convert to **physical memory** location: access byte at offset in frame



# A (Simplified) Page Table



# Example: A Mini Page Table

---

Assume we have a 64 bytes ( $2^6$ ) of physical memory

Assume we want pages of 4 bytes ( $2^2$ )

How long should our addresses be?

6 bits

How many offset bits should we assign?


2 bits

How many virtual pages can we have?

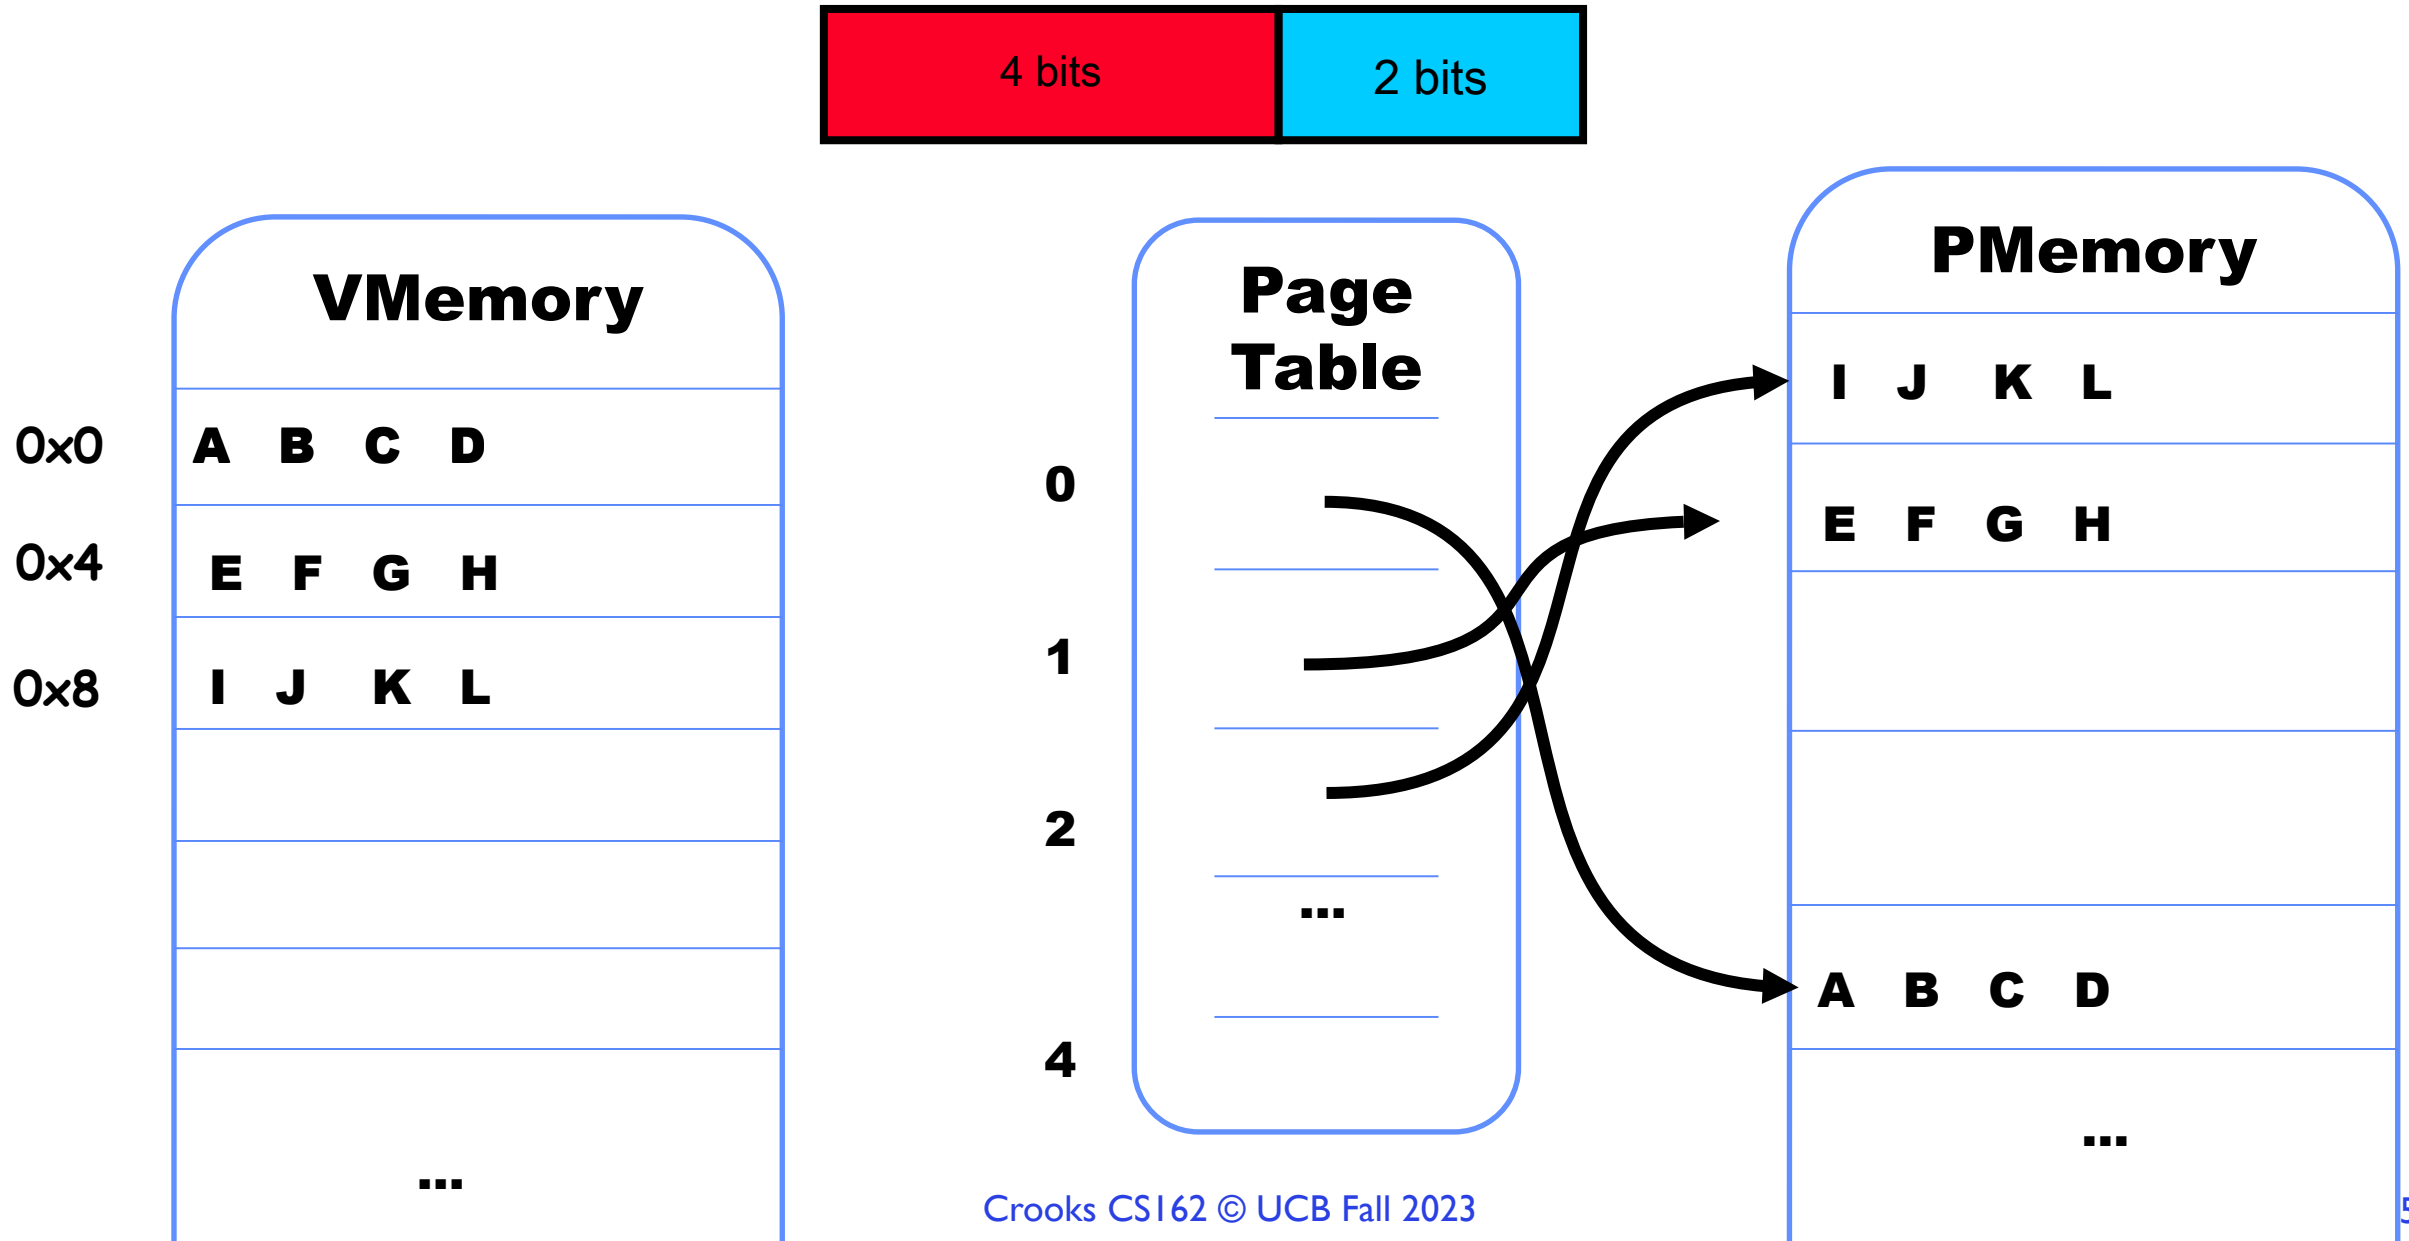
6 bit addresses, 2 bit for offsets, 4 bits for VPN.

$2^4 = 16$  pages

# Example: A Mini Page Table




# Example: A Mini Page Table



# Example: A Mini Page Table

0x9 =



# Step 1: Extract Virtual Page Number

**0x9 =**

0010

01

**VMemory**

**Page  
Table**

**PMemory**

0x0

**A B**

0x4

**E F**

0x8

**I J**

**Virtual Page Number:**

**0010 => 2.**

**Access Index 2 of Page Table**

**2**

...

**4**

**A B C D**

...

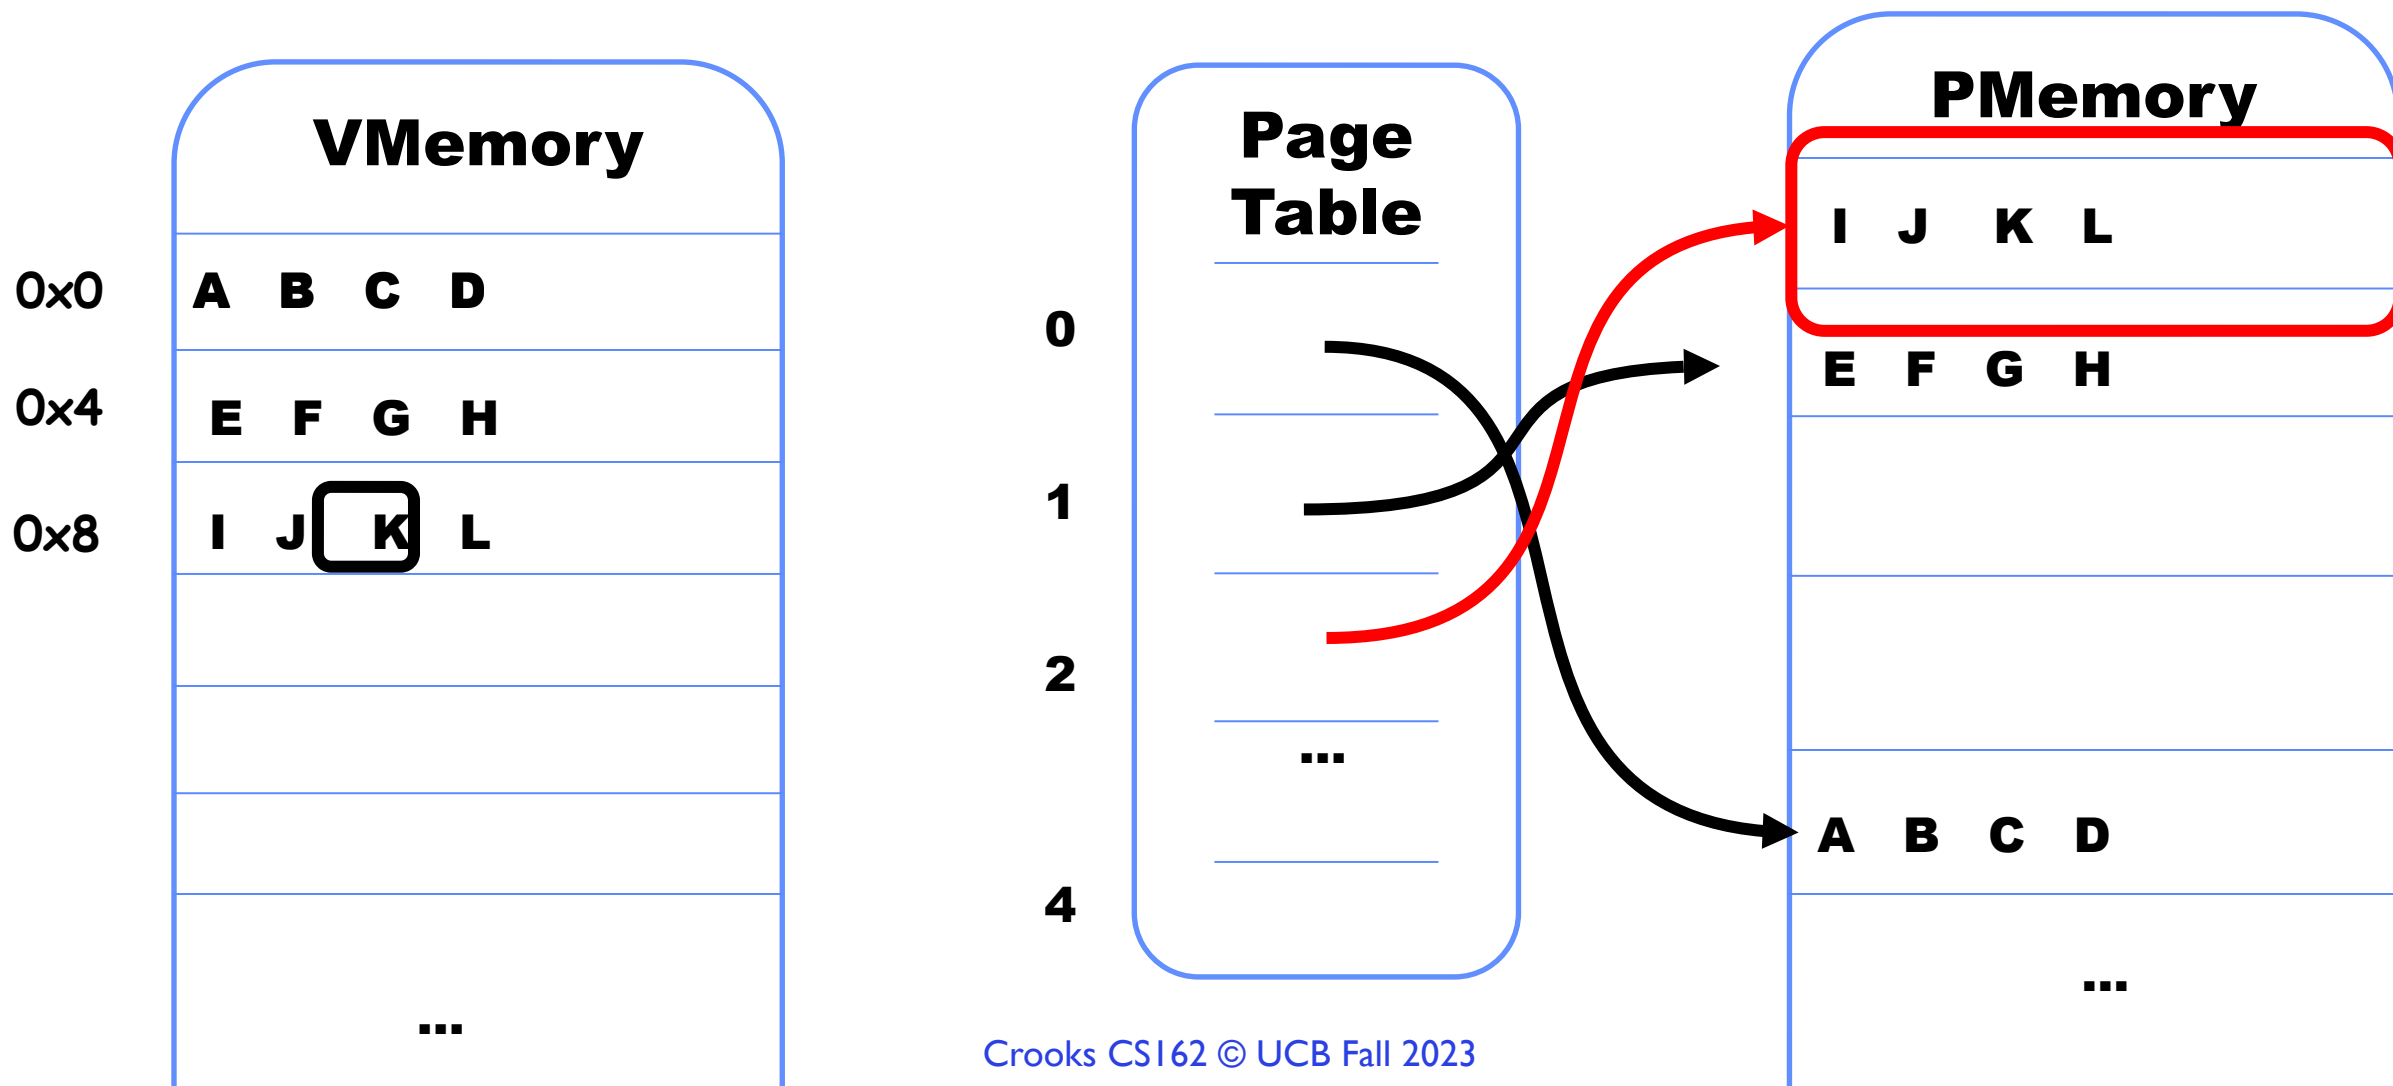
...

## Step 2: Identify Physical Page Number

0x9 =

0010

01



## Step 3: Extract Frame Offset

0x9 =

0010

01

### VMemory

0x0

A B C D

0x4

E F G H

0x8

I J **K** L

...

### Page Table

0

1

2

...

4

### PMemory

I J K L

E F G H

A B C D

...



## Step 3: Extract Frame Offset

**0x9 =**

0010

01

**VMemory**

**Page  
Table**

**PMemory**

0x0

**A B**

**Offset:**

0x4

**E F**

**01 => 1.**

0x8

**I J**

**Access Byte 1 of Frame**

2

...

4

**A B C D**

...

...

## Step 3: Extract Frame Offset

0x9 =

0010

01

### VMemory

0x0

A B C D

0x4

E F G H

0x8

I J **K** L

...

### Page Table

0

1

2

...

4

### PMemory

I J **K** L

E F G H

A B C D

...

## Step 4: Convert to Physical Address

**0x9 =**

0010

01

### VMemory

0x0

**A B C D**

0x4

**E F G H**

0x8

**I J K**

### Page Table

0

1

### PMemory

**I J K L**

**E F G H**

**Physical Page Number \* Page Size  
+ Offset**


**=**

**0 \* 4 + 1 = 1**

...

...

# What is a page table entry? (32 bits)



- P: Present (same as “valid” bit in other architectures)
- W: Writeable
- U: User accessible
- PWT: Page write transparent: external cache write-through
- PCD: Page cache disabled (page cannot be cached)
- A: Accessed: page has been accessed recently
- D: Dirty: page has been modified recently
- PS: Page Size

**Size of page table entry:  
PFN (20 bits) + 12 bits for  
access control/caching**

**4 bytes**

# The Great Power of the PTE

---

## Demand Paging

Keep only active pages in memory  
Place others on disk and mark their PTEs invalid

## Copy-on-Write

UNIX fork gives *copy* of parent address space to child. Use combination of page sharing + marking pages as read-only

## Zero Fill On Demand

New data pages must carry no information  
Mark PTEs as invalid; page fault on use gets zeroed page

## Data Breakpoints


For debugger, mark instruction page as read-only. Will trigger page-fault when try to execute

# Paging & Sharing

---

Processes share a page by each mapping a page of their own virtual address space to the same frame

Use protection bits for fine-sharing



# Where is page sharing used ?

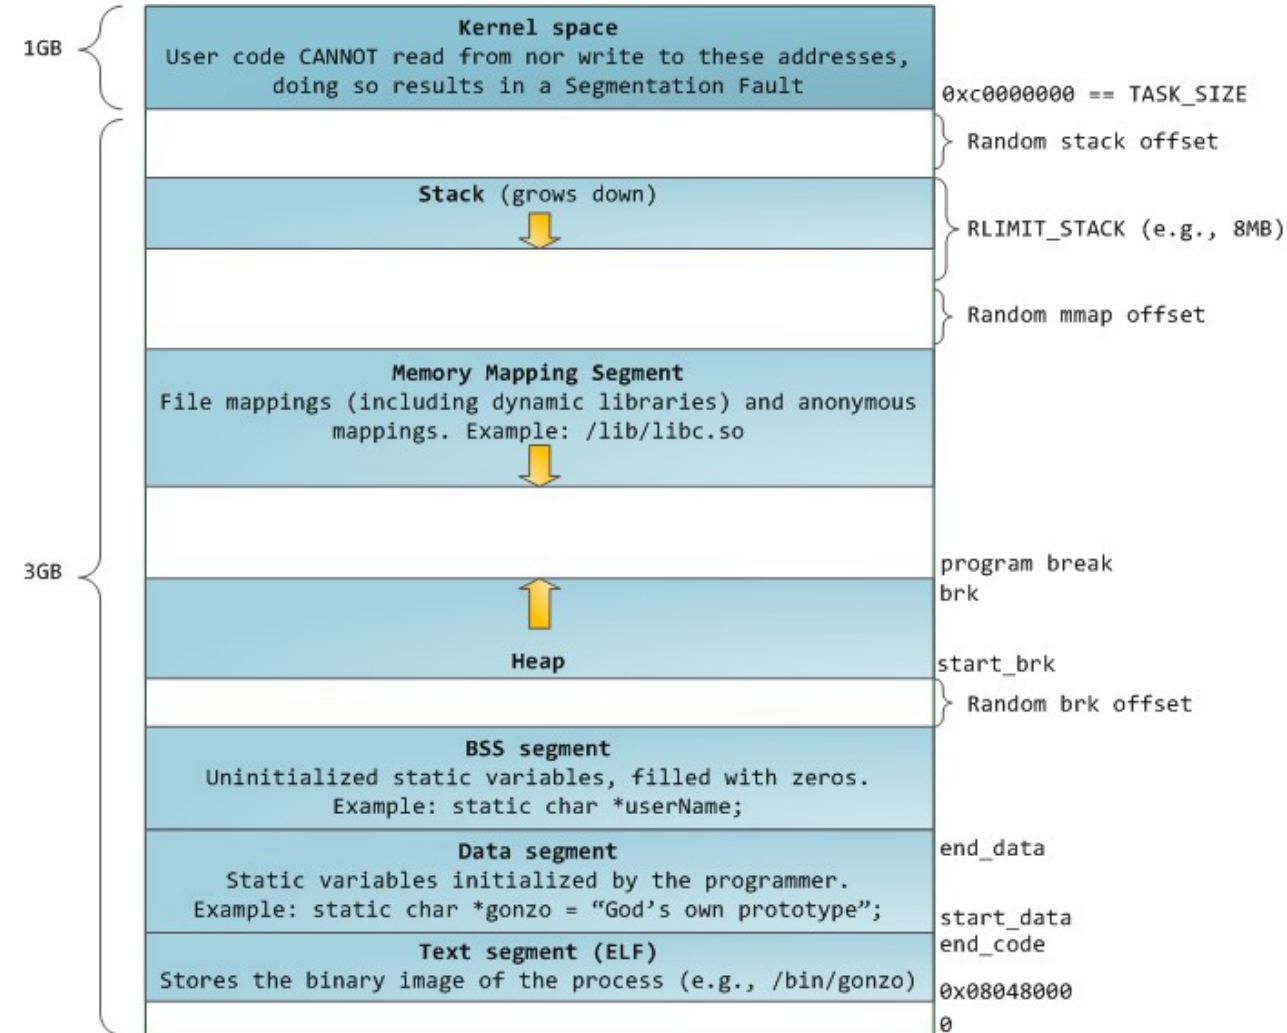
---

Kernel region of every process has the same page table entries

Different processes running same binary!  
Do not need to duplicate code segments

Shared-memory segments between different processes

# Memory Layout for Linux 32-bit





# An aside: Meltdown

---

From the paper:

Meltdown is a novel attack that allows overcoming memory isolation completely by providing a simple way for **any user process to read the entire kernel memory of the machine it executes on**, including all physical memory mapped in the kernel region. Meltdown does not exploit any software vulnerability, i.e., it works on all major operating systems.

```
1. raise_exception();  
2. // the line below is never reached  
3. access(probe_array[data * 4096]);]
```



# Are we done?

---

How big can a page table get on x86 (32 bits)?

4KB page  $\Rightarrow 2^{12}$   
 $2^{32}/2^{12} \Rightarrow 2^{20}$  pages  
 $2^{20} * 4 \text{ bytes} = 4 \text{ MB (approx.)}$   
That's a lot per process!!

How big can a page table get on x86 (64 bits)?

4KB page  $\Rightarrow 2^{12}$   
 $2^{64}/2^{12} \Rightarrow 2^{52}$  pages  
 $2^{52} * 8 \text{ bytes} = 36 \text{ petabytes (approx.)}$   
That's a lot per process!!

# Limitations of paging

---

## Space overhead

With a 64-bit address space, size of page table can be huge

## Time overhead

Accessing data now requires two memory accesses must also access page table, to find mapped frame

## Internal Fragmentation

4KB pages

# The Secret to the Whole of CS

---

Batching

Caching

Indirection

Specialised Hardware



# Sparsity

---

Address space is **sparse**, i.e. has holes that are not mapped to physical memory

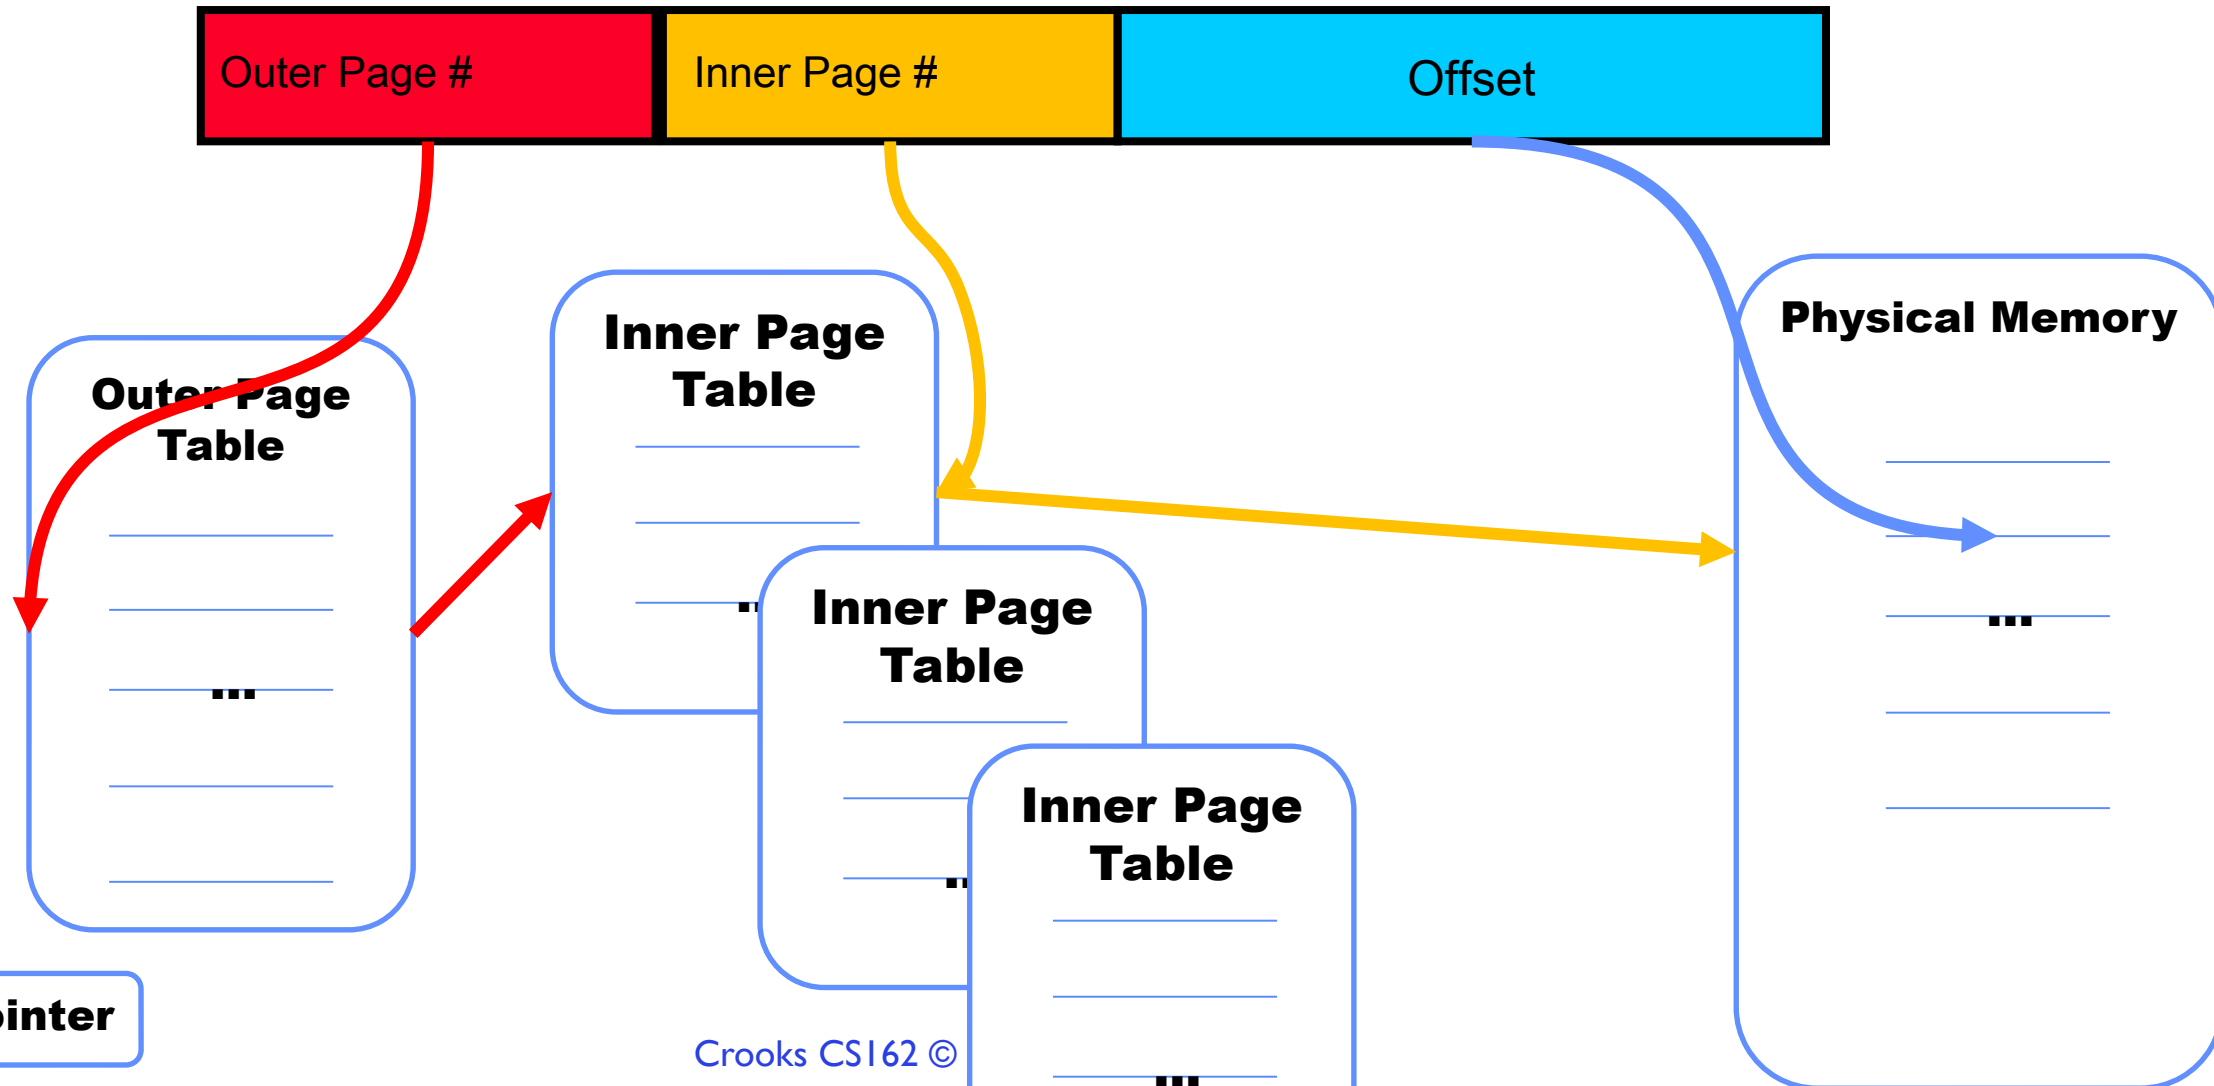
Most this space is taken up by page tables mapped to nothing

Process has access to full  $2^{64}$  bytes  
(virtually)

Physically, that would be 17,179,869,184  
gigabytes


# Paging the page table: 2-level paging

## Tree of Page Tables



## V2: What is a page table entry? (32 bits)

---



- P: Present (same as “valid” bit in other architectures)
- W: Writeable
- U: User accessible
- PWT: Page write transparent: external cache write-through
- PCD: Page cache disabled (page cannot be cached)
- A: Accessed: page has been accessed recently
- D: Dirty: page has been modified recently
- PS: Page Size

# Paging the page table: 2-level paging

---

## Tree of Page Tables



Number of  
top-level  
pages

Ensure that  
fits on a single  
page

Defines size of a page



# Paging the page table: 2-level paging

---

## Tree of Page Tables



10 bits

10 bits

4 KB  
12 bits

Want to make sure that  
inner page table fits in a  
page!

$$2^{12}/2^2 = 2^{10}$$

# Example: x86 classic 32-bit address translation

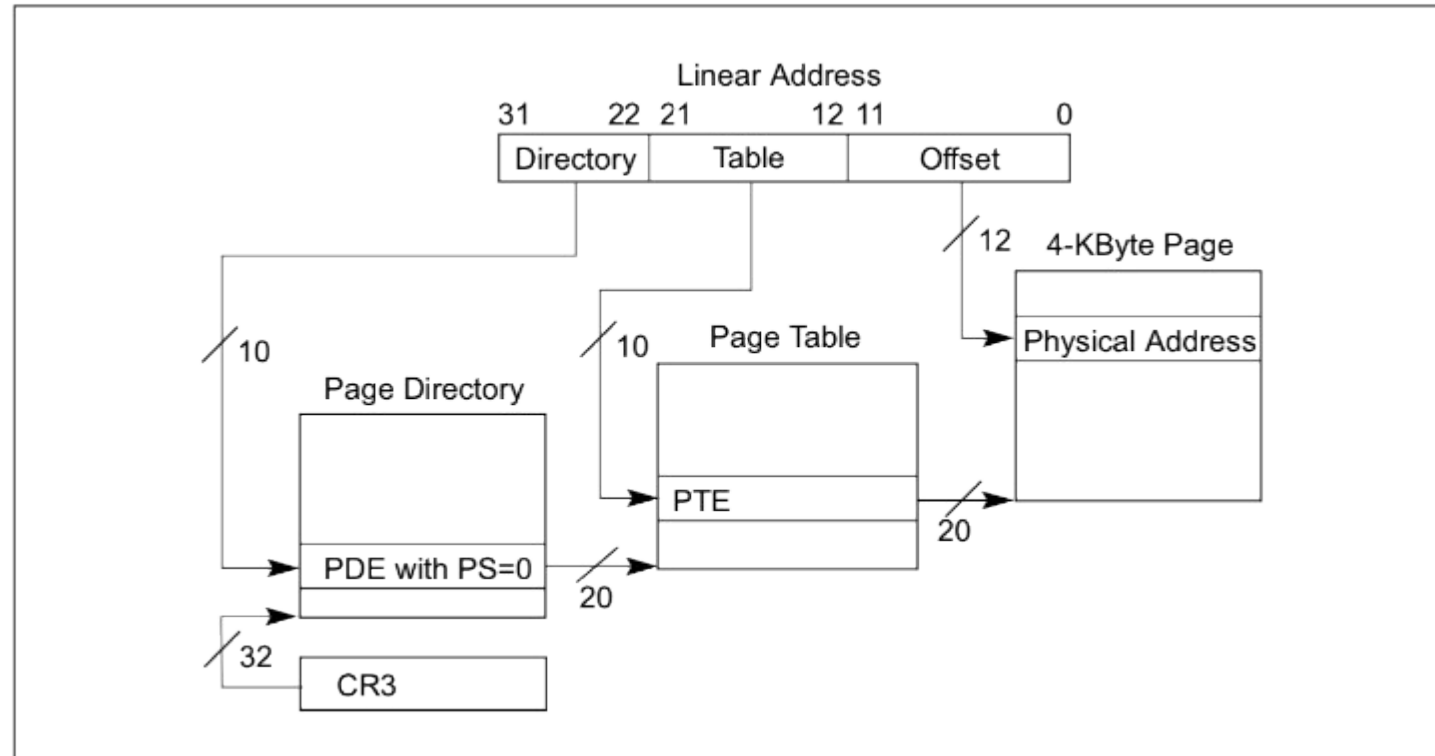
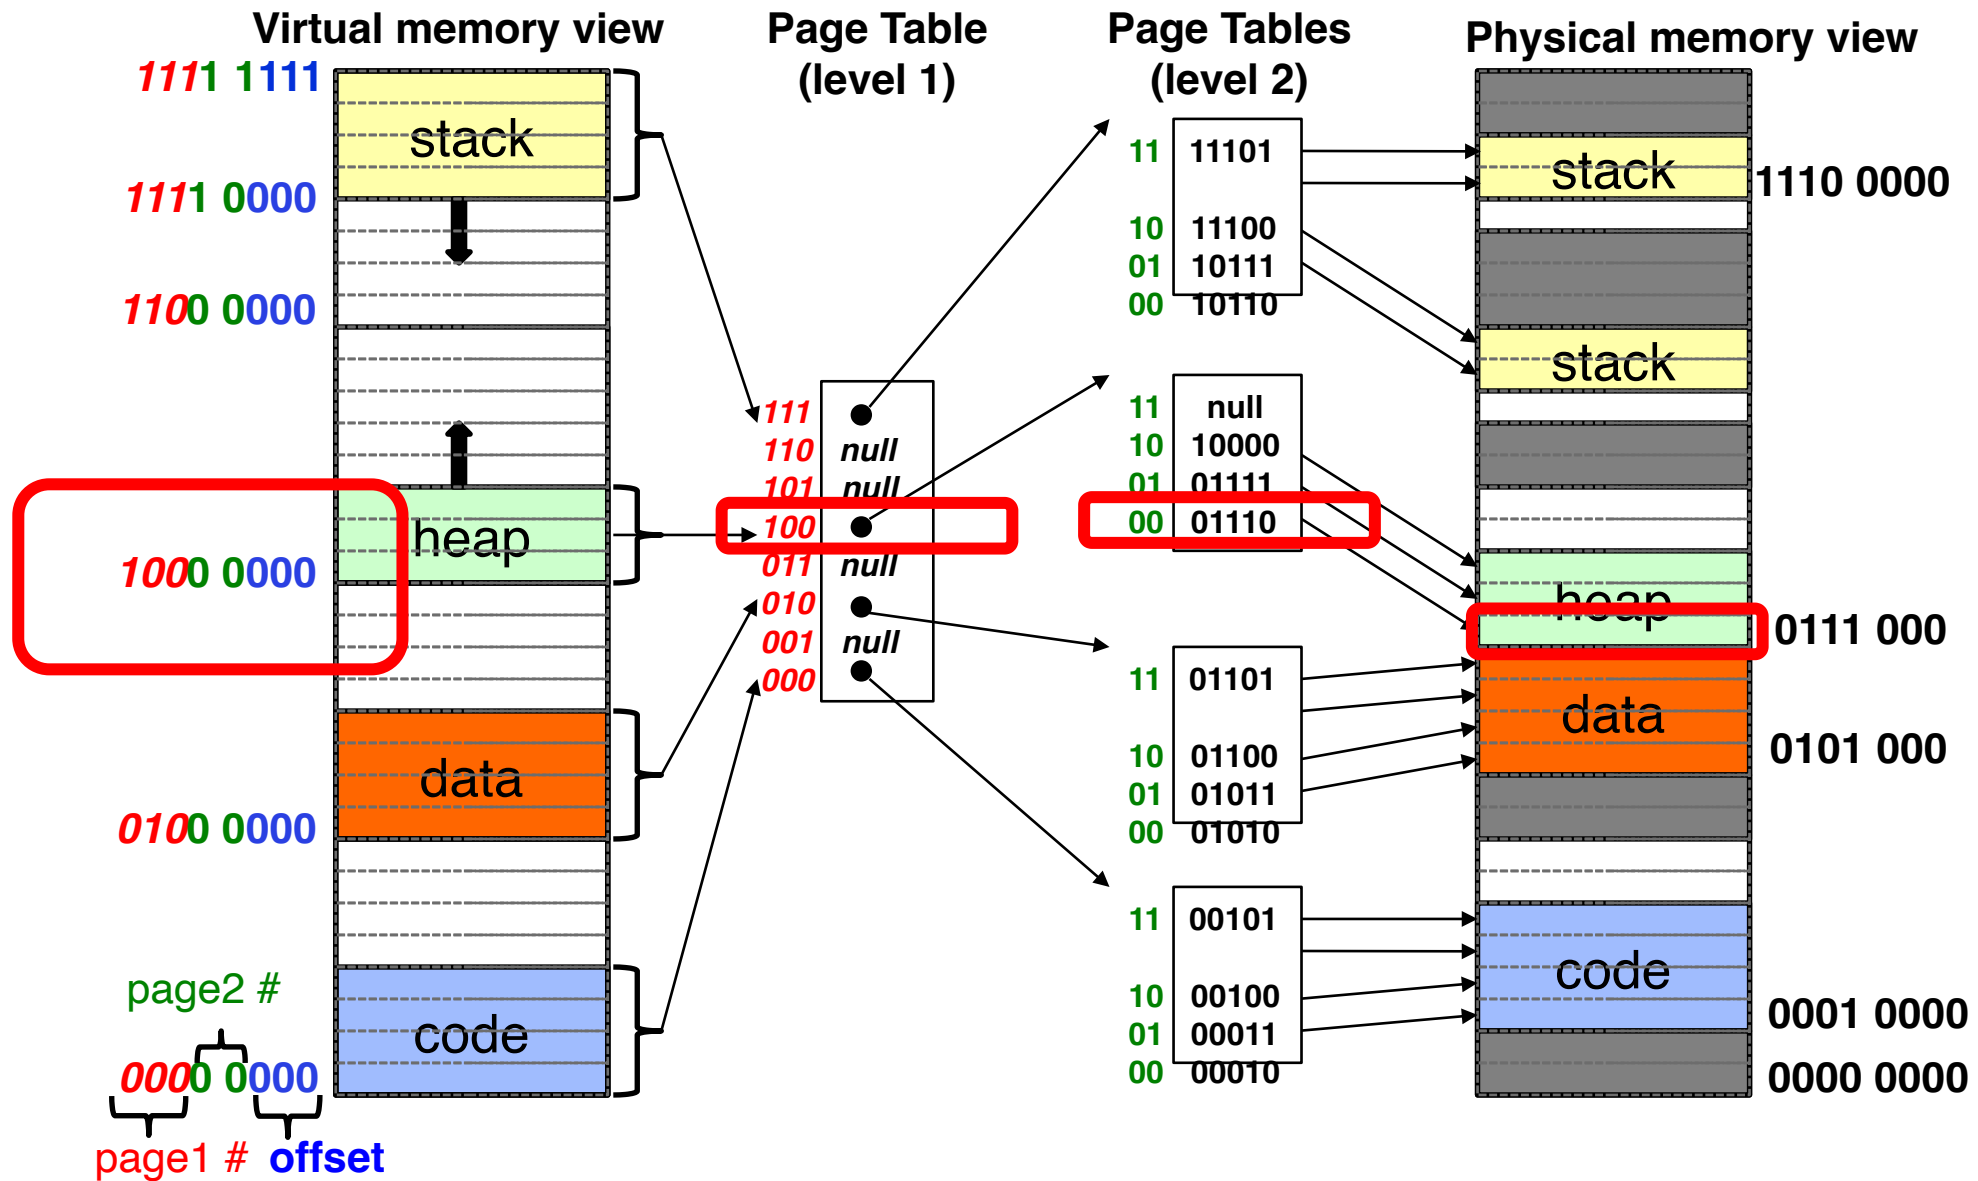


Figure 4-2. Linear-Address Translation to a 4-KByte Page using 32-Bit Paging

Top-level page-table: [Page Directory](#)

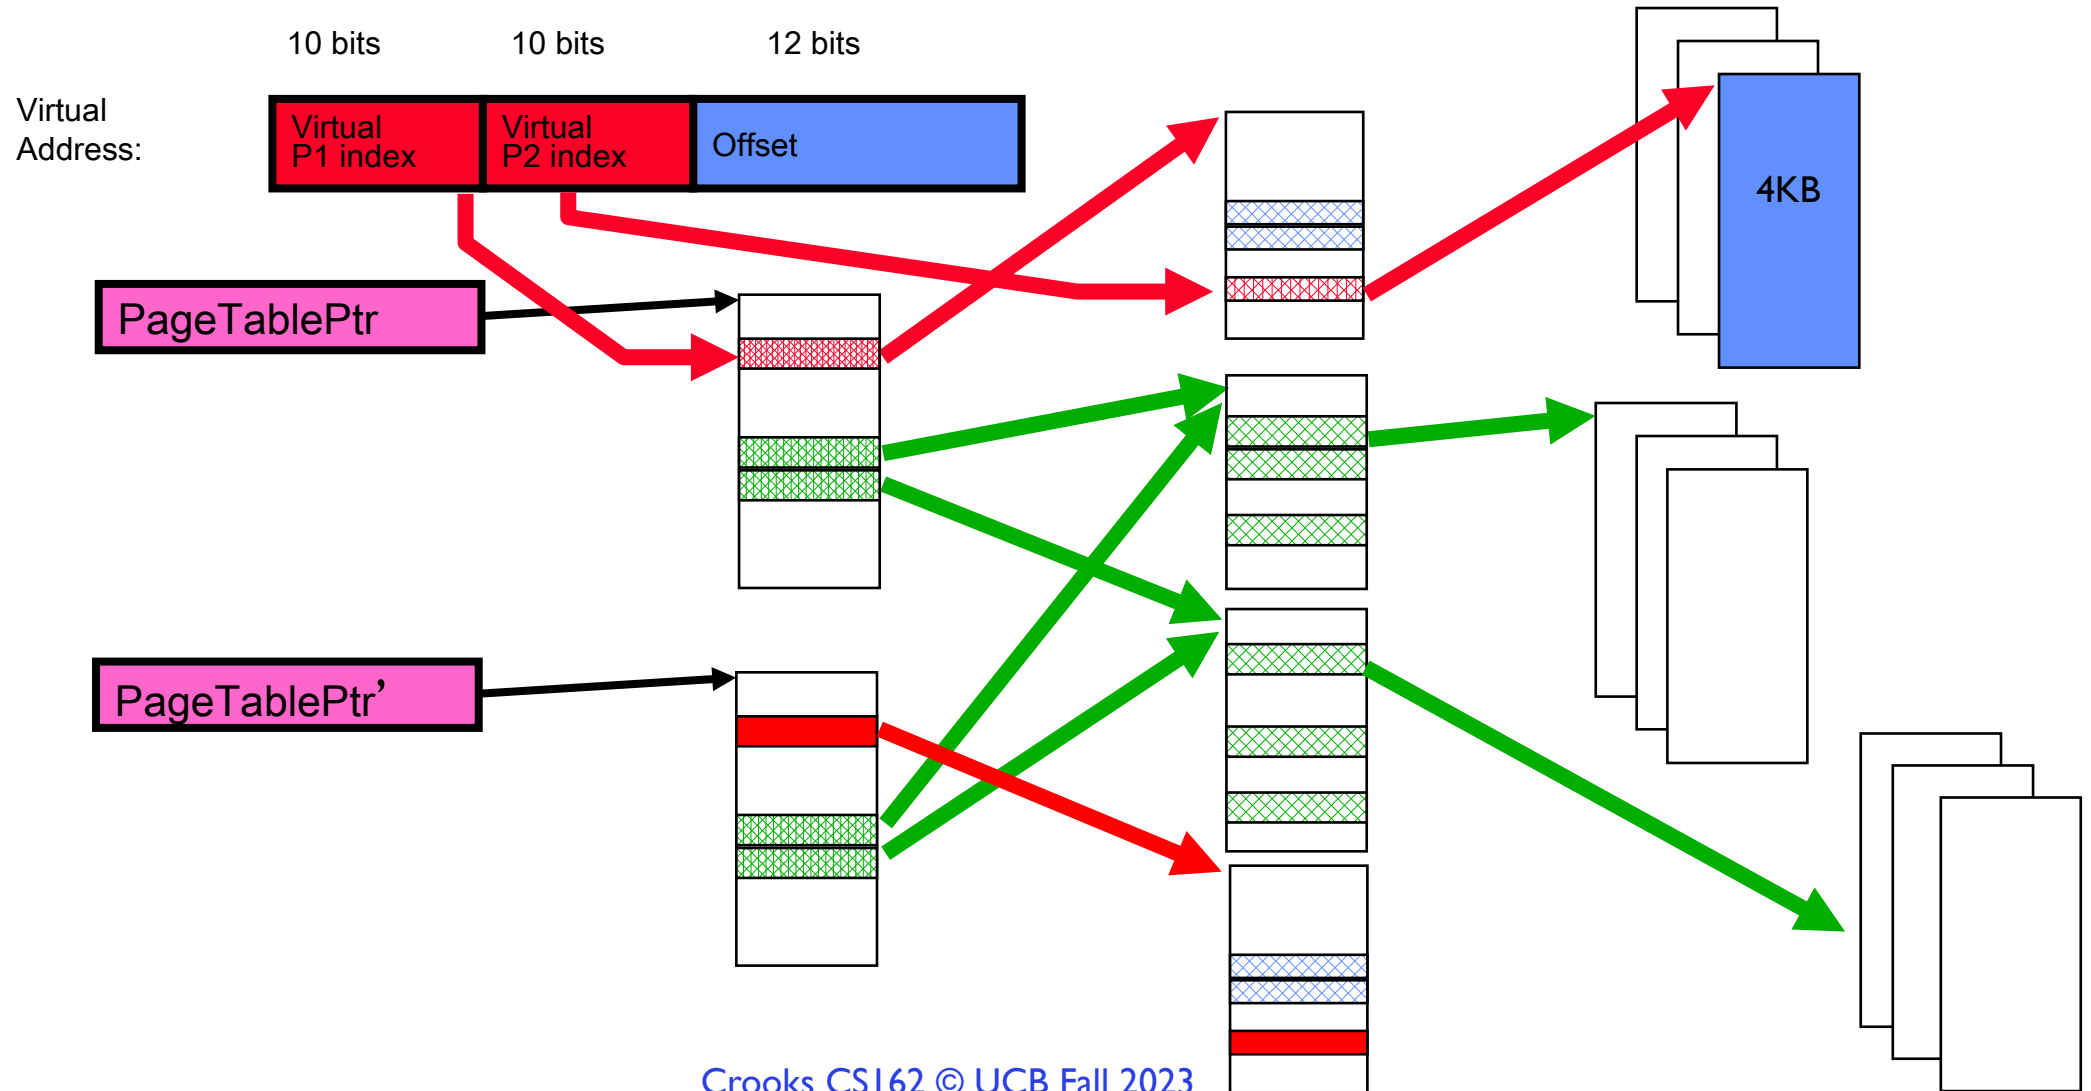
Inner page-table: [Page Directory Entries](#)

# Example Address Space View



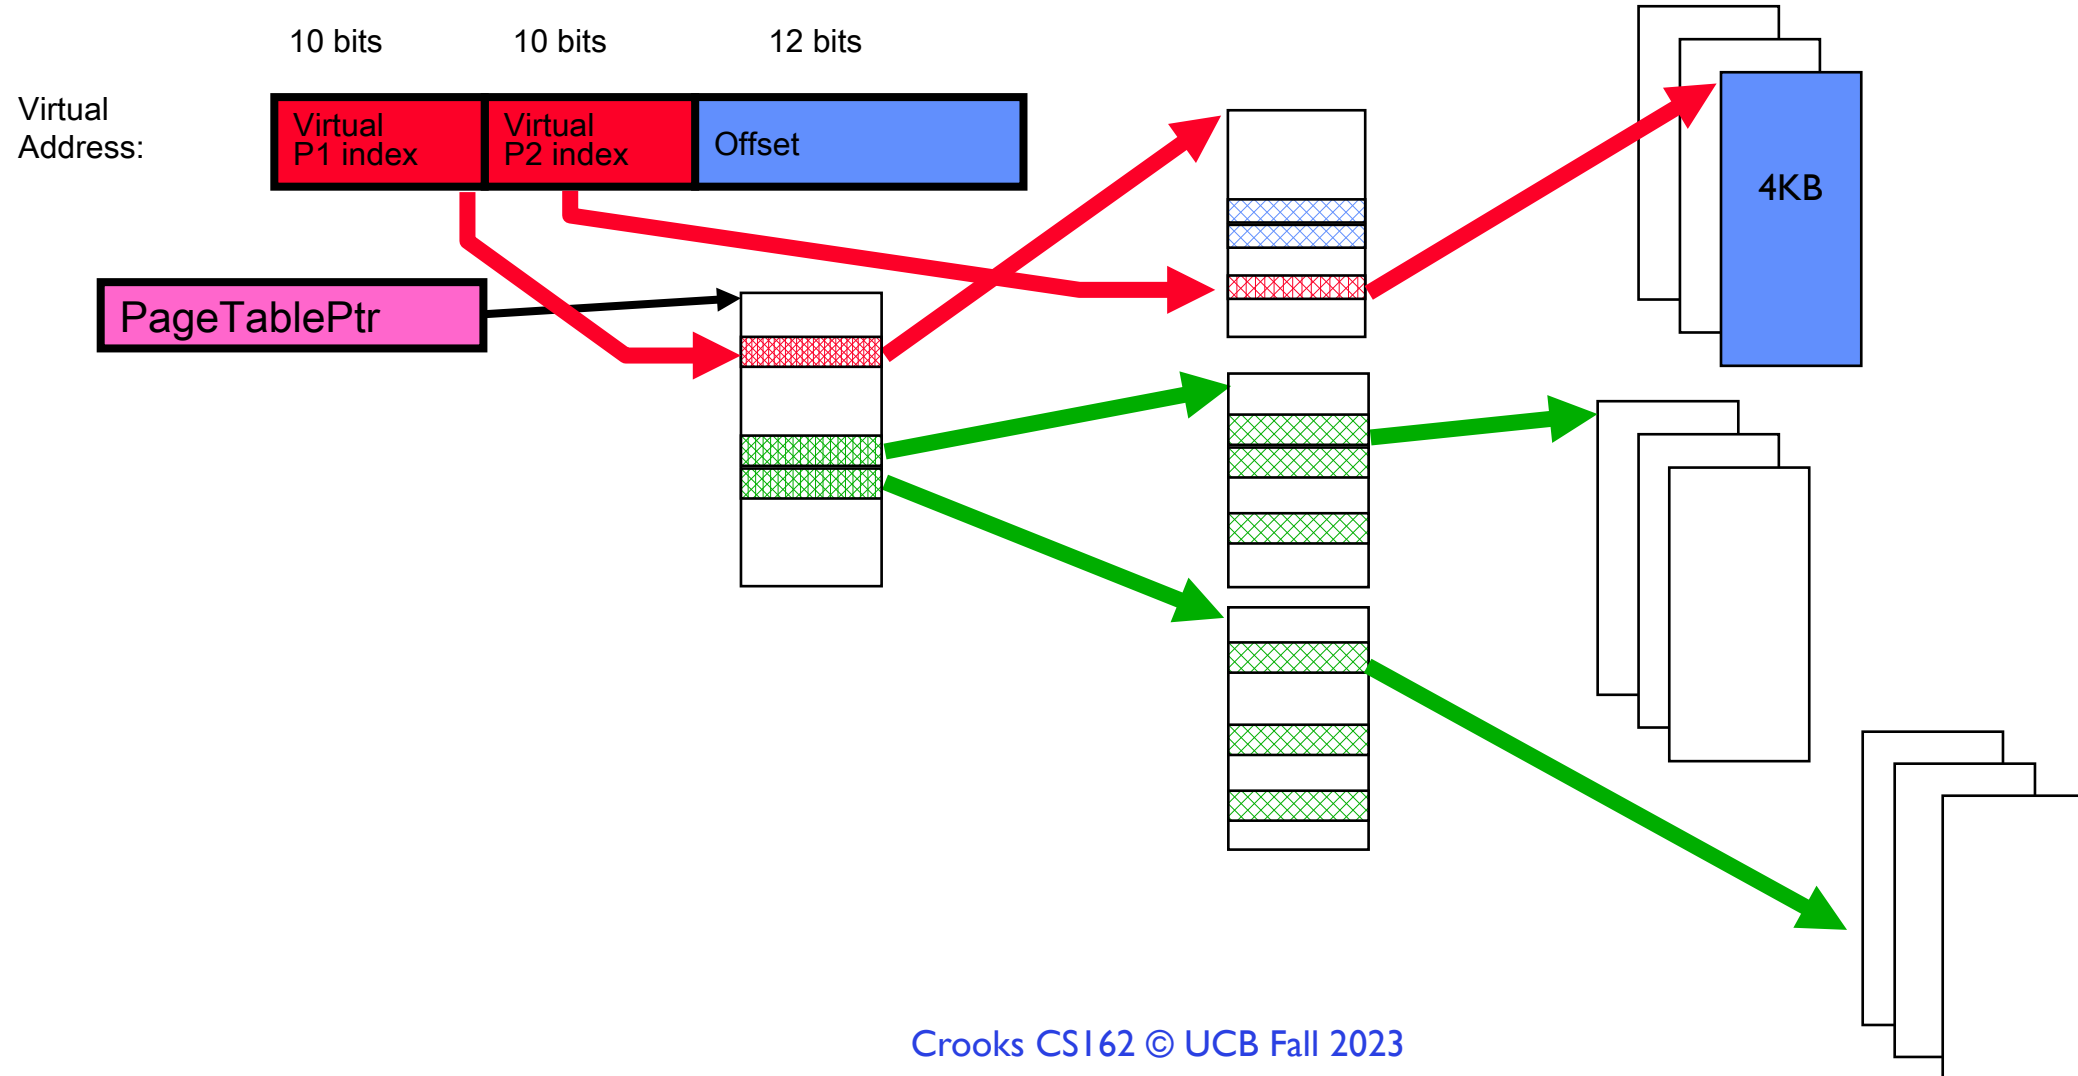
# Sharing with multilevel page tables

Entire regions of the address space can be efficiently shared



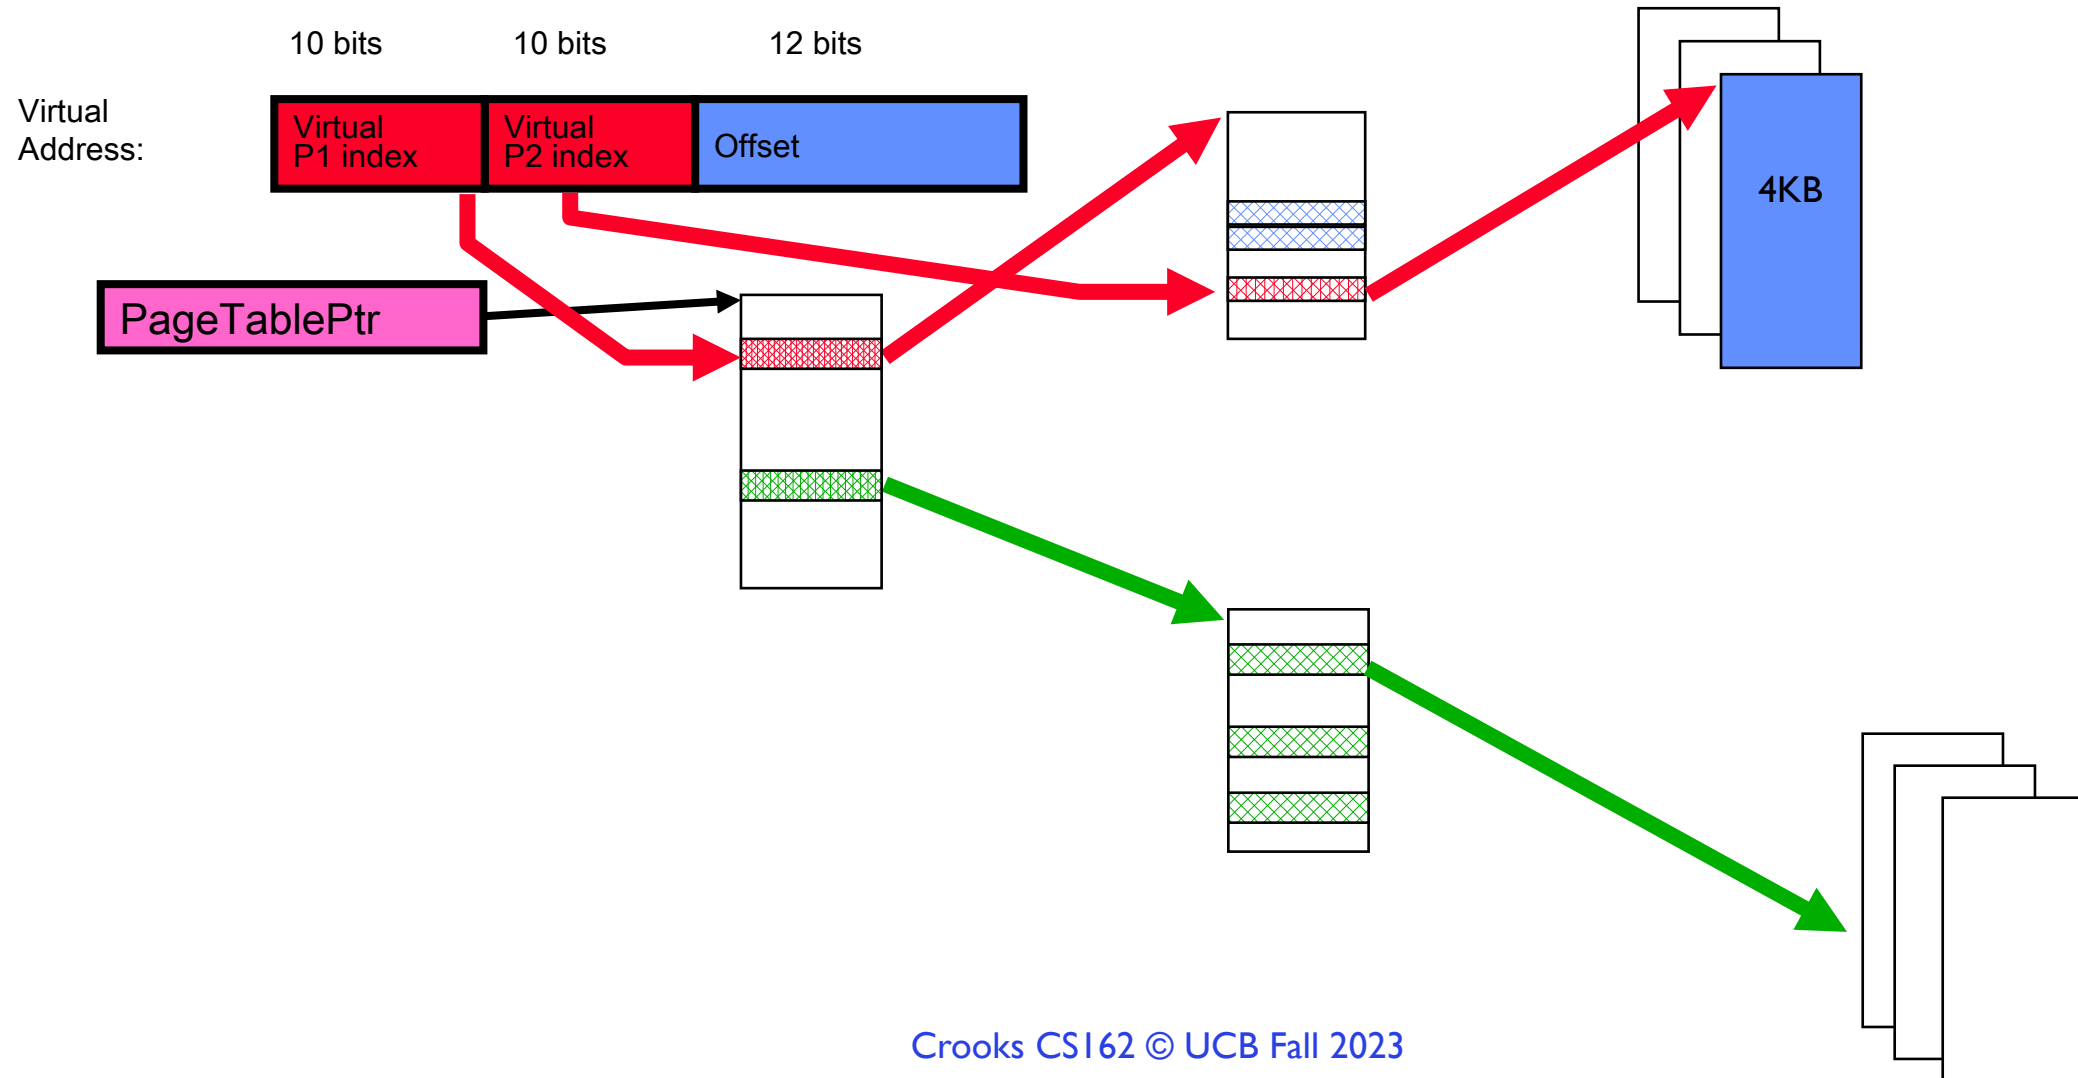
# Marking entire regions as invalid!

If region of address space unused, can mark entire inner region as invalid



# Marking entire regions as invalid!

If region of address space unused, can mark entire inner region as invalid



# Has this helped?

---

Assuming 10/10/12 split:

Size of Page Table

Outer:  $(2^{10} * 4 \text{ bytes}) +$

Inner:  $2^{10} * (2^{10} * 4 \text{ bytes})$

Overhead of indirection! BUT Marking inner pages as invalid helps when address spaces are sparse

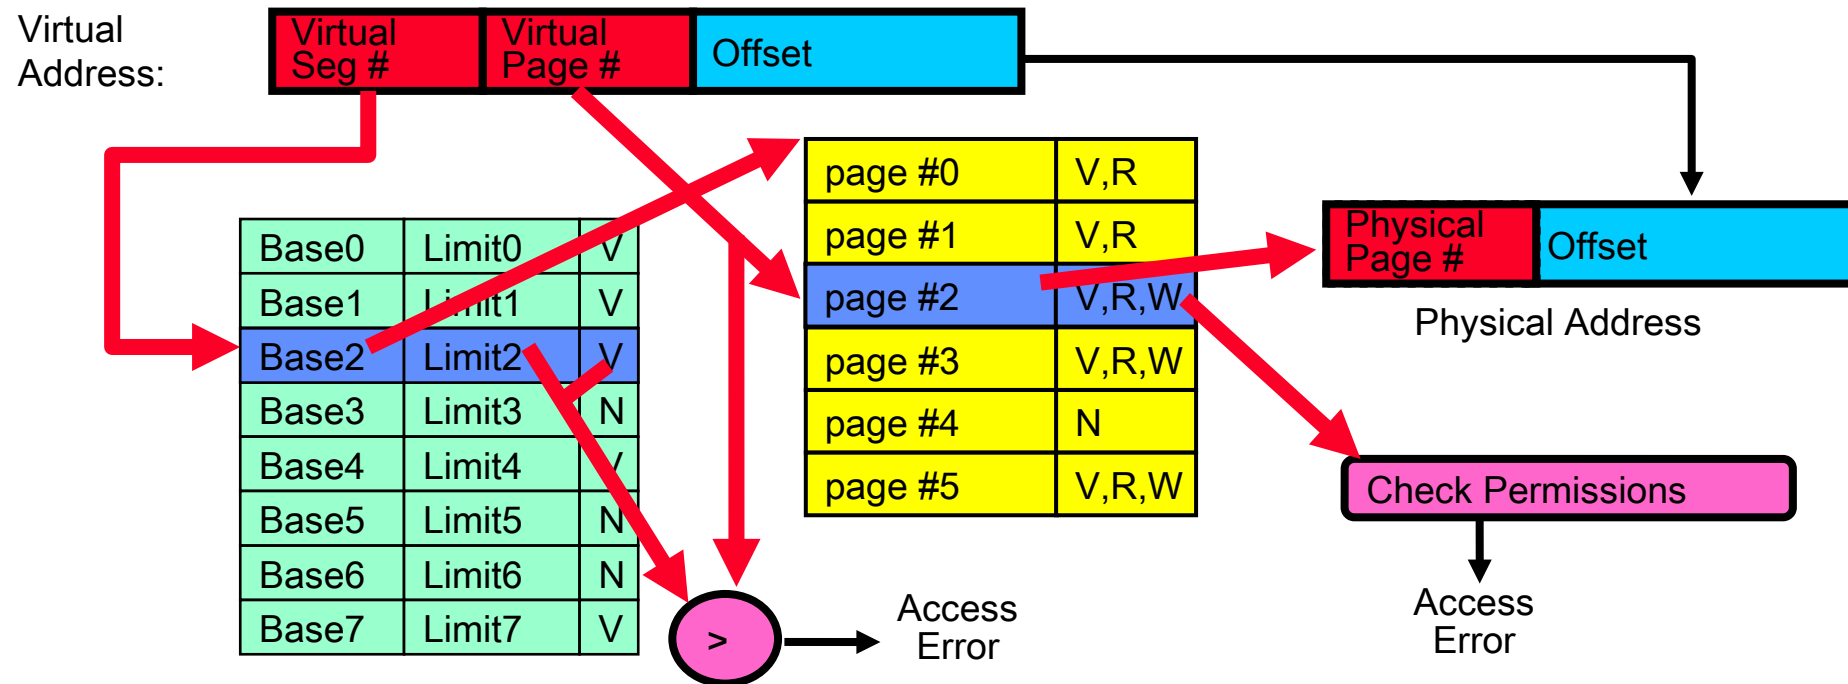
Downside: now have to do

two memory accesses for translation

# Paged Segmentation

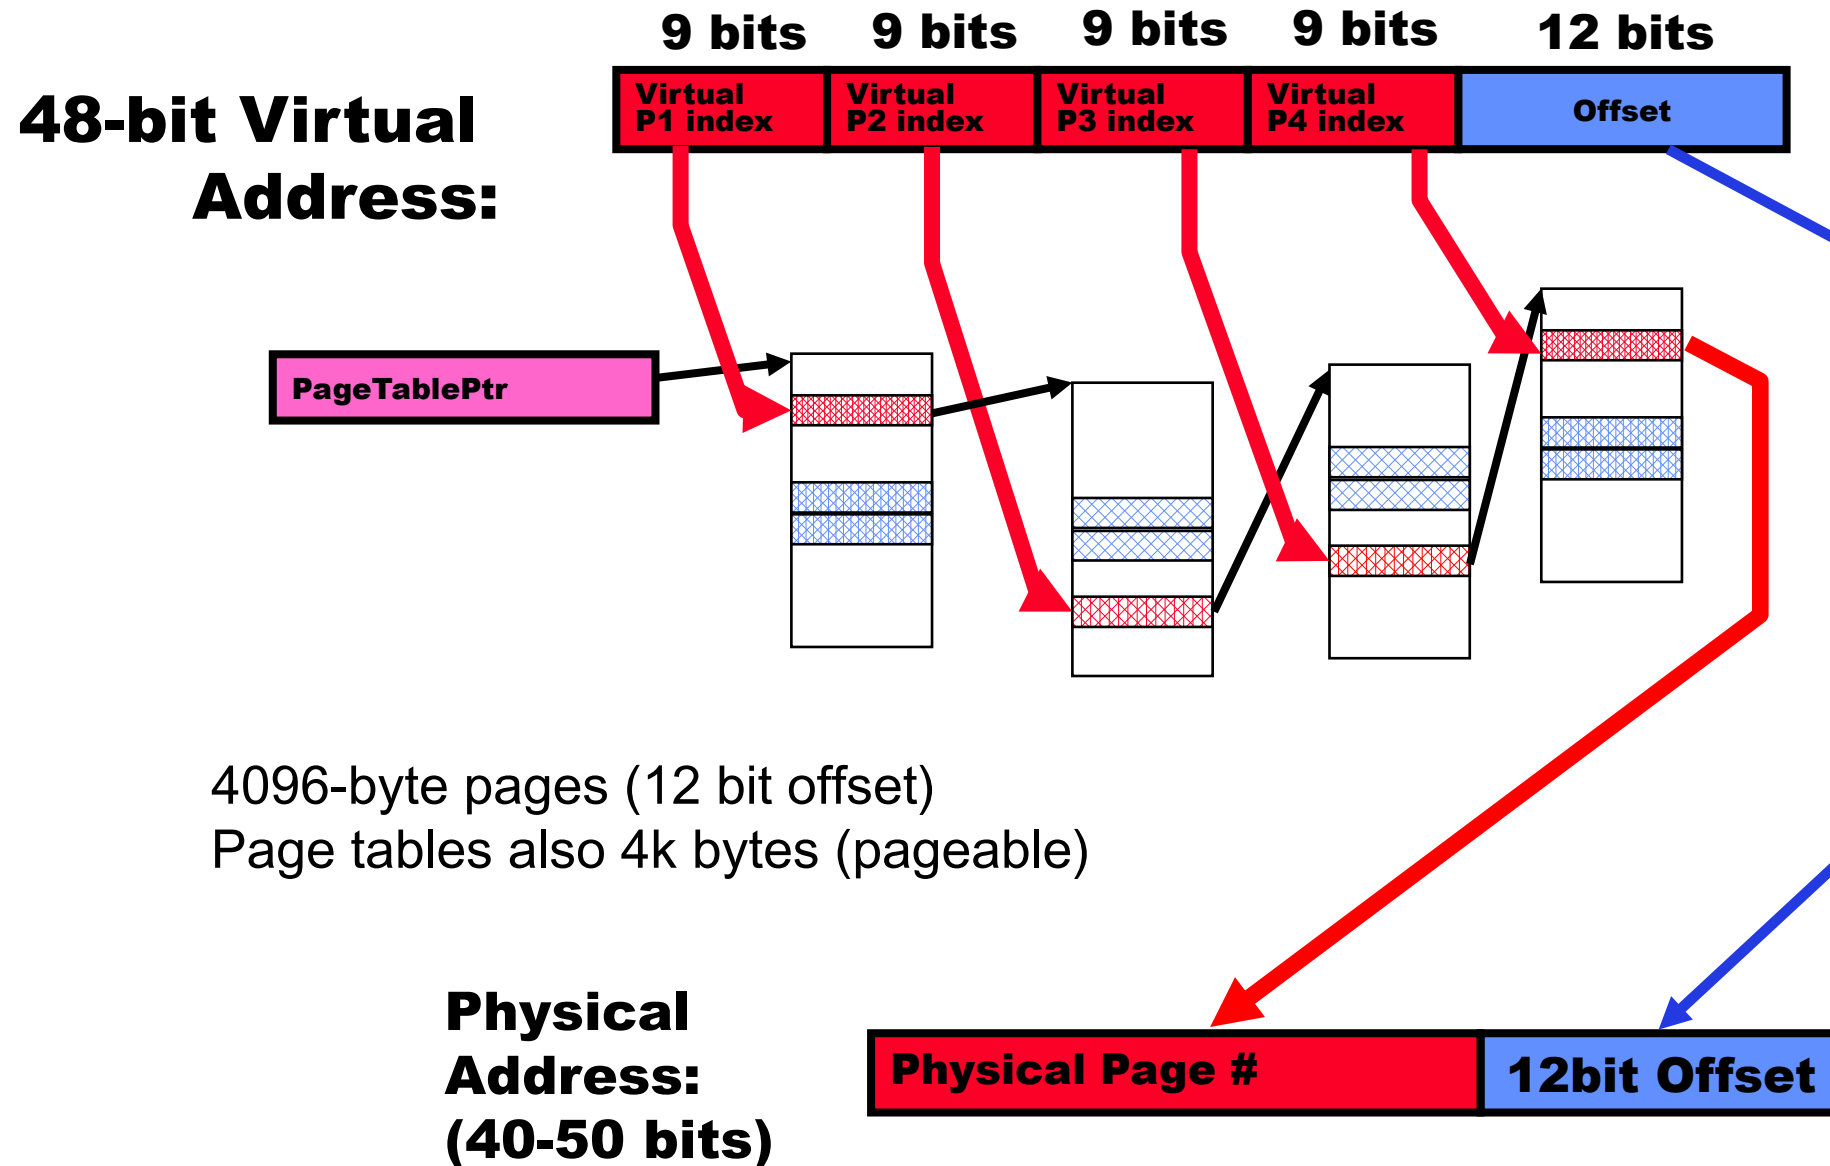
Use segments for top level. Paging within each segment.

Used in x86 (32 bit).  
Code Segment, Data Segment, etc.





# X86 64 bits has a four-level page table!



# Inverted Page Table

---

A single page table that has an entry for each physical page of the system

Each entry contains process ID + which virtual page maps to physical page

Physical memory much smaller than virtual memory

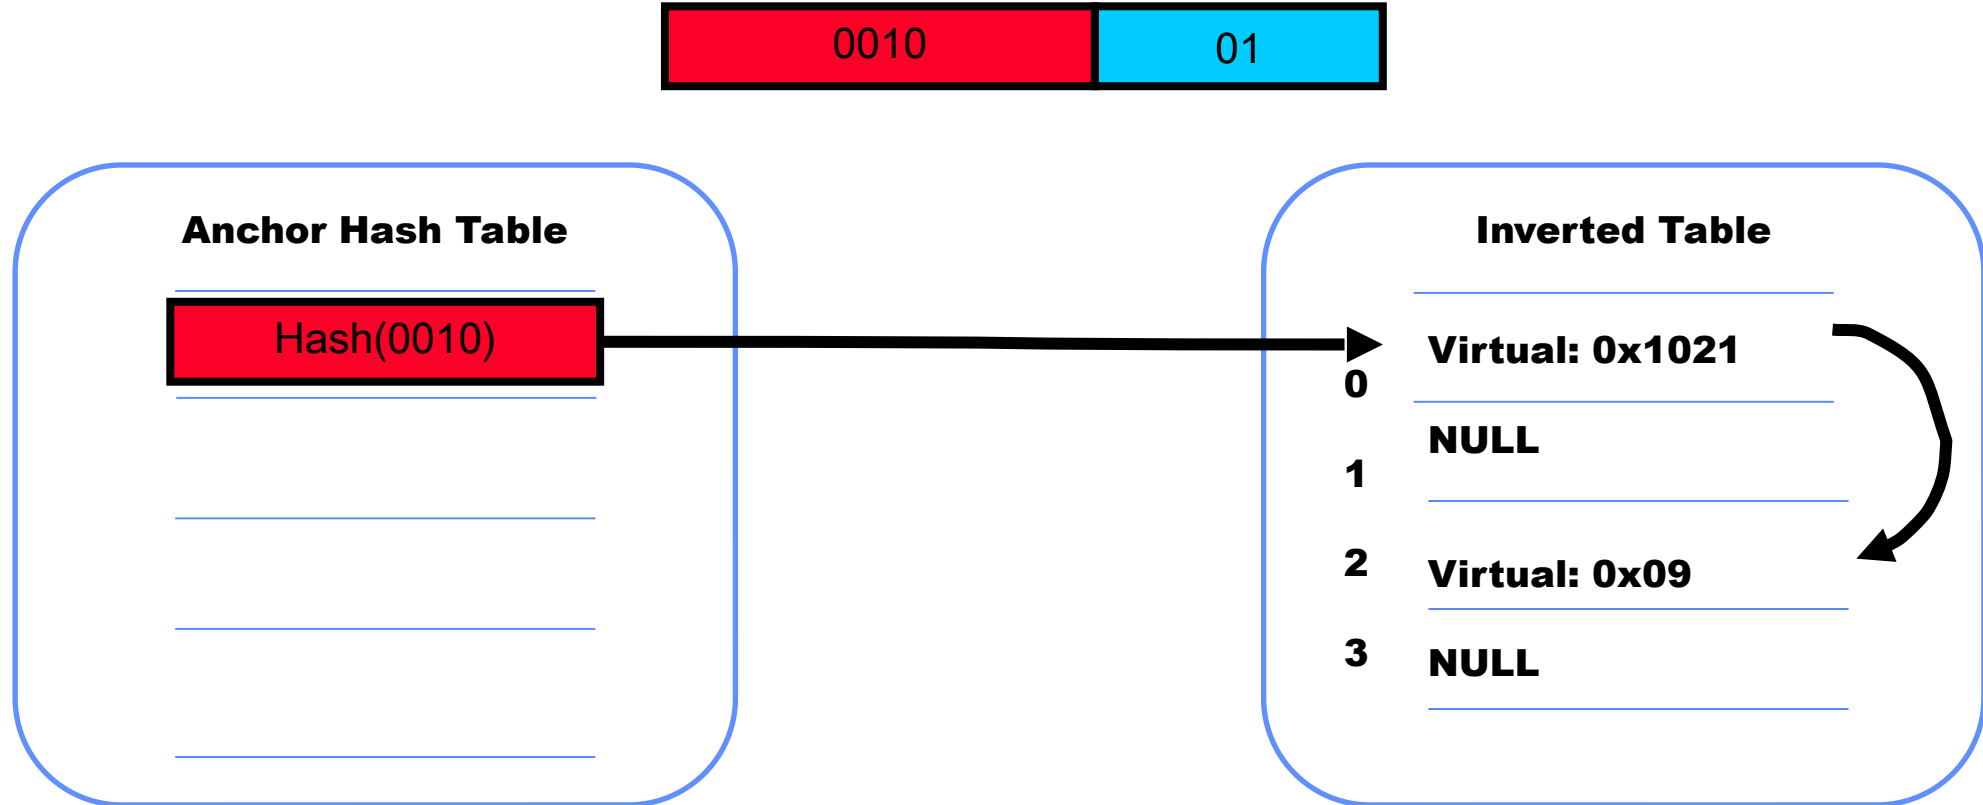
Size proportional to **size of physical memory**

Inverted Table	
<b>0</b>	<b>Virtual: 0x1021</b>
<b>1</b>	<b>NULL</b>
<b>2</b>	<b>Virtual: 0x0123</b>
<b>3</b>	<b>NULL</b>

# Inverted Page Table

Don't we have it backwards?

Add a hash table. Virtual memory can only map to specific physical frames




# Address Translation Comparison

---

	Advantages	Disadvantages
Simple Segmentation	Fast context switching (segment map maintained by CPU)	External fragmentation
Paging (Single-Level)	No external fragmentation Fast and easy allocation	Large table size (~ virtual memory) Internal fragmentation
Paged Segmentation	Table size ~ # of pages in virtual memory Fast and easy allocation	Multiple memory references per page access
Multi-Level Paging		
Inverted Page Table	Table size ~ # of pages in physical memory	Hash function more complex No cache locality of page table

# How is the Translation Accomplished?

---



MMU must translate virtual address to physical address on every instruction fetch, load or store

What does the MMU need to do to translate an address?  
Read, check, and update PTE  
(set accessed bit/dirty bit on write)

# How can we speedup translation?

---

MMU must make at least 2 memory reads to walk page table. Slow!

Use specialized hardware to  
cache virtual-physical memory translations!

Introducing the Translation Lookaside Buffer (TLB)

# Recall: CS61c Caching Concept

---

**Cache:** a repository for copies that can be accessed more quickly than the original


Only good if:  
Frequent case frequent enough and  
Infrequent case not too expensive

Important measure  
Average Access time =  
(Hit Rate x Hit Time) + (Miss Rate x Miss Time)

# Recall: In Machine Structures (eg. 61C) ...

---

Caching is the key to memory system performance





## Recall: In Machine Structures (eg. 61C) ...

---

$$\begin{aligned} & \text{Average Memory Access Time (AMAT)} \\ &= (\text{Hit Rate} \times \text{HitTime}) + (\text{Miss Rate} \times \text{MissTime}) \\ & \text{Where HitRate} + \text{MissRate} = 1 \end{aligned}$$

$$\text{HitRate} = 90\% \Rightarrow \text{AMAT} = (0.9 \times 1) + (0.1 \times 101) = 11 \text{ ns}$$

$$\text{HitRate} = 99\% \Rightarrow \text{AMAT} = (0.99 \times 1) + (0.01 \times 101) = 2.01 \text{ ns}$$

$\text{MissTime}_{L_1}$  includes

$$\text{HitTime}_{L_1} + \text{MissPenalty}_{L_1} \equiv \text{HitTime}_{L_1} + \text{AMAT}_{L_2}$$

# Why Does Caching Help? Locality!

---

## Temporal Locality (Locality in Time):

Keep recently accessed data items closer to processor

## Spatial Locality (Locality in Space):

Move contiguous blocks to the upper levels

# Recall: Memory Hierarchy


---

Take advantage of the principle of locality to:

- 1) Present the illusion of having as much memory as in the cheapest technology
- 2) Provide average speed similar to that offered by the fastest technology

Recall: fast but small/expensive. Slow but large!

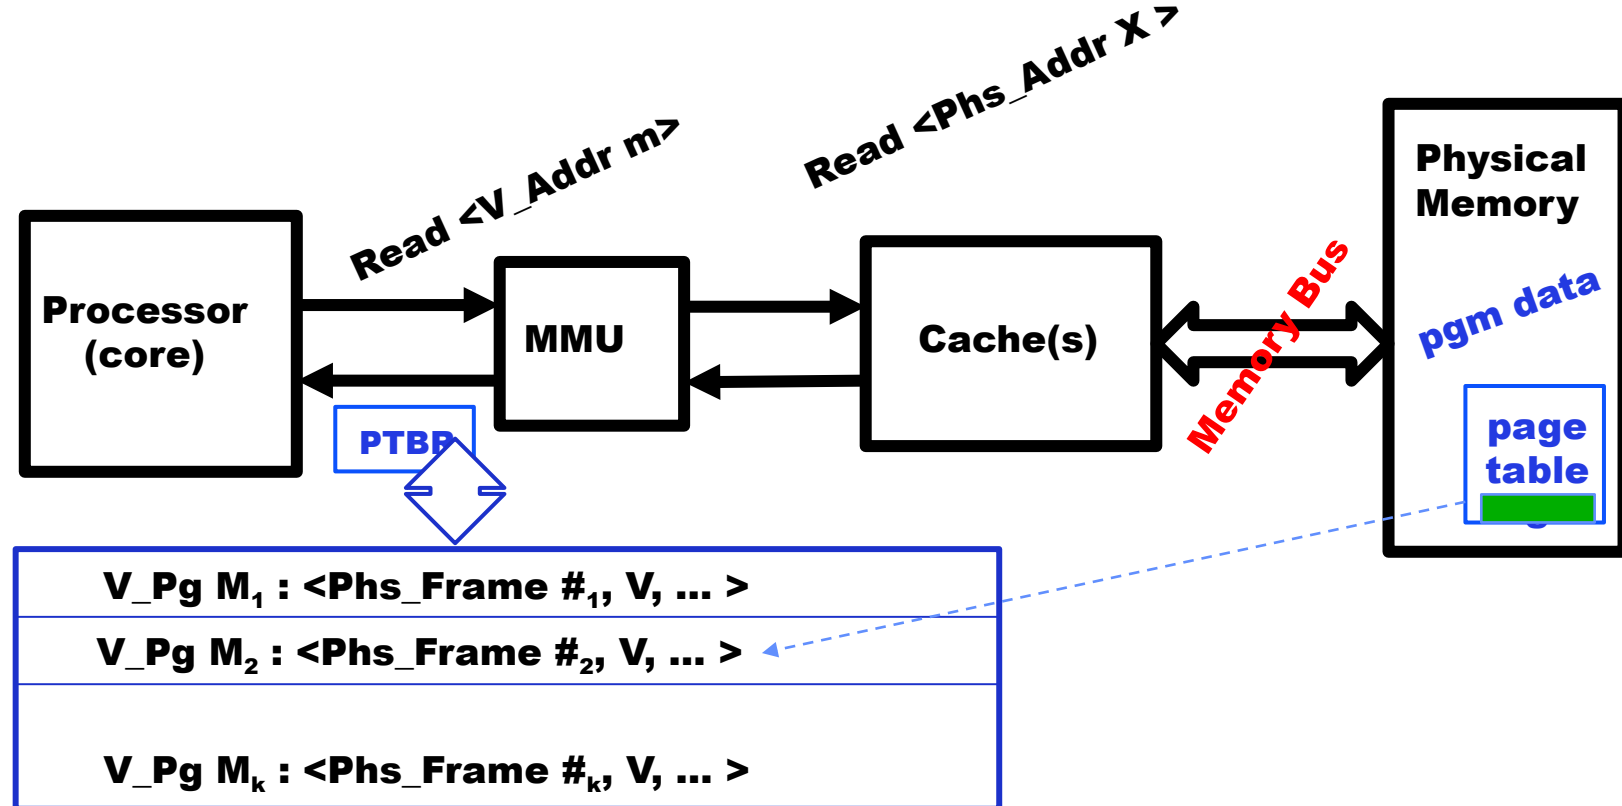
# Recall: Memory Hierarchy



# How do we make Address Translation Fast?

Cache results of recent translations !

Cache Page Table Entries using Virtual Page # as the key



# Translation Look-Aside Buffer

---

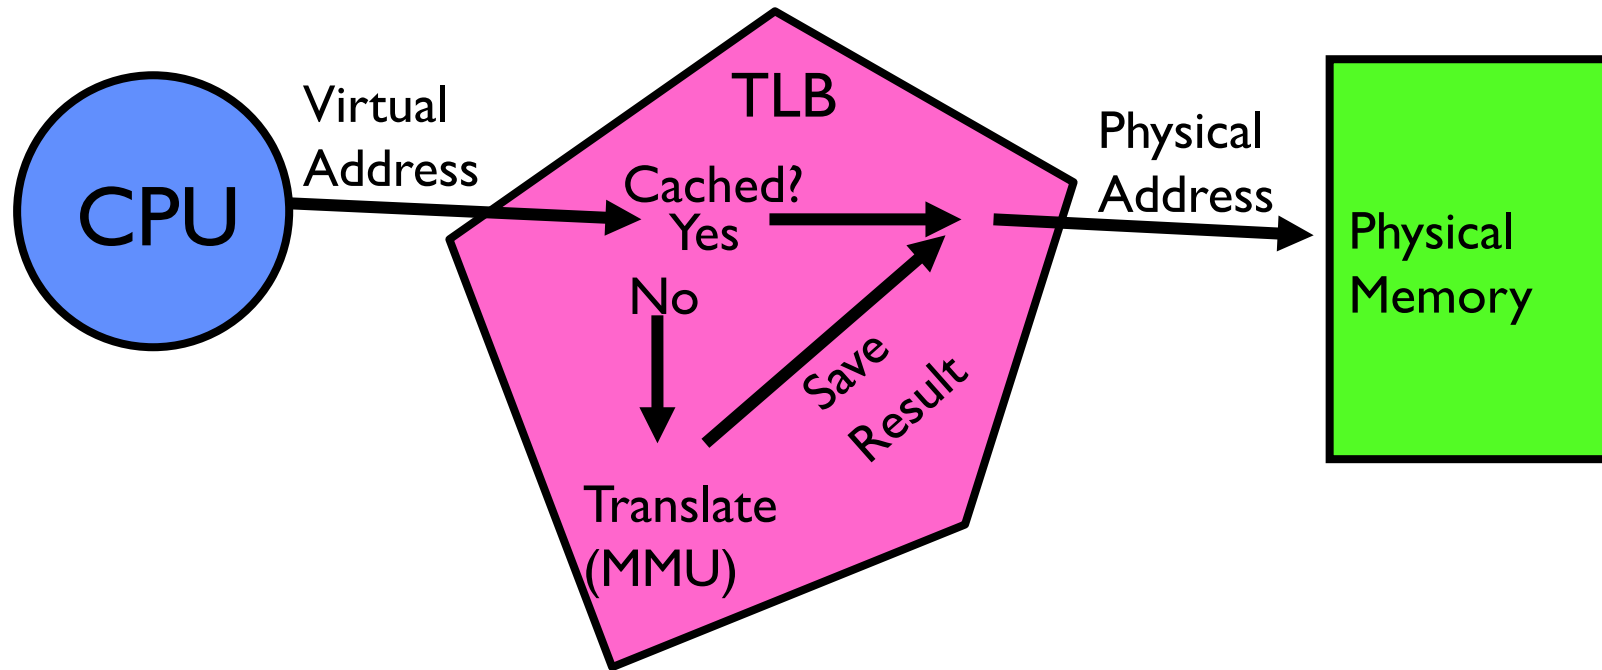
Record recent Virtual Page # to Physical Frame # translation

If present, have the physical address without reading any of the page tables !!!

Caches the end-to-end result

# Caching Applied to Address Translation

---



Does page locality exist?

Instruction accesses spend a lot of time on the same page  
(since accesses sequential)

Stack accesses have definite locality of reference