

Computer Systems II

Li Lu

Room 605, CaoGuangbiao Building

li.lu@zju.edu.cn

<https://person.zju.edu.cn/lynnluli>



About the Class




Talk about this course

- What have you learned from the previous course?
 - Digital Logic, basic instruction set design, CPU design (mainly on single-cycle CPU), etc.
- What will be covered in this course and what you can get from this course?
 - Move forward one step to learn more complex CPU design
 - Begin to explore the principle of Operating Systems
 - Know not only what by also why



How to Prepare for the Class

- Textbook (Computer Organization and Design: The Hardware/Software Interface RISC-V Edition; Operating System Concepts)
- References
- Teaching Components
 - Lectures
 - Labs



Course Topics

- Instruction Classification and Design Principle ~ 1 week
- Concept, Category, Architecture and Design of Pipeline CPU ~ 2 weeks
- Hazard of Pipeline CPU ~ 2 weeks
- Software/Hardware Interfaces ~ 1 week
- Introduction of OS ~ 2 weeks
- Interrupt ~ 1 week
- Process and Thread ~ 2 weeks
- Scheduling, Synchronization and Deadlock ~ 3 weeks
- Final Review ~ 1 week



Instructor & TA

- Instructors
 - Li Lu 卢立 (li.lu@zju.edu.cn)
 - Yajin Zhou 周亚金 (yajin_zhou@zju.edu.cn)
- TAs
 - Yangye Zhou 周扬叶 (1076192792@qq.com)
 - Jingjing Wang 王晶晶 (3200104880@zju.edu.cn)



Course Organization

- Lectures
 - Monday (Zijingang West 1-504)
 - 2:15 PM – 3:50 PM
 - Thursday (Zijingang West 1-504)
 - 2:15 PM – 3:50 PM
- Labs
 - Monday (Zijingang lab room)
 - 4:15 PM – 5:50 PM (Only ODD weeks)
 - Thursday (Zijingang lab room)
 - 4:15 PM – 5:50 PM (Every week)
 - No group, please work alone



Course Grading

- Homework Assignment 10%
 - TBD
- Projects 60%
 - Lab 0 – CPU Design Review
 - Lab 1 – Pipeline CPU Design with Stall
 - Lab 2 – Hazard and Forwarding
 - Lab 3 – Kernel Boot
 - Lab 4 – Interrupt
 - Lab 5 – Simple Scheduling
 - Lab 6 – Running OS on CPU
- Final Exam 30%



Course Policy

- Academic integrity
 - We will strictly enforce the university, college, and department policies against academic dishonesty
 - **Plagiarism in any form will not be tolerated!**
- Unless otherwise noted, work turned in should reflect your independent capabilities
 - If unsure, note / cite sources and help
- Late work penalized **5%/day**
 - No penalty for documented emergency or by prior arrangement in special circumstances



Systems I Review


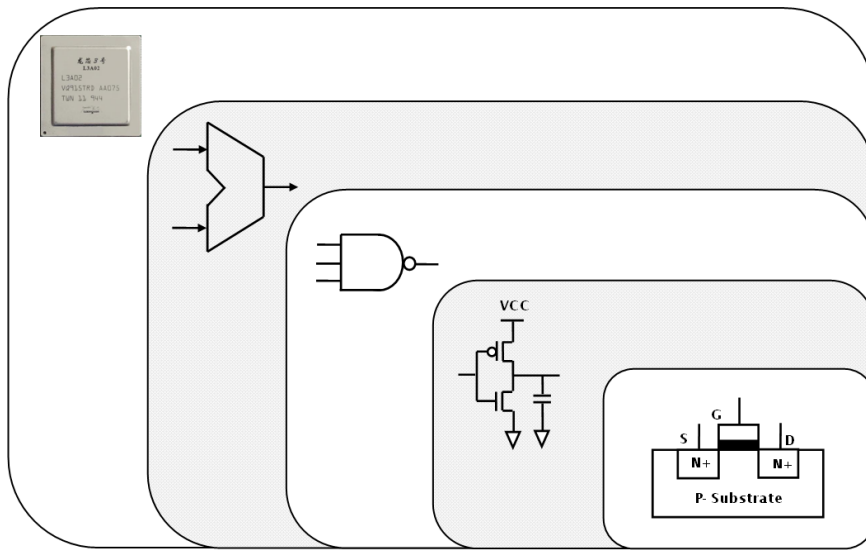


Begin with Answering the Questions

- Why use binary in the computers?
- What is the architecture of computers?
- What is the inside of processor (CPU)?
- Why do? Why using the CPU to solve the problem?
- What are the basic principles of CPU/ISA design?
- How to design a CPU? And how to improve it?



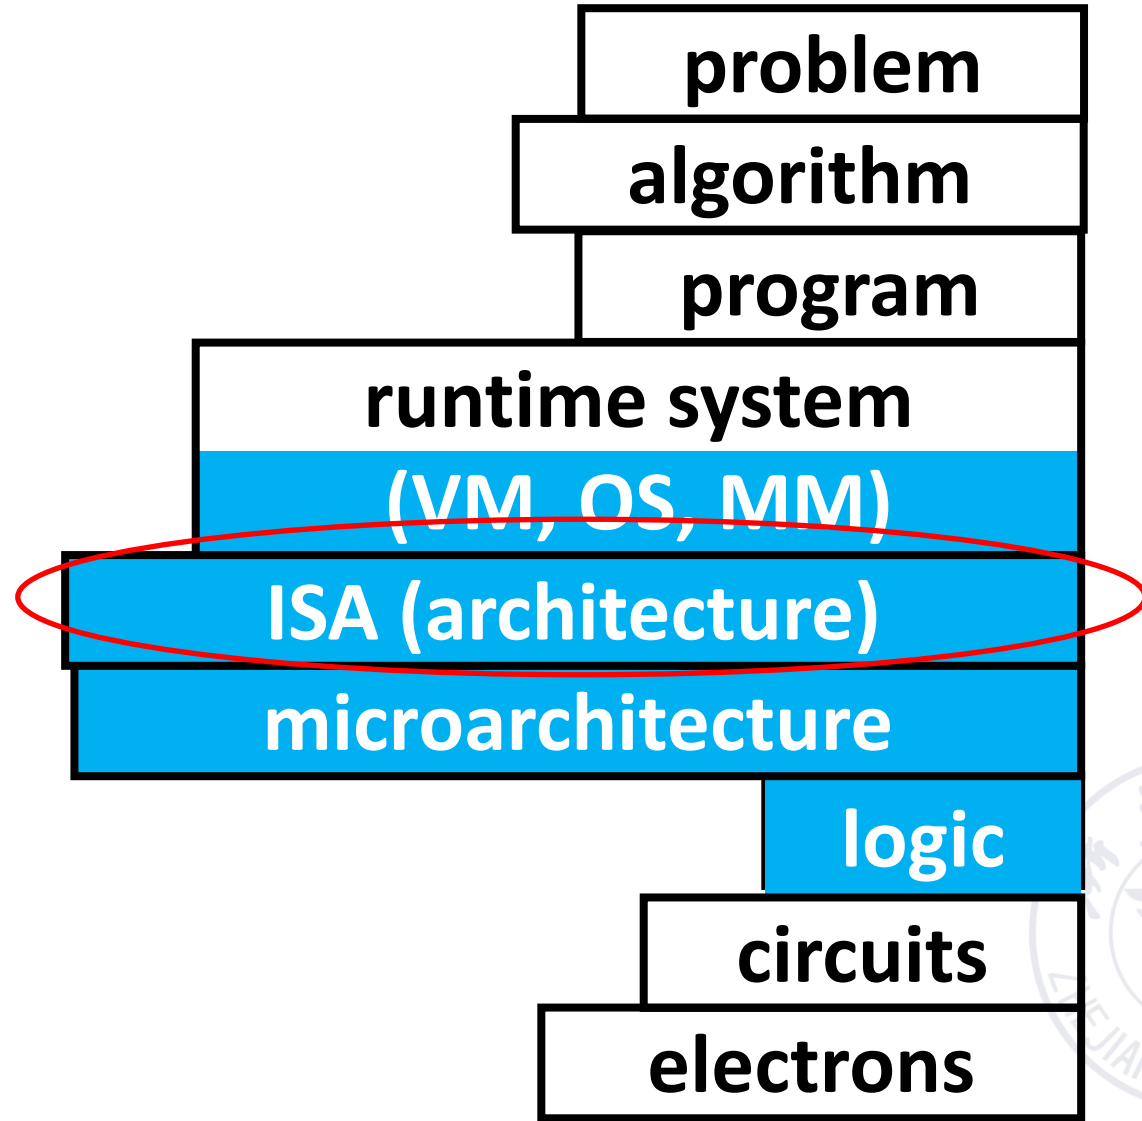
Why use binary in the computers?



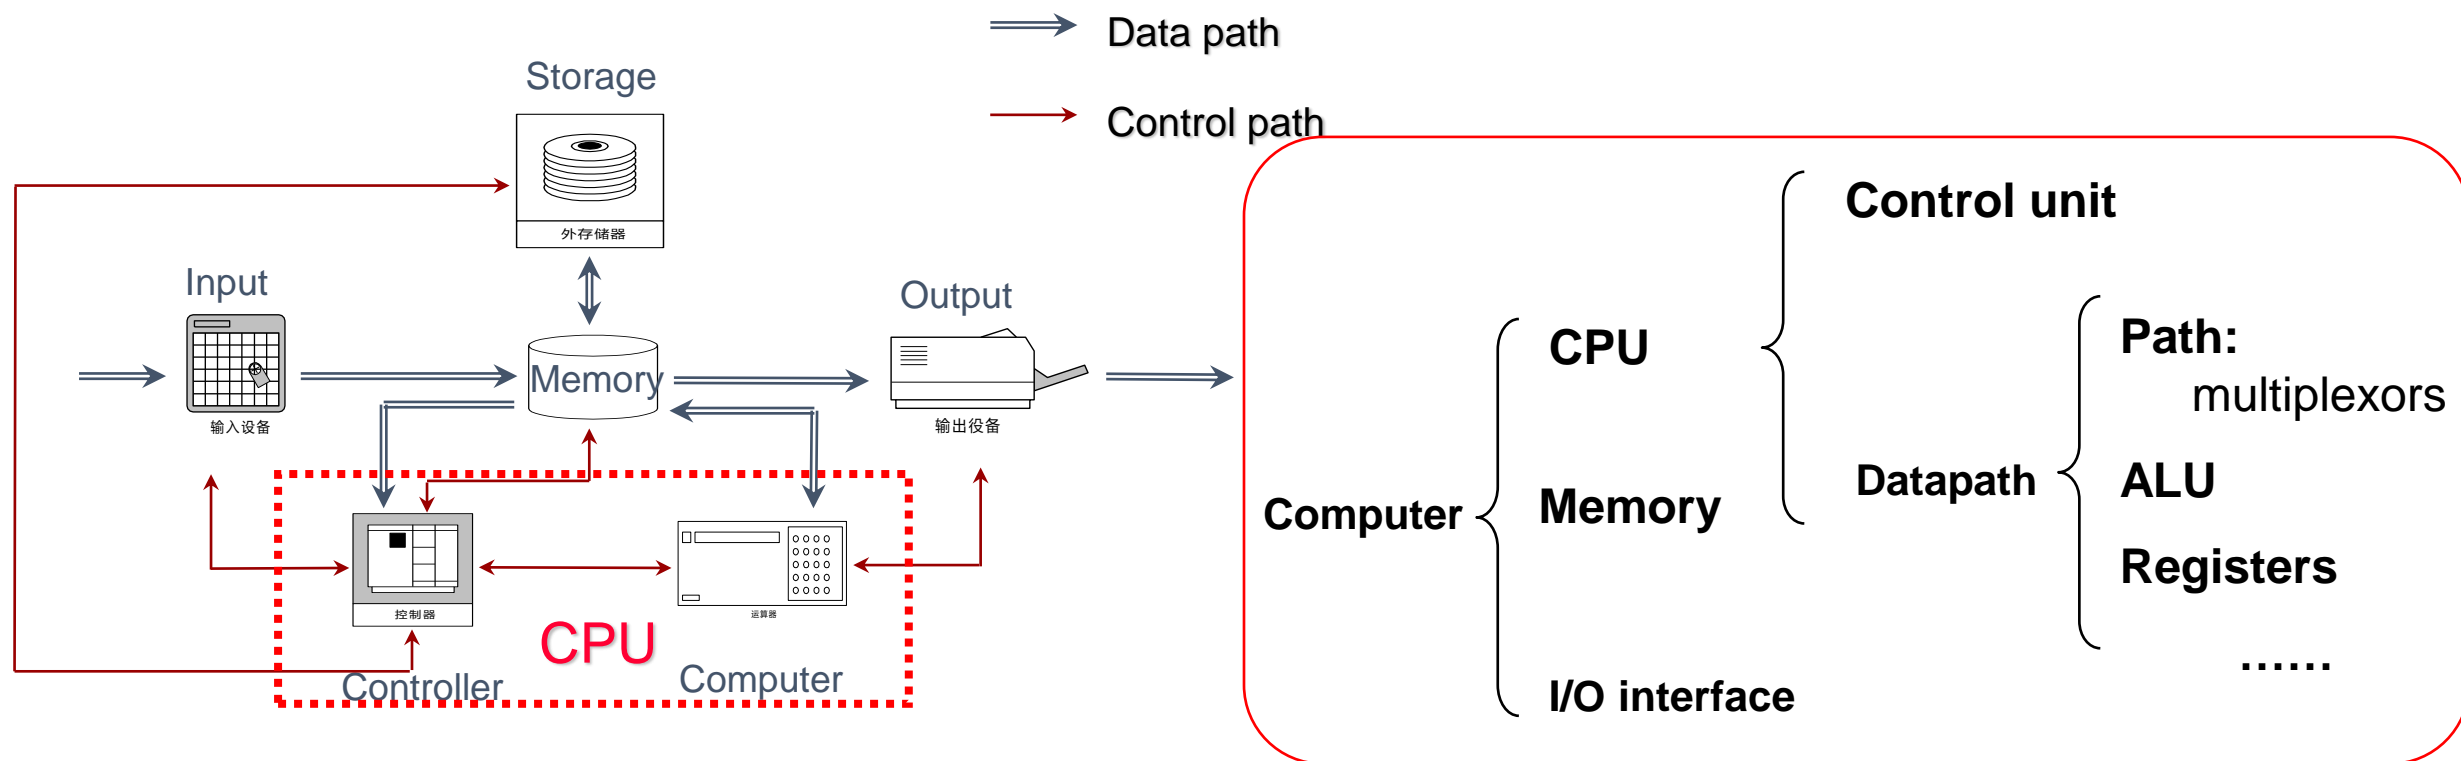
- The computer is composed of electronic components, and binary is the easiest to realize.

What is the architecture of computers?

Computer Architecture!



What is the architecture of computers?



- Von Neumann structure: data and programs are in memory.
- CPU takes instructions and data from memory for operation and puts the results into memory.

Von Neumann Structure



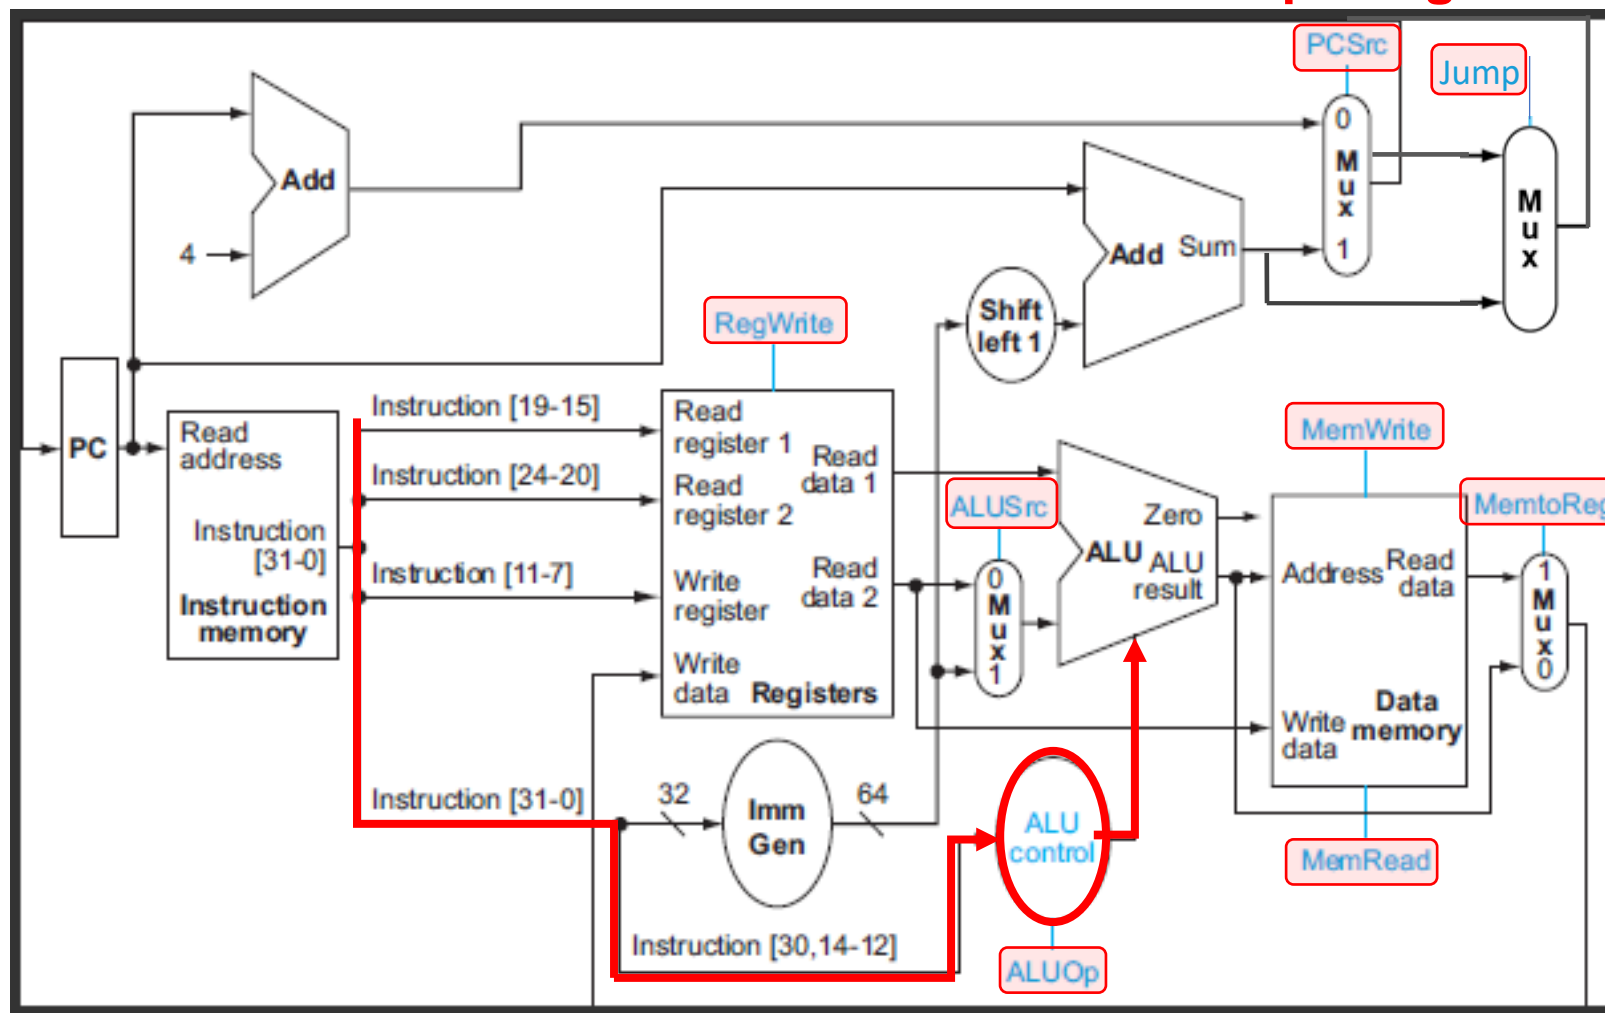
What is the inside of Processor (CPU)?

- Datapath: performs operations on data
- Control: sequences datapath, memory, ...
- Cache memory
 - Small fast SRAM memory for immediate access to data



What is the inside of Processor (CPU)?

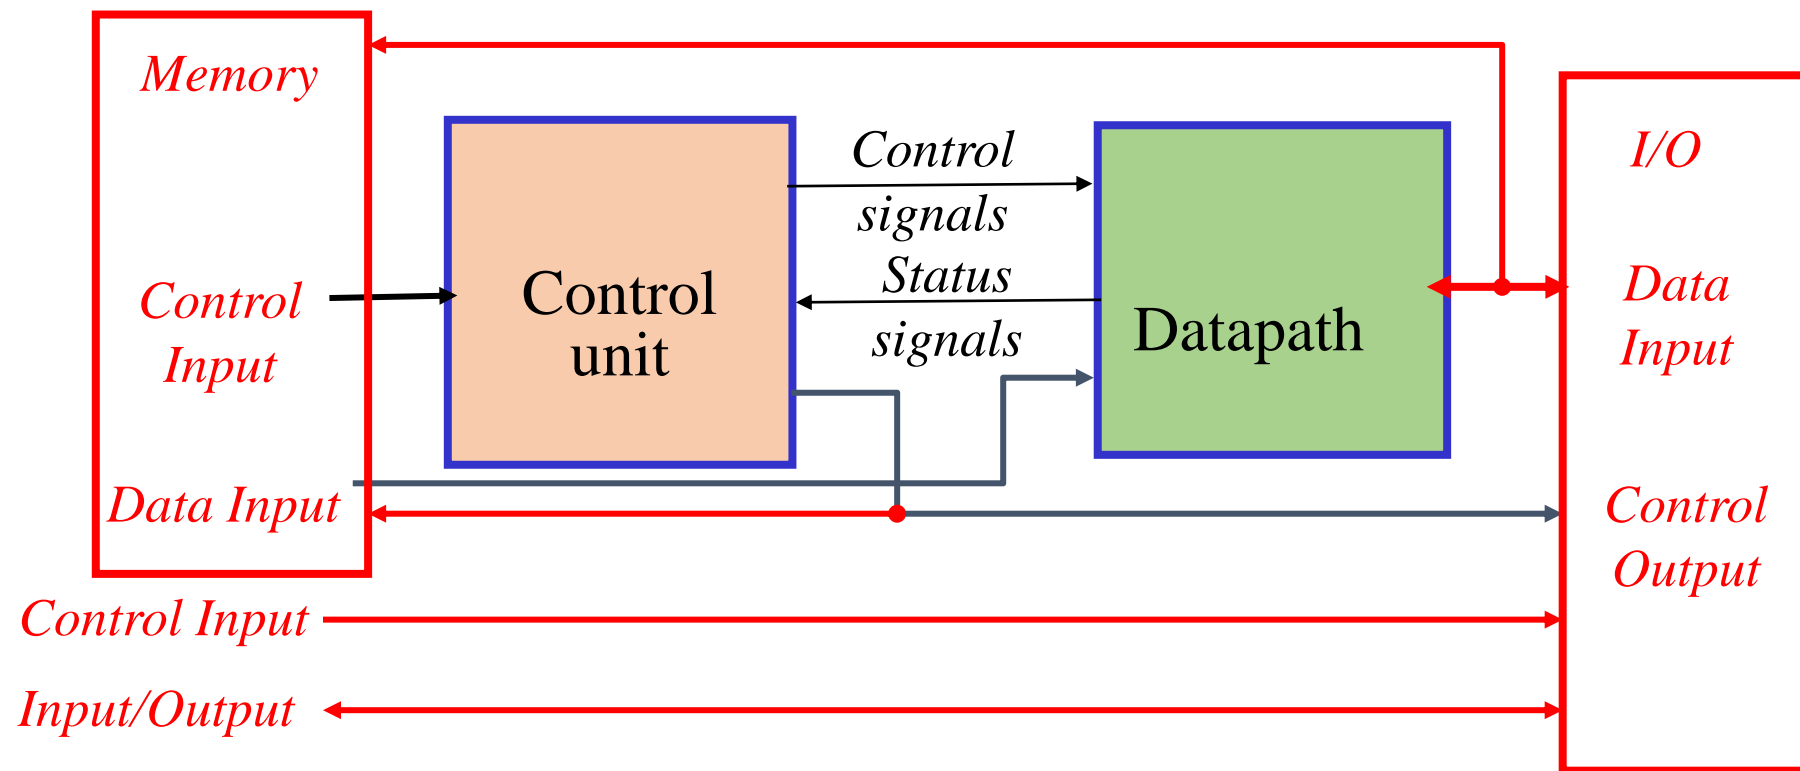
Output signals



Why do?

Why using the CPU to solve the problem?

- Design methodology : General design



Program state machine (PSM)




Why do?

Why using the CPU to solve the problem?

- Better way? Turing Machine

- 1. An infinitely long paper tape TAPE
- 2. A read/write head HEAD
- 3. A set of control rules TABLE
- 4. A state register



Turing Model: Marvin • Minsky (1967)

- Ideal vs. Universal

- Ideal: cannot implement, because no unlimited tape exists
- Universal: Long enough tape with head and tail connected, to replace the infinite long paper tape



Why do?

Why using the CPU to solve the problem?

- CPU
 - General digital system
 - A Turing Machine
 - Implemented by Register Transmission Control Technology
 - Datapath
 - The component of processor that performs arithmetic operations
 - Control
 - The component of processor that commands the datapath, memory, and I/O device according to the instructions of the program



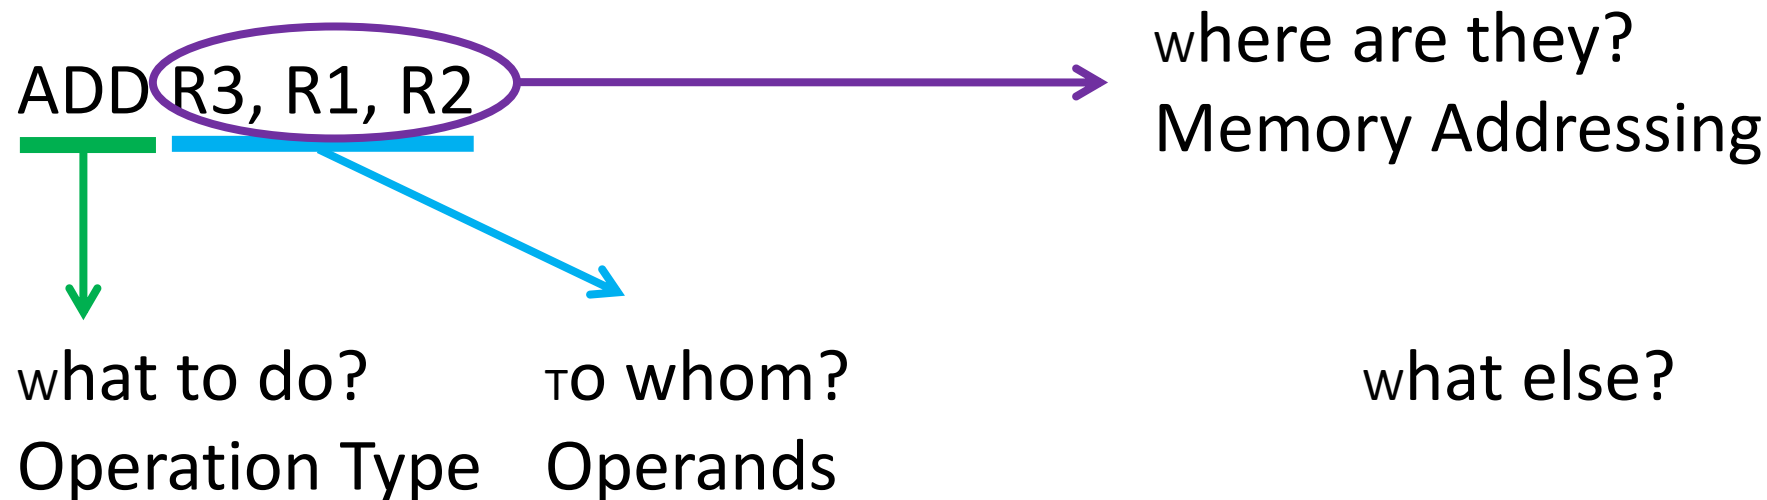
What are the basic principles of CPU/ISA design?

- Instruction Set
 - The instructions repertoire of a computer
 - Different computers have different instruction sets
 - with many aspects in common
 - Early computers had very simple instruction sets
 - with simplified implementation
 - Many modern computers also have simple instruction sets



What are the basic principles of CPU/ISA design?

- Instruction Set **Architecture**



What are the basic principles of CPU/ISA design?

- Instruction Set Architecture (ISA)

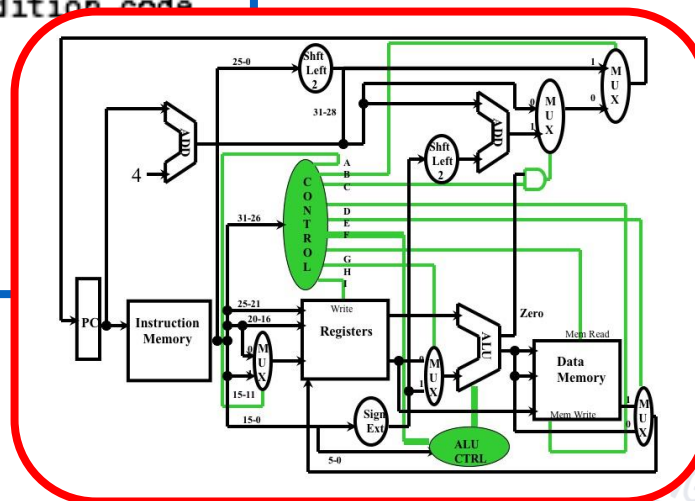
```

1  #include <stdio.h>
2  #include <stdlib.h>
3  void read(int *p);
4  int findmax(int *p);
5  #define N 10
6  int m,n;

```

Programmer-visible instruction set

	LOAD	R1,&a	; R1 <-- contents of 'a'
	LOAD	R2,&a	; R2 <-- contents of 'a'
	TEST	R1,R2	; compare R1 and R2, set condition code
	JNE	@L1	; goto L1 if not equal
	ADD	R1,R2	; R1 <-- R1 + R2
	TEST	R1,R2	; compare R1 and R2, set condition code
	JGE	@L2	; goto L2 if R1 >= R2
	JMP	@END	; goto END
@L1	ADD	R1, R2	; R1 <-- R1 + R2
	JMP	@END	; goto END
@L2	ADD	R1, R2	; R1 <-- R1 + R2
@END	SUB	R2, R3	; R2 <-- R2 - R3



How to Design a CPU?

- Understand ISA
 - Taking RISC-V as example
- A RISC-V ISA is defined as a base integer ISA
 - The base integer ISAs are very similar to that of the early RISC processors (such as MIPS)
 - No branch delay slots
 - Support for optional variable-length instruction encodings
- Goal: A ***standard free*** and ***open*** architecture for industry implementations



RISC-V ISA

32-bit Instruction Formats

31	27	26	25	24	20	19	15	14	12	11	7	6	0
funct7				rs2		rs1		funct3		rd		opcode	
imm[11:0]				rs1		funct3		rd		opcode			
imm[11:5]				rs2		rs1		funct3		imm[4:0]		opcode	
imm[12:10:5]				rs2		rs1		funct3		imm[4:1:11]		opcode	
imm[31:12]				rd		opcode							
imm[20:10:11:19:12]				rd		opcode							

16-bit (RVC) Instruction Formats

CR	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
CI	funct4				rd/rs1				rs2				op			
CSS	funct3		imm		rd/rs1				imm				op			
CIW	funct3		imm						rs2				op			
CL	funct3		imm								rd'		op			
CS	funct3		imm			rs1'			imm		rd'		op			
CB	funct3		imm			rs1'			imm		rs2'		op			
CJ	funct3		offset			rs1'			offset						op	
	funct3		jump target										op			

Base Integer Instructions: RV32I and RV64I						RV Privileged Instructions							
Category	Name	Fmt	RV32I Base		+RV64I	Category	Name	Fmt	RV mnemonic				
Shifts	Shift Left Logical	R	SLL	rd,rs1,rs2	SLLW rd,rs1,rs2	Trap	Mach-mode trap return	R	MRET				
	Shift Left Log. Imm.	I	SLLI	rd,rs1,shamt	SLLIW rd,rs1,shamt		Supervisor-mode trap return	R	SRET				
	Shift Right Logical	R	SRL	rd,rs1,rs2	SRLW rd,rs1,rs2		Interrupt Wait for Interrupt	R	WFI				
	Shift Right Log. Imm.	I	SRLI	rd,rs1,shamt	SRLIW rd,rs1,shamt		MMU Virtual Memory FENCE	R	SFENCE.VMA	rs1,rs2			
	Shift Right Arithmetic	R	SRA	rd,rs1,rs2	SRAW rd,rs1,rs2		Examples of the 60 RV Pseudoinstructions						
	Shift Right Arith. Imm.	I	SRAI	rd,rs1,shamt	SRAIW rd,rs1,shamt		Branch = 0 (BEQ rs,x0,imm)	J	BEQZ	rs,imm			
Arithmetic	ADD	R	ADD	rd,rs1,rs2	ADDW rd,rs1,rs2	Jump (uses JAL x0,imm)	J	J	imm				
	ADD Immediate	I	ADDI	rd,rs1,imm	ADDIW rd,rs1,imm	MoVe (uses ADDI rd,rs,0)	R	MV	rd,rs				
	SUBtract	R	SUB	rd,rs1,rs2	SUBW rd,rs1,rs2	RETurn (uses JALR x0,0,ra)	I	RET					
	Load Upper Imm	U	LUI	rd,imm									
Add Upper Imm to PC	U	AUIPC	rd,imm										
Logical	XOR	R	XOR	rd,rs1,rs2		Optional Compressed (16-bit) Instruction Extension: RV32C	Category	Name	Fmt	RVC	RISC-V equivalent		
	XOR Immediate	I	XORI	rd,rs1,imm			Loads	Load Word	CL	C.LW	rd',rs1',imm	LW	rd',rs1',imm*4
	OR	R	OR	rd,rs1,rs2				Load Word SP	CI	C.LWSP	rd,imm	LW	rd,sp,imm*4
	OR Immediate	I	ORI	rd,rs1,imm				Float Load Word SP	CL	C.FLW	rd',rs1',imm	FLW	rd',rs1',imm*8
	AND	R	AND	rd,rs1,rs2				Float Load Word	CI	C.FLWSP	rd,imm	FLW	rd,sp,imm*8
	AND Immediate	I	ANDI	rd,rs1,imm				Float Load Double	CL	C.FLD	rd',rs1',imm	FLD	rd',rs1',imm*16
Compare	Set <	R	SLT	rd,rs1,rs2				Float Load Double SP	CI	C.FLDSP	rd,imm	FLD	rd,sp,imm*16
	Set < Immediate	I	SLTI	rd,rs1,imm			Stores	Store Word	CS	C.SW	rs1',rs2',imm	SW	rs1',rs2',imm*4
	Set < Unsigned	R	SLTU	rd,rs1,rs2				Store Word SP	CSS	C.SWSP	rs2,imm	SW	rs2,sp,imm*4
	Set < Imm Unsigned	I	SLTIU	rd,rs1,imm				Float Store Word	CS	C.FSW	rs1',rs2',imm	FSW	rs1',rs2',imm*8
Branches	Branch =	B	BEQ	rs1,rs2,imm				Float Store Double	CS	C.FSD	rs1',rs2',imm	FSD	rs1',rs2',imm*16
	Branch ≠	B	BNE	rs1,rs2,imm				Float Store Double SP	CSS	C.FSDSP	rs2,imm	FSD	rs2,sp,imm*16
	Branch <	B	BLT	rs1,rs2,imm			Arithmetic	ADD	CR	C.ADD	rd,rs1	ADD	rd,rd,rs1
	Branch ≥	B	BGE	rs1,rs2,imm				ADD Immediate	CI	C.ADDI	rd,imm	ADDI	rd,rd,imm
	Branch < Unsigned	B	BLTU	rs1,rs2,imm				ADD SP Imm * 16	CI	C.ADDI16SP	x0,imm	ADDI	sp,sp,imm*16
	Branch ≥ Unsigned	B	BGEU	rs1,rs2,imm				ADD SP Imm * 4	CIW	C.ADDI4SPN	rd',imm	ADDI	rd',sp,imm*4
Jump & Link	J&L	J	JAL	rd,imm	SUB		CR	C.SUB	rd,rs1	SUB	rd,rd,rs1		
	Jump & Link Register	I	JALR	rd,rs1,imm	AND		CR	C.AND	rd,rs1	AND	rd,rd,rs1		
Synch	Synch thread	I	FENCE		AND Immediate		CI	C.ANDI	rd,imm	ANDI	rd,rd,imm		
	Synch Instr & Data	I	FENCE.I		OR		CR	C.OR	rd,rs1	OR	rd,rd,rs1		
Environment	CALL	I	ECALL		eXclusive OR		CR	C.XOR	rd,rs1	AND	rd,rd,rs1		
	BREAK	I	EBREAK		MoVe		CR	C.MV	rd,rs1	ADD	rd,rs1,x0		
					Load Immediate		CI	C.LI	rd,imm	ADDI	rd,x0,imm		
					Load Upper Imm		CI	C.LUI	rd,imm	LUI	rd,imm		
Control Status Register (CSR)	Read/Write	I	CSRWR	rd,csr,rs1			Shifts	Shift Left Imm	CI	C.SLLI	rd,imm	SLLI	rd,rd,imm
	Read & Set Bit	I	CSRRS	rd,csr,rs1				Shift Right Ar. Imm.	CI	C.SRAI	rd,imm	SRAI	rd,rd,imm
	Read & Clear Bit	I	CSRRC	rd,csr,rs1				Shift Right Log. Imm.	CI	C.SRLI	rd,imm	SRLI	rd,rd,imm
	Read/Write Imm	I	CSRRI	rd,csr,imm		Branches	Branch=0	CB	C.BEQZ	rs1',imm	BEQ	rs1',x0,imm	
	Read & Set Bit Imm	I	CSRRII	rd,csr,imm			Branch≠0	CB	C.BNEZ	rs1',imm	BNE	rs1',x0,imm	
	Read & Clear Bit Imm	I	CSRRCI	rd,csr,imm		Jump	Jump	CJ	C.J	imm	JAL	x0,imm	
Loads	Load Byte	I	LB	rd,rs1,imm			Jump Register	CR	C.JR	rd,rs1	JALR	x0,rs1,0	
	Load Halfword	I	LH	rd,rs1,imm		Jump & Link	J&L	CJ	C.JAL	imm	JAL	ra,imm	
	Load Byte Unsigned	I	LBU	rd,rs1,imm		Jump & Link Register	CR	C.JALR	rs1	JALR	ra,rs1,0		
	Load Half Unsigned	I	LHU	rd,rs1,imm		System Env. BREAK	CI	C.EBREAK		EBREAK			
Load Word	I	LW	rd,rs1,imm		+RV64I	Optional Compressed Extention: RV64C							
Stores	Store Byte	S	SB	rs1,rs2,imm		LWU	rd,rs1,imm	All RV32C (except C.JAL, 4 word loads, 4 word stores) plus:					
	Store Halfword	S	SH	rs1,rs2,imm		LD	rd,rs1,imm	ADD Word (C.ADDW) Load Doubleword (C.LD)					
	Store Word	S	SW	rs1,rs2,imm				ADD Imm. Word (C.ADDIW) Load Doubleword SP (C.LDSP)					
								SUBtract Word (C.SUBW) Store Doubleword (C.SD)					
								Store Doubleword SP (C.SDSP)					

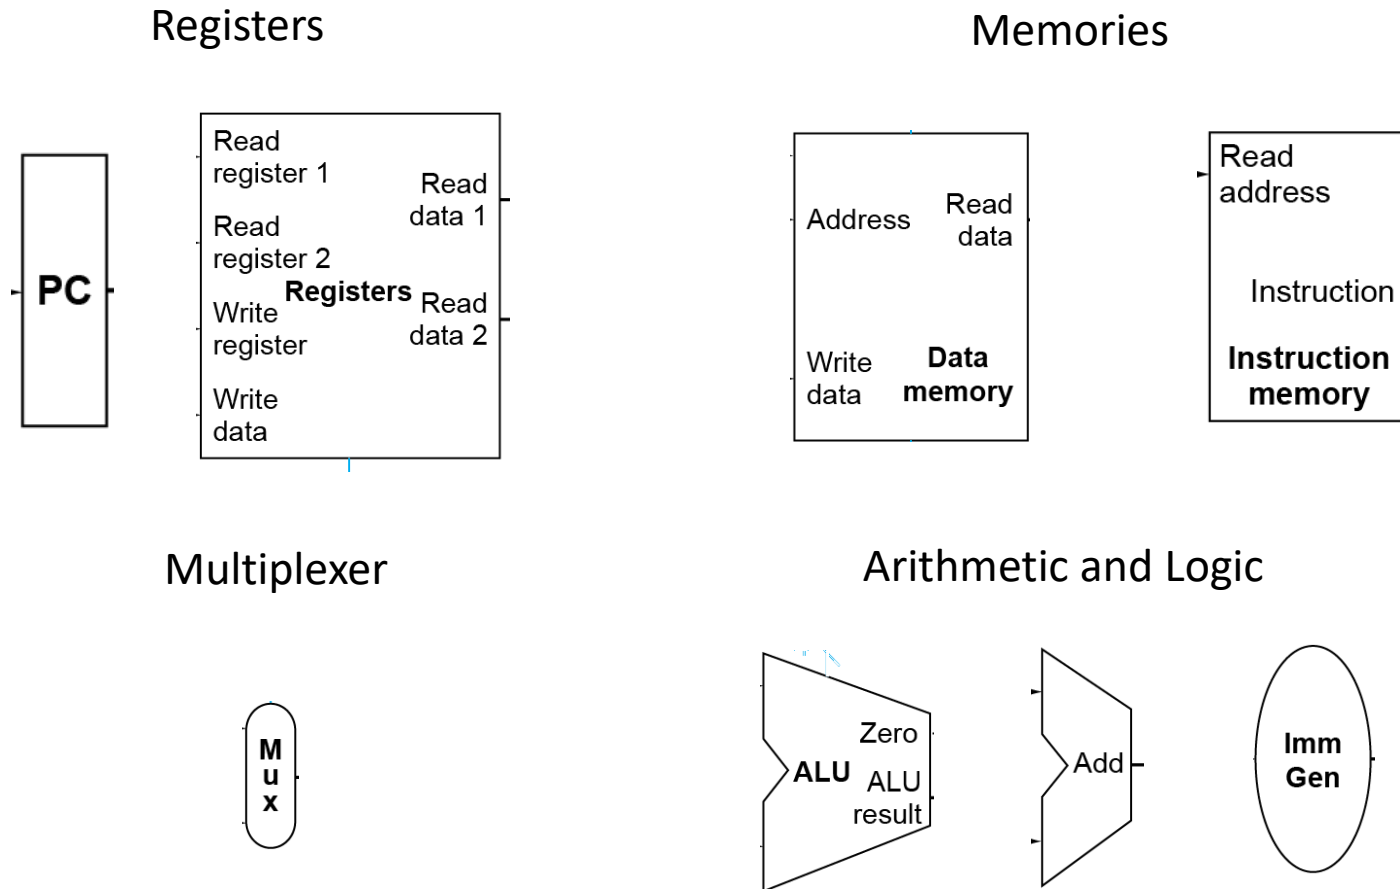
How to Design a CPU?

- Understand instruction execution of CPU
 - **Fetch :**
 - Take instructions from the instruction memory
 - Modify PC to point the next instruction
 - **Instruction decoding & Read Operand:**
 - Will be translated into machine control command
 - Reading Register Operands, whether or not to use
 - **Executive Control:**
 - Control the implementation of the corresponding ALU operation
 - **Memory access:**
 - Write or Read data from memory
 - Only ld/sd
 - **Write results to register:**
 - If it is R-type instructions, ALU results are written to rd
 - If it is I-type instructions, memory data are written to rd
 - **Modify PC** for branch instructions



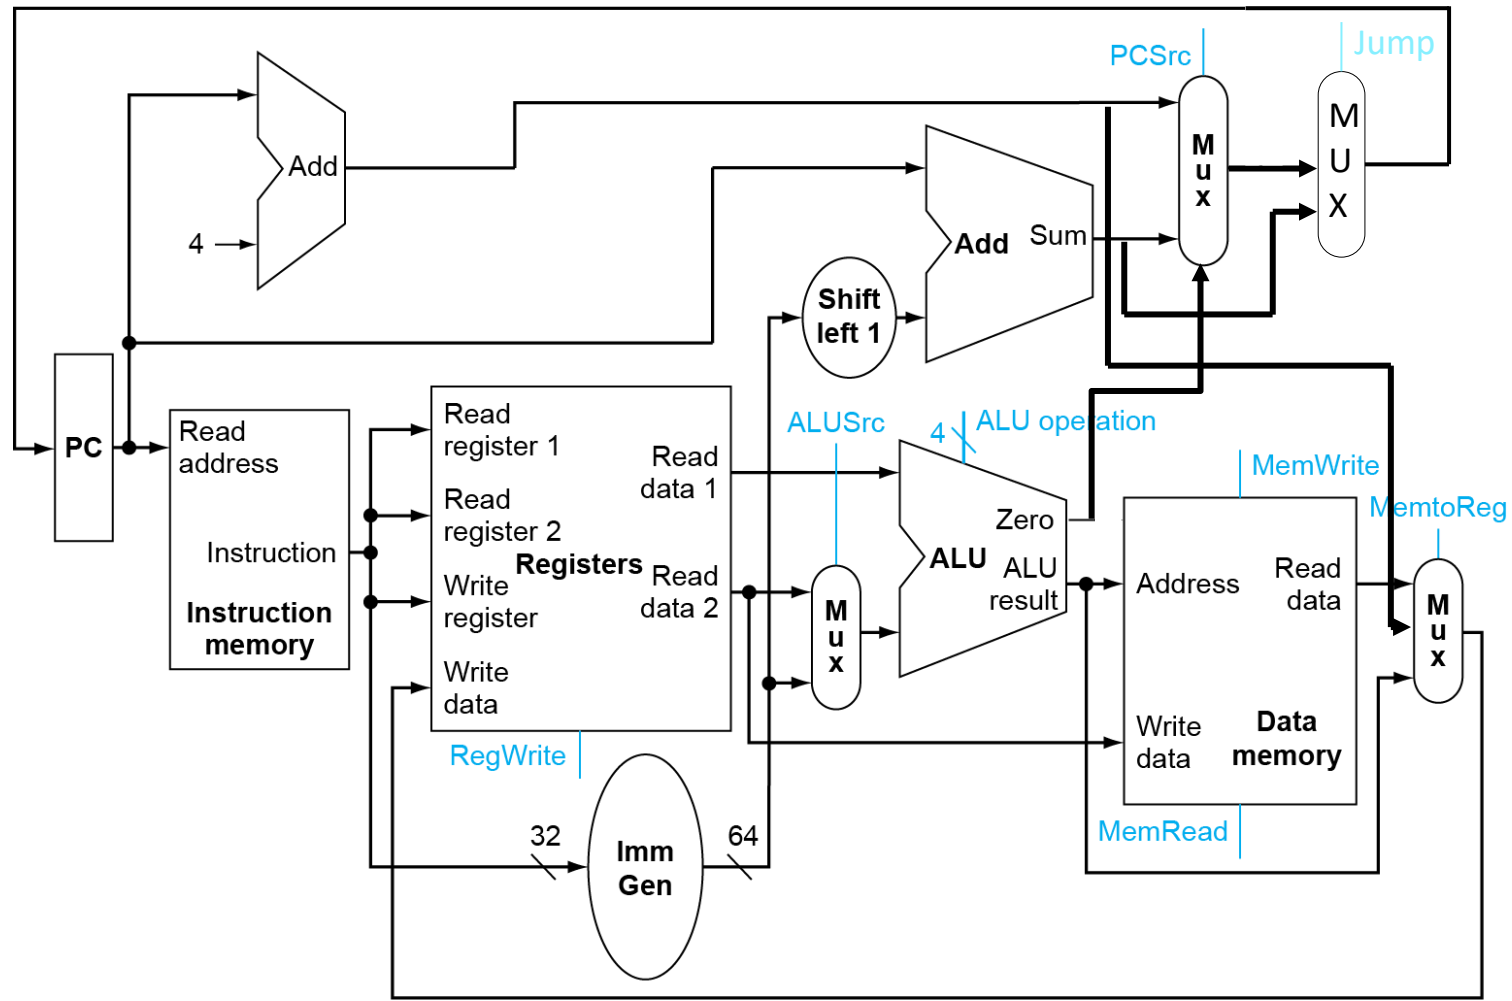
How to Design a CPU?

- Knowing the elements



How to Design a CPU?


- Construct the datapath



How to Design a CPU?

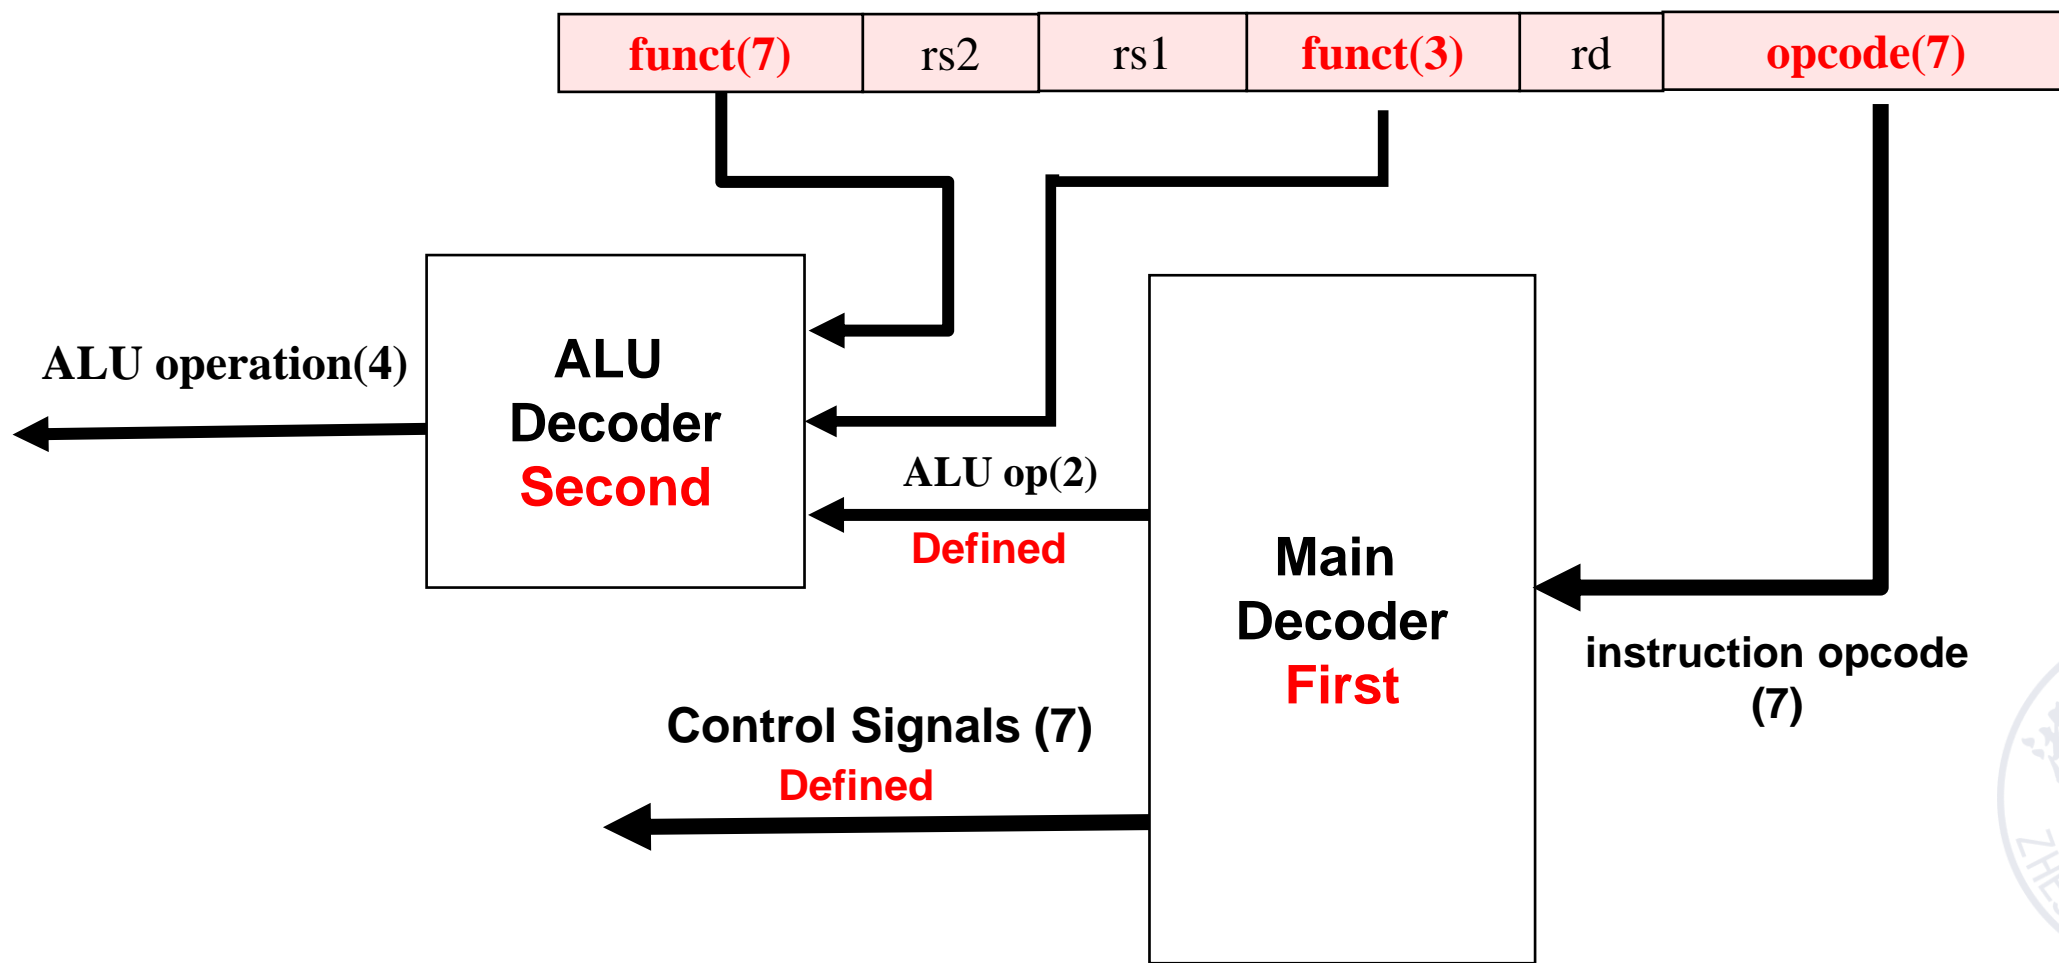
- Construct the controller
 - **Information** comes from the 32 bits of the instruction
 - Selecting the **operations** to perform (ALU, read/write, etc.)
 - Controlling the **flow of data** (multiplexor inputs)
 - ALU's operation based on **instruction type** and **function** code

Input		Output								
Instruction	OPCode	ALUSrcB	Memto-Reg	Reg Write	Mem Read	Mem Write	Branch	Jump	ALU Op1	ALU Op0
R-format	0110011	0	00	1	0	0	0	0	1	0
Ld(I-Type)	0000011	1	01	1	1	0	0	0	0	0
sd(S-Type)	0100011	1	x	0	0	1	0	0	0	0
beq(B-Type)	1100111	0	x	0	0	0	1	0	0	1
Jal(J-Type)	1101111	x	10	1	0	0	0	1	x	x



How to Design a CPU?

- Construct the controller



Why not Single-Cycle?

- Longest delay determines clock period
 - Critical path: load instruction
 - Instruction memory → register file → ALU → data memory → register file
- Waste of area. If the instruction needs to use some functional unit multiple times.
 - E.g., the instruction 'mult' needs to use the ALU repeatedly. So, the CPU will be very large
- Any solution for performance improvement?

Instr	Instr fetch	Register read	ALU op	Memory access	Register write	Total time
ld	200ps	100 ps	200ps	200ps	100 ps	800ps
sd	200ps	100 ps	200ps	200ps		700ps
R-type	200ps	100 ps	200ps		100 ps	600ps
beq	200ps	100 ps	200ps			500ps



How to Improve the CPU?

- Reduce the number of instructions
 - Make instructions that *do* more (CISC)
 - User better compilers
- Use less cycles to perform the instruction
 - Simpler instructions (RISC)
 - Use multiple units/ALUs/cores in parallel
- Increase the clock frequency
 - Find a *newer* technology to manufacture
 - Redesign time critical components
 - Adopt multi-cycle

Multi-Cycle CPU Design

- Single-Cycle microarchitecture
 - + Simple
 - Clock cycle time limited by longest instruction (lw)
 - Two adders/ALUs and two memories
- Multi-Cycle microarchitecture
 - + Shorter clock cycle period
 - + Simpler instructions run faster
 - + Reuse expensive hardware on multiple cycles
 - Sequencing overhead paid many times

Basic Principles of Multi-Cycle CPU Design

- Separate the execution of instructions into several stages with the same period
 - Hard to achieve the same period, thus aims to be *almost the same*
- Each stage occupies one clock cycle
- Each clock cycle at most completes a memory access, register access, or ALU operation
- Execution result of the previous clock cycle needs to be stored in specific sequential logic components
- Clock cycle period depends on the most complex operation

Why not Multi-Cycle?

- Even Longer Time for Instruction Execution
 - E.g., 92.5s (Single-cycle) < 133.9s (Multi-cycle)
 - Cannot achieve expected performance compared with single-cycle
- But provide a new perspective for CPU design
 - Finer-grained execution in one clock cycle
- We will improve performance by *pipelining*

Instruction Set Principles



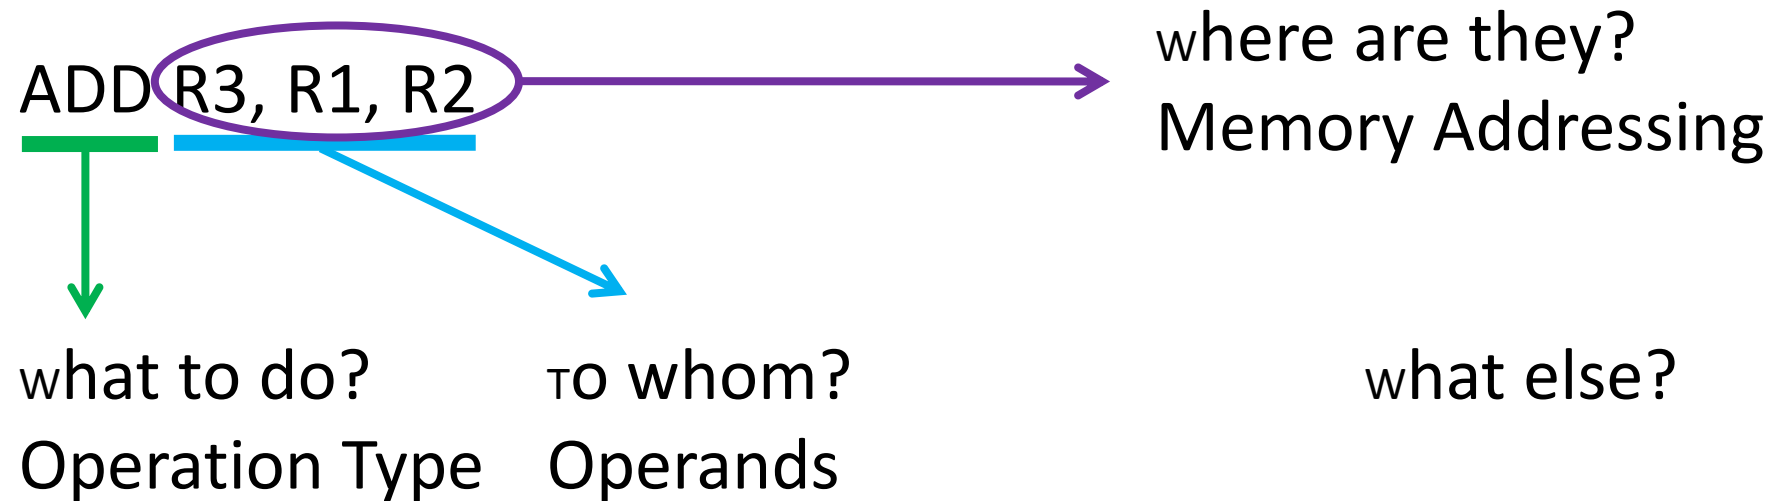
What are the basic principles of ISA design?

- Instruction Set
 - The instructions repertoire of a computer
 - Different computers have different instruction sets
 - with many aspects in common
 - Early computers had very simple instruction sets
 - with simplified implementation
 - Many modern computers also have simple instruction sets



What are the basic principles of ISA design?

- Instruction Set **Architecture**



What are the basic principles of ISA design?

- Instruction Set Architecture (ISA)

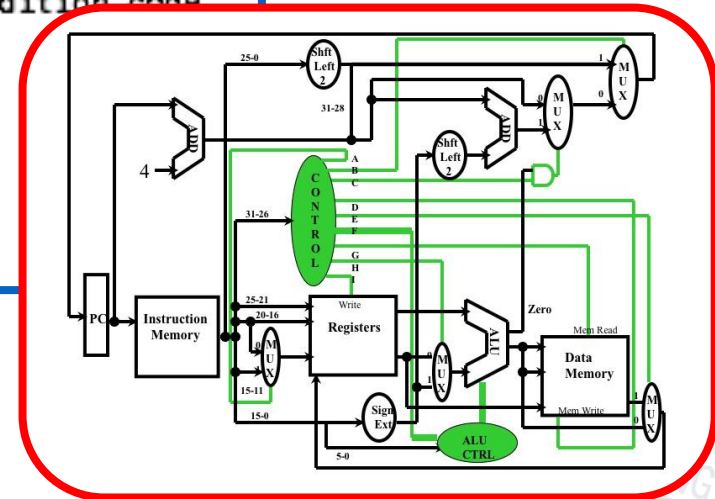
```
1  #include <stdio.h>
2  #include <stdlib.h>
3  void read(int *p);
4  int findmax(int *p);
5  #define N 10
6  int m,n;
```

Programmer-visible instruction set

```

LOAD      R1,&a      ; R1 <-- contents of 'a'
LOAD      R2,&a      ; R2 <-- contents of 'a'
TEST      R1,R2      ; compare R1 and R2, set condition code
JNE       @L1        ; goto L1 if not equal
ADD       R1,R2      ; R1 <-- R1 + R2
TEST      R1,R2      ; compare R1 and R2, set condition code
JGE       @L2        ; goto L2 if R1 >= R2
JMP       @END       ; goto END
@L1      ADD      R1, R2      ; R1 <-- R1 + R2
JMP       @END       ; goto END
@L2      ADD      R1, R2      ; R1 <-- R1 + R2
@END     SUB      R2, R3      ; R2 <-- R2 - R3

```

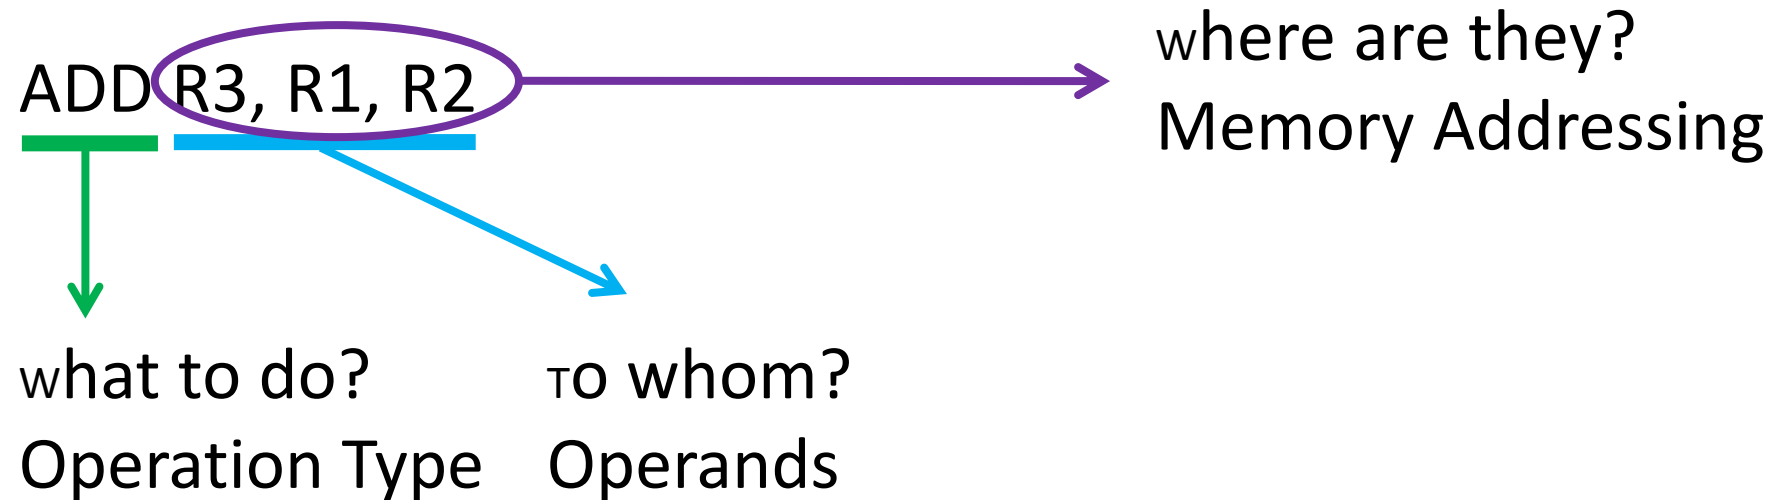


What are the basic principles of ISA design?

- Compatibility
- Versatility
- High efficiency
- Security



ISA Classification Basis




ISA Classification Basis

- The types of internal storage:

- Stack
- Accumulator
- Register

In processor, stores data fetched from memory or cache




ISA Classes

- Stack architecture
- Accumulator architecture
- General-purpose register architecture (GPR)



ISA Classes: Stack Architecture

- Implicit Operands
on the **T**op **O**f the **S**tack (TOS)
- $C = A + B$ (memory locations)
Push A
Push B
Add
Pop C



ISA Classes: Stack Architecture

- Implicit Operands
on the **T**op **O**f the **S**tack (TOS)


- $C = A + B$ (memory locations)

Push A

Push B

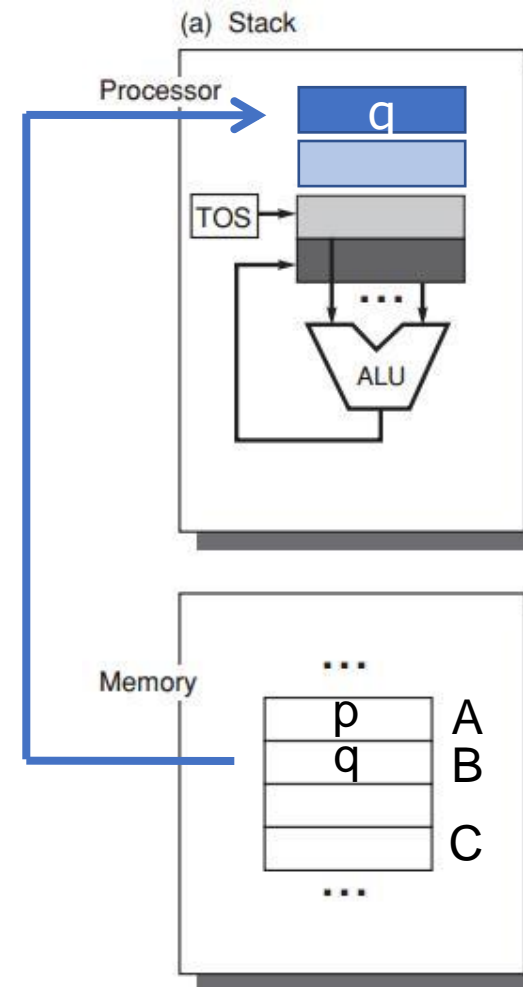
Add

Pop C



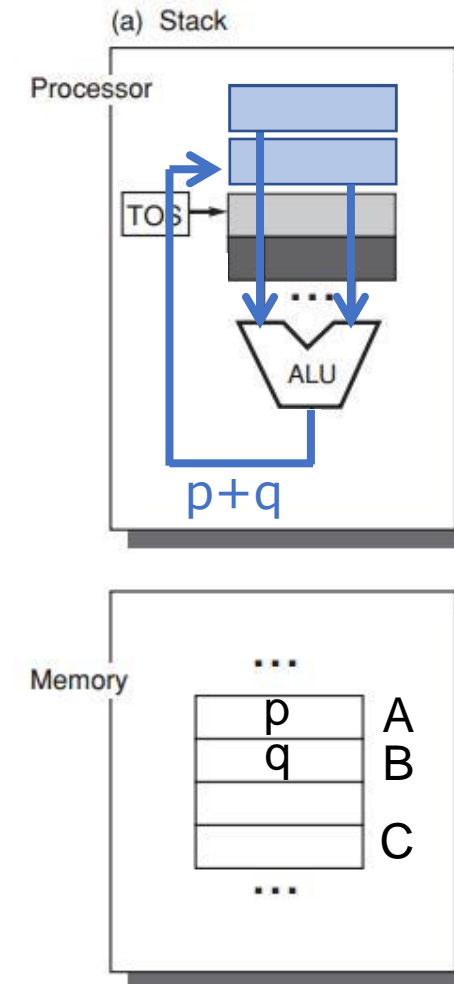
ISA Classes: Stack Architecture

- Implicit Operands
on the **T**op **O**f the **S**tack (TOS)
- $C = A + B$ (memory locations)
Push A
Push B
Add
Pop C




ISA Classes: Stack Architecture

- Implicit Operands
on the **T**op **O**f the **S**tack (TOS)
- First operand removed from
second op replaced by the result
- $C = A + B$ (memory locations)
Push A
Push B
Add
Pop C



ISA Classes: Stack Architecture


- Implicit Operands
on the **T**op **O**f the **S**tack (TOS)
- First operand removed from
second op replaced by the result
- $C = A + B$ (memory locations)
Push A
Push B
Add
Pop C



ISA Classes: Stack Architecture


- Example:

$$98 * (12 + 45)$$




ISA Classes: Accumulator Architecture

- One implicit operand: the accumulator
one explicit operand: mem location
- $C = A + B$
Load A
Add B
Store C
- Accumulator is both an implicit input operand and a result




ISA Classes: Accumulator Architecture

- One implicit operand: the accumulator
one explicit operand: mem location
- $C = A + B$
Load A
Add B
Store C
- Accumulator is both an implicit input operand and a result




ISA Classes: Accumulator Architecture

- One implicit operand: the accumulator
one explicit operand: mem location
- $C = A + B$
Load A
Add B
Store C
- Accumulator is both an implicit input operand and a result



ISA Classes: Accumulator Architecture

- One implicit operand: the accumulator
one explicit operand: mem location
- $C = A + B$
Load A
Add B
Store C
- Accumulator is both an implicit input operand and a result



ISA Classes: General-Purpose Register Arch

- Only explicit operands
 - registers
 - memory locations
- Operand access:
 - direct memory access
 - loaded into temporary storage first

