

# CS162

## Operating Systems and Systems Programming

### Lecture 7

## Concurrency

Professor Natacha Crooks

<https://cs162.org/>

# Correctness Requirements

---

Threaded programs must work for all interleavings of thread instruction sequences

Cooperating threads inherently non-deterministic and non-reproducible

Really hard to debug unless carefully designed!

# The Importance of Milk

---



# The Importance of Milk

---

Great thing about OS's – analogy between problems in OS and problems in real life

Help you understand real life problems better

But, computers are much stupider than people

# Solve with a lock?

---

Lock prevents someone from doing something


- Lock before entering critical section
- Unlock when leaving
- Wait if locked



Fix the milk problem by putting a key on the refrigerator

Lock it and take key if you are going to go buy milk

Fixes too much: roommate angry if only wants OJ



# Too Much Milk: Correctness Properties

---

What are the correctness properties for the “Too much milk” problem???

- Never more than one person buys
- Someone buys if needed

First attempt: Restrict ourselves to use only atomic load and store operations as building blocks

# Too Much Milk: Solution #1


---

Use a note to avoid buying too much milk:

- Leave a note before buying (kind of “lock”)
- Remove note after buying (kind of “unlock”)
- Don’t buy if note (wait)

Suppose a computer tries this  
(remember, only memory read/write are atomic)

```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove note;  
    }  
}
```



# Too Much Milk: Solution #1

---

<u>Thread A</u>	<u>Thread B</u>
if (noMilk) {	if (noMilk) {
if (noNote) {	if (noNote) {
leave Note;	
buy Milk;	
remove Note;	
}	
}	
	leave Note;
	buy Milk;
	}
	}
	remove Note;



# Too Much Milk: Solution #1


---

Still too much milk **but only occasionally!**

Thread can get context switched after checking milk and note but before buying milk!

Solution makes problem worse since fails **intermittently**

- Makes it really hard to debug...
- Must work despite what the dispatcher does!



# Too Much Milk: Solution #1½

---

Let's try to fix this by placing note first

```
leave Note;  
if (noMilk) {  
    if (noNote) {  
        buy milk;  
    }  
}  
remove Note;
```

What happens here?

- Well, with human, probably nothing bad
- With computer: no one ever buys milk

# Too Much Milk Solution #2

---

How about labeled notes?

- Now we can leave note before checking

Algorithm looks like this:

<u>Thread A</u>	<u>Thread B</u>
leave note A;	leave note B;
if (noNote B) {	if (noNoteA) {
if (noMilk) {	if (noMilk) {
buy Milk;	buy Milk;
}	}
}	}
remove note A;	remove note B;

## Too Much Milk Solution #2

---

Possible for neither thread to buy milk

- Context switches at exactly the wrong times can lead each to think that the other is going to buy

Really insidious:

- Extremely unlikely this would happen, but will at worse possible time
  - Probably something like this in UNIX

# Too Much Milk Solution #2: problem!

---

*I'm not getting milk, You're getting milk*

This kind of lockup is called “starvation!”

# Too Much Milk Solution #3

---

<u>Thread A</u>	<u>Thread B</u>
leave note A;	leave note B;
while (note B) {\X	if (noNote A) {\Y
do nothing;	if (noMilk) {
}	buy milk;
if (noMilk) {	}
buy milk;	}
}	remove note B;
remove note A;	

# Too Much Milk Solution #3

---

Both can guarantee that:

- It is safe to buy, or
- Other will buy, ok to quit

At X:

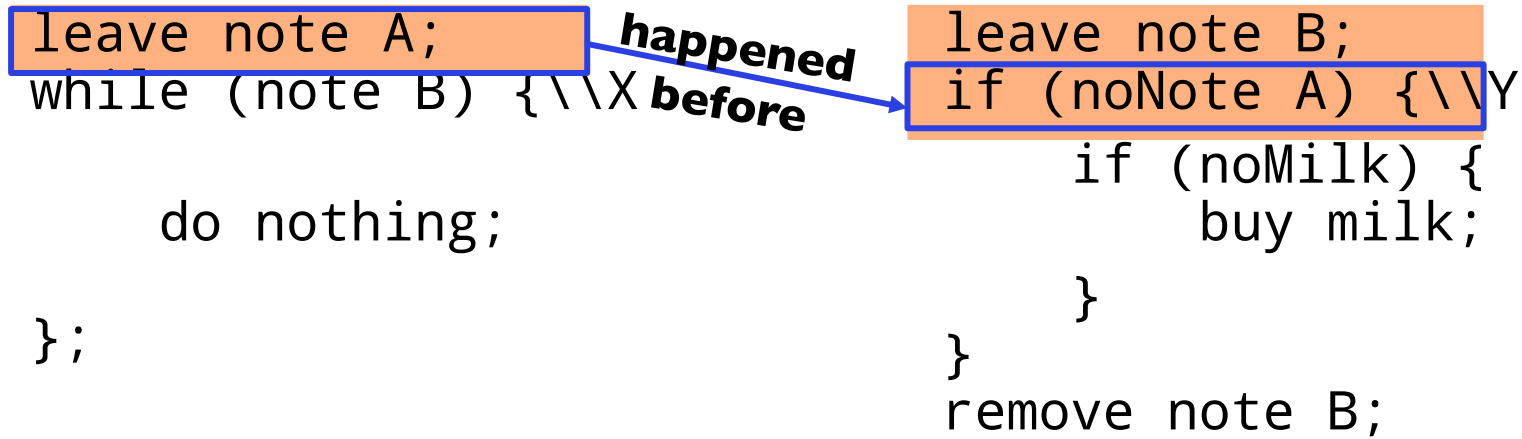
- If no note B, safe for A to buy,
- Otherwise wait to find out what will happen

At Y:

- If no note A, safe for B to buy
- Otherwise, A is either buying or waiting for B to quit

# Case 1

- “leave note A” happens before “if (noNote A)”



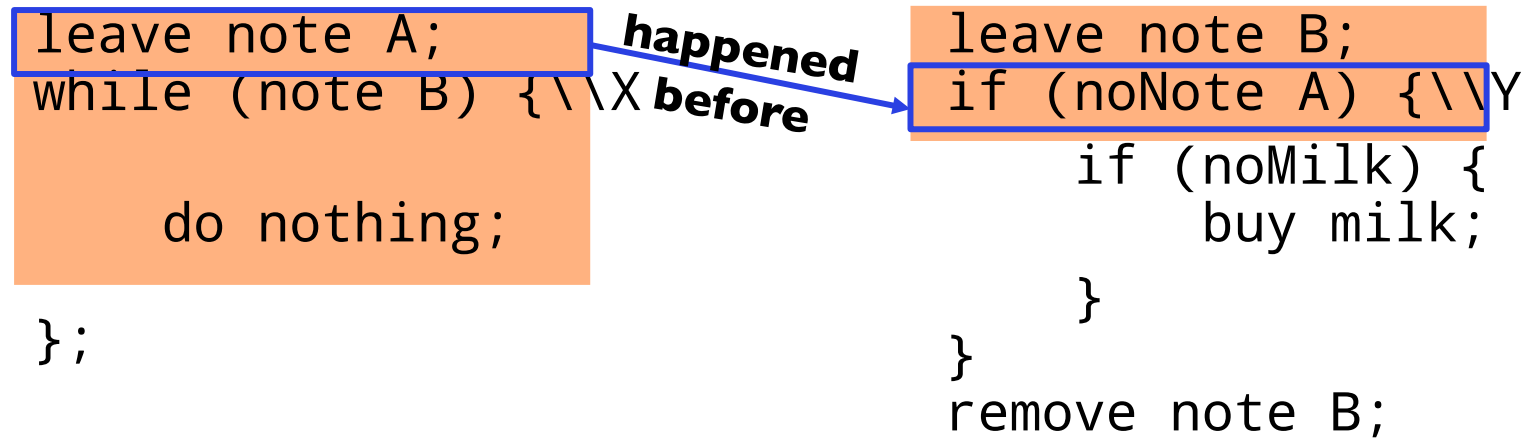
```
leave note A;  
while (note B) {  
    do nothing;  
};  
  
if (noMilk) {  
    buy milk;  
}  
remove note A;
```

```
leave note B;  
if (noNote A) {  
    if (noMilk) {  
        buy milk;  
    }  
}  
remove note B;
```



# Case 1

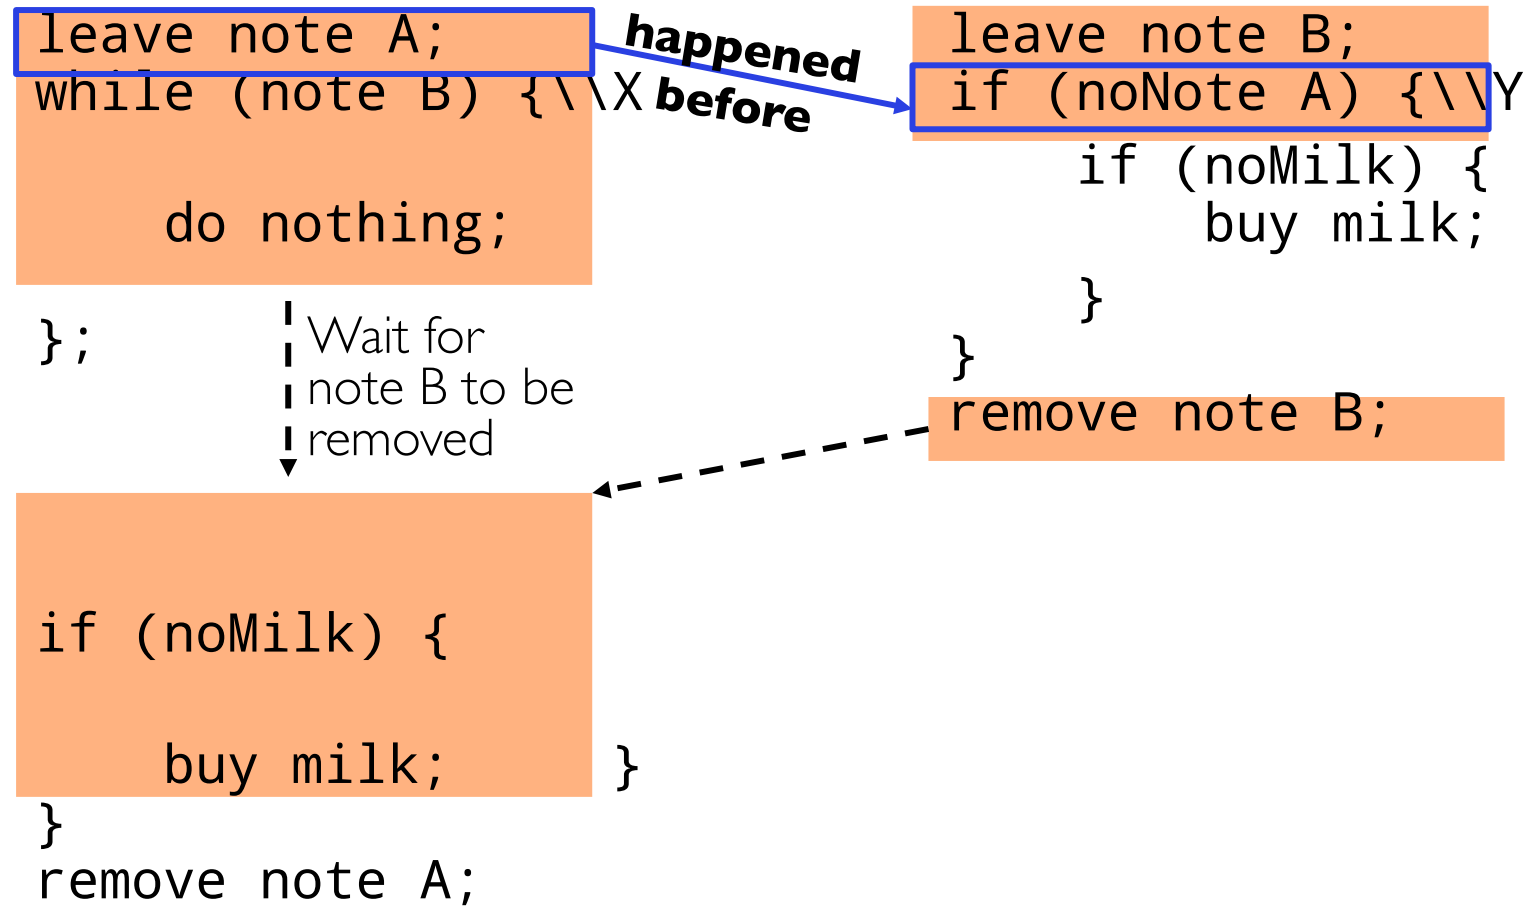
- “leave note A” happens before “if (noNote A)”



```
if (noMilk) {  
    buy milk;  
}  
remove note A;
```

# Case 1

- “leave note A” happens before “if (noNote A)”



## Case 2

- “if (noNote A)” happens before “leave note A”

The diagram illustrates the execution order of two code snippets. A blue arrow labeled "happened before" points from the "if (noNote A)" line in the right snippet to the "leave note A;" line in the left snippet, indicating that the condition check in the right snippet occurs before the leave action in the left snippet.

```
leave note A;
while (note B) {\X

    do nothing;

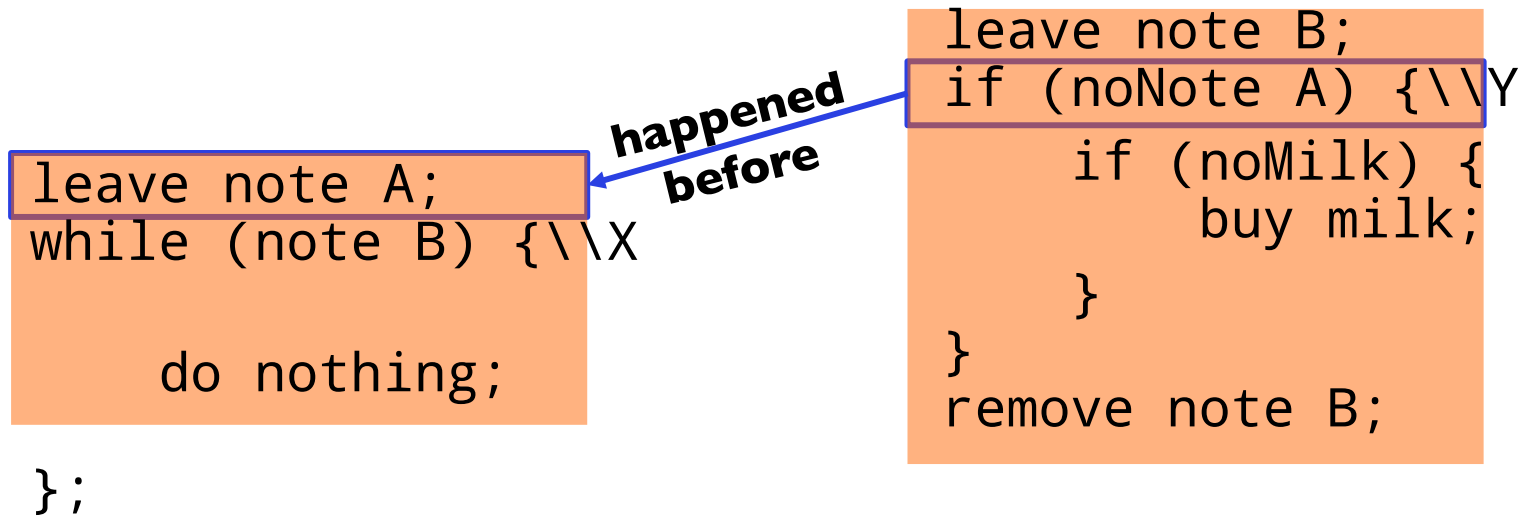
};

if (noMilk) {
    buy milk;
}
remove note A;
```

```
leave note B;
if (noNote A) {\Y
    if (noMilk) {
        buy milk;
    }
}
remove note B;
```

## Case 2

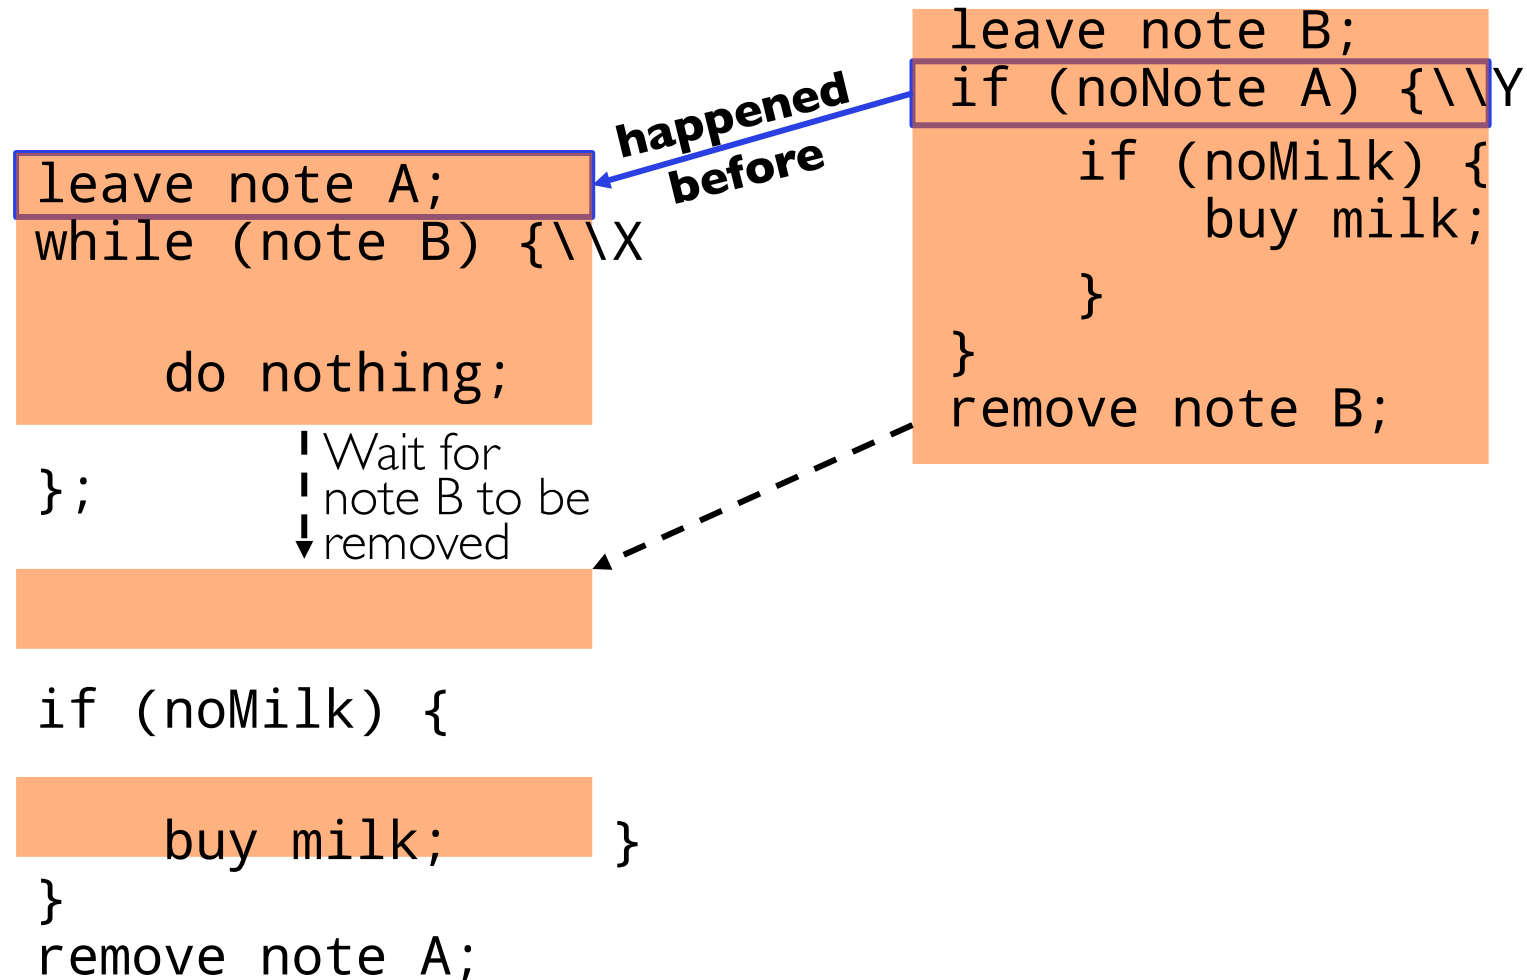
- “if (noNote A)” happens before “leave note A”



```
if (noMilk) {  
    buy milk;    }  
}  
remove note A;
```

## Case 2

- “if (noNote A)” happens before “leave note A”



---

## This Generalizes to Threads...

Leslie Lamport's "Bakery  
Algorithm" (1974)

Computer  
Systems

G. Bell, D. Siewiorek,  
and S.H. Fuller, Editors

---

### A New Solution of Dijkstra's Concurrent Programming Problem

Leslie Lamport  
Massachusetts Computer Associates, Inc.

---

**A simple solution to the mutual exclusion problem is  
presented which allows the system to continue to operate**

# Solution #3 discussion

---

Solution #3 works, but it's really unsatisfactory

- Really complex – even for this simple an example
  - » Hard to convince yourself that this really works
- A's code is different from B's – what if lots of threads?
  - » Code would have to be slightly different for each thread
- While A is waiting, it is consuming CPU time
  - » This is called “busy-waiting”

# Too Much Milk: Solution #4?

---

Recall our target lock interface:

- `acquire(&milklock)` – wait until lock is free, then grab
- `release(&milklock)` – Unlock, waking up anyone waiting
- These must be atomic operations – if two threads are waiting for the lock and both see it's free, only one succeeds to grab the lock

Then, our milk problem is easy:

```
    acquire(&milklock);  
    if (nomilk)  
        buy milk;  
    release(&milklock);
```



# Where are we going with synchronization?

---

Programs	Shared Programs
Higher-level API	Locks   Semaphores   Monitors   Send/Receive
Hardware	Load/Store   Disable Ints   Test&Set   Compare&Swap

Implement various higher-level synchronization primitives using atomic operations

# How to Implement Locks?

---

Prevents someone from doing something

Lock before entering critical section and  
before accessing shared data

Unlock when leaving, after accessing shared data



# Hardware Lock Instruction?

---

Is this a good idea?

What about putting a task to sleep?

What is the interface between the hardware and scheduler?

Complexity?

- » Done in the Intel 432

- » Each feature makes HW more complex and slow

# How about disabling interrupts?

---

Can we build multi-instruction atomic operations?

Recall: dispatcher gets control in two ways.

- » Internal: Thread does something to relinquish the CPU
- » External: Interrupts cause dispatcher to take CPU

On a uniprocessor, can avoid context-switching by:

- » Avoiding internal events (although virtual memory tricky)
- » Preventing external events by disabling interrupts

# How about disabling interrupts?

---

Naïve implementation of locks:

LockAcquire { disable Ints; }

LockRelease { enable Ints; }

Problems with this approach?

# How about disabling interrupts?

---

Consider following:

```
LockAcquire ( ) ;  
While (TRUE) { ; }
```

Real-Time system—no guarantees on timing!

Critical Sections might be arbitrarily long


What happens with I/O or other important events?

“Reactor about to meltdown. Help?”

# Disabling Interrupts – But more smartly

---

Key idea: maintain a lock variable and impose mutual exclusion only during operations on that variable

`int value = FREE;` 

```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts?  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

```
Release() {  
    disable interrupts;  
    if (anyone on wait queue) {  
        take thread off wait queue  
        Place on ready queue;  
    } else {  
        value = FREE;  
    }  
    enable interrupts;  
}
```

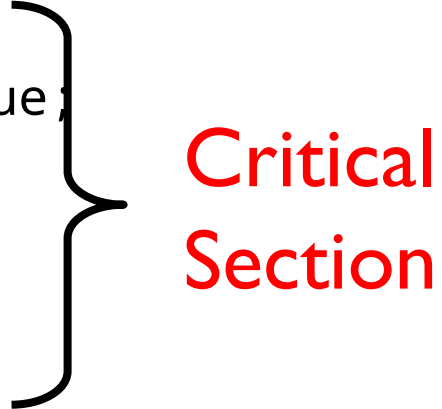
# New Lock Implementation: Discussion

---

Why do we need to disable interrupts at all?

- Avoid interruption between checking and setting lock value
- Otherwise two threads could think that they both have lock

```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts?  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```



**Critical  
Section**

Note: unlike previous solution, the critical section (inside `Acquire()`) is very short



# Interrupt Re-enable in Going to Sleep

---

What about re-enabling ints when going to sleep?

```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

# Interrupt Re-enable in Going to Sleep

---

```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
Enable Position → put thread on wait queue;  
                    Go to sleep();  
    } else {  
                    value = BUSY;  
    }  
    enable interrupts;  
}
```

Before Putting thread on the wait queue?

# Interrupt Re-enable in Going to Sleep

---

What about re-enabling ints when going to sleep?

```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

Enable Position →

After putting the thread on the wait queue?

# Interrupt Re-enable in Going to Sleep

---

What about re-enabling ints when going to sleep?

```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
Enable Position →    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```


After putting the thread on the wait queue?

# How to Re-enable After Sleep()?

---

In scheduler, since interrupts are disabled when you call sleep:

- Responsibility of the next thread to re-enable ints
- When the sleeping thread wakes up, returns to acquire and re-enables interrupts



# Atomic Read-Modify-Write Instructions

---

Problems with previous solution:

- Can't give lock implementation to users
- Doesn't work well on multiprocessor

Alternative: **atomic instruction sequences**

- These instructions read a value and write a new value atomically
- **Hardware** is responsible for implementing this correctly
  - » on both uniprocessors (not too hard)
  - » and multiprocessors (requires help from cache coherence protocol)
- Unlike disabling interrupts, can be used on both uniprocessors and multiprocessors

# Examples of Read-Modify-Write

---

- ```
test&set (&address) {  
    result = M[address];  
    M[address] = 1;  
    return result;  
}
```

```
/* most architectures */  
// return result from "address" and  
// set value at "address" to 1
```
- ```
swap (&address, register) {  
    temp = M[address];  
    M[address] = register;  
    register = temp;  
}
```


```
/* x86 */  
// swap register's value to  
// value at "address"
```
- ```
compare&swap (&address, reg1, reg2) { /* x86 (returns old value), 68000  
*/  
    if (reg1 == M[address]) {  
        M[address] = reg2;  
        return success;  
    } else {  
        return failure;  
    }  
}
```

```
// If memory still == reg1,  
// then put reg2 => memory  
  
// Otherwise do not change memory
```

# Using of Compare&Swap for queues

---

```
addToQueue(&object) {  
    do { // repeat until no conflict  
        ld r1, M[root] // Get ptr to current  
head      st r1, M[object] // Save link in new  
object  
    } until (compare&swap(&root,r1,object));  
}
```





# Implementing Locks with test&set

---

Simple lock that doesn't require entry into the kernel:

```
// (Free) Can access this memory location from user space!
int mylock = 0; // Interface: acquire(&mylock);
                //                release(&mylock);

acquire(int *thelock) {
    while (test&set(thelock)); // Atomic operation!
}

release(int *thelock) {
    *thelock = 0; // Atomic operation!
}
```

# Implementing Locks with test&set

---

## Simple explanation:

- If lock is free, test&set reads 0 and sets lock=1, so lock is now busy. It returns 0 so while exits.
- If lock is busy, test&set reads 1 and sets lock=1 (no change) It returns 1, so while loop continues.
- When we set thelock = 0, someone else can get lock.

## Busy-Waiting: thread consumes cycles while waiting

- For multiprocessors: every test&set() is a write, which makes value ping-pong around in cache (using lots of network BW)

# Problem: Busy-Waiting for Lock

---



## Positives for this solution

- Machine can receive interrupts
- User code can use this lock
- Works on a multiprocessor

## Negatives

- This is very inefficient as thread will consume cycles waiting
- Waiting thread may take cycles away from thread holding lock (no one wins!)
- Homework/exam solutions should avoid busy-waiting!

# Better Locks using test&set

---

Idea: only busy-wait to atomically check lock value



```
int guard = 0; // Global Variable!
int mylock = FREE; // Interface: acquire(&mylock);
                        //                      release(&mylock);
```

```
acquire(int *thelock) {
    // Short busy-wait time
    while (test&set(guard));
    if (*thelock == BUSY) {
        put thread on wait queue;
        go to sleep() & guard = 0;
        // guard == 0 on wakeup!
    } else {
        *thelock = BUSY;
        guard = 0;
    }
}
```

```
release(int *thelock) {
    // Short busy-wait time
    while (test&set(guard));
    if anyone on wait queue {
        take thread off wait queue
        Place on ready queue;
    } else {
        *thelock = FREE;
    }
    guard = 0;
```

# Linux futex: Fast Userspace Mutex

```
#include <linux/futex.h>
#include <sys/time.h>

int futex(int *uaddr, int futex_op, int val,
          const struct timespec *timeout );
```

*uaddr* points to a 32-bit value in user space

*futex\_op*

- FUTEX\_WAIT – if *val* == \**uaddr* sleep till FUTEX\_WAKE
  - » **Atomic** check that condition still holds after we disable interrupts (in kernel!)
- FUTEX\_WAKE – wake up at most *val* waiting threads
- FUTEX\_FD, FUTEX\_WAKE\_OP, FUTEX\_CMP\_REQUEUE: More interesting operations!

*timeout*

- ptr to a *timespec* structure that specifies a timeout for the op

# Linux futex: Fast Userspace Mutex

---

```
#include <linux/futex.h>
#include <sys/time.h>

int futex(int *uaddr, int futex_op, int val,
          const struct timespec *timeout );
```

Interface to the kernel sleep() functionality!

- Let thread put themselves to sleep – conditionally!

futex is not exposed in libc; it is used within the implementation of pthreads

- Can be used to implement locks, semaphores, monitors, etc...

# Example: First try: T&S and futex

---

```
int mylock = 0; // Interface: acquire(&mylock);
                //                release(&mylock);

acquire(int *thelock) {                release(int *thelock) {
    while (test&set(thelock)) {        thelock = 0; // unlock
        futex(thelock, FUTEX_WAIT, 1);  futex(&thelock, FUTEX_WAKE, 1);
    }
}                                     }
```

Sleep interface by using futex – no busywaiting

No overhead to acquire lock

Every unlock has to call kernel to potentially wake someone up – even if none

## Example: Try #2: T&S and futex

---

```
bool maybe = false;
int mylock = 0; // Interface:
acquire(&mylock, &maybe_waiters);
//
```

```
release(&mylock, &maybe_waiters);
```

```
acquire(int *thelock, bool *maybe) {
    while (test&set(thelock)) {
        // Sleep, since lock busy!
        *maybe = true;
        futex(thelock, FUTEX_WAIT, 1);
```

```
        // Make sure other sleepers not stuck
        *maybe = true;
    }
}
```

```
release(int *thelock, bool *maybe) {
    thelock = 0;
    if (*maybe) {
        *maybe = false;
        // Try to wake up someone
        futex(&value, FUTEX_WAKE, 1);
    }
}
```

This is syscall-free in the uncontended case

- Temporarily falls back to syscalls if multiple waiters, or concurrent acquire/release
  - But it can be considerably optimized!
- See “[Futexes are Tricky](#)” by Ulrich Drepper



# Where are we going with synchronization?

|                  |                                                     |
|------------------|-----------------------------------------------------|
| Programs         | Shared Programs                                     |
| Higher-level API | Locks   Semaphores   Monitors   Send/Receive        |
| Hardware         | Load/Store   Disable Ints   Test&Set   Compare&Swap |

- We are going to implement various higher-level synchronization primitives using atomic operations
  - Everything is pretty painful if only atomic primitives are load and store
  - Need to provide primitives useful at user-level

# Higher-level Primitives than Locks

---

Goal of last couple of lectures:


- What is right abstraction for synchronizing threads that share memory?
- Want as high a level primitive as possible

Synchronization is a way of coordinating multiple concurrent activities that are using shared state

- This lecture and the next presents some ways of structuring sharing

# Producer-Consumer with a Bounded Buffer

---



## Problem Definition

- Producer(s) put things into a shared buffer
  - Consumer(s) take them out
- Need synchronization to coordinate producer/consumer

Don't want producer and consumer to have to work in lockstep, so put a fixed-size buffer between them


- Need to synchronize access to this buffer
  - Producer needs to wait if buffer is full
  - Consumer needs to wait if buffer is empty

# Producer-Consumer with a Bounded Buffer

---

Example 1: GCC compiler

– `cpp` | `cc1` | `cc2` | `as` | `ld`




Example 2: Coke machine

- Producer can put limited number of Cokes in machine
- Consumer can't take Cokes out if machine is empty

Others: Web servers, Routers, ....

# Circular Buffer Data Structure (sequential case)

```
typedef struct buf {  
    int write_index;  
    int read_index;  
    <type> *entries[BUFSIZE];  
} buf_t;
```



Insert: write & bump write ptr (enqueue)

Remove: read & bump read ptr (dequeue)

How to tell if Full (on insert) Empty (on remove)?

And what do you do if it is?

What needs to be atomic?

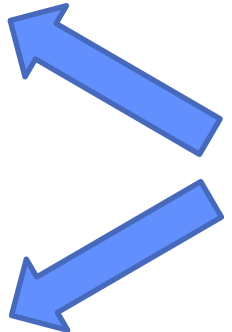
# Circular Buffer – first cut

---

mutex buf\_lock = <initially unlocked>

```
Producer(item) {  
    acquire(&buf_lock);  
    while (buffer full) {}; // Wait for a free  
slot  
    enqueue(item);  
    release(&buf_lock);  
}
```

```
Consumer() {  
    acquire(&buf_lock);  
    while (buffer empty) {}; // Wait for arrival  
    item = dequeue();  
    release(&buf_lock);  
    return item  
}
```



Will we ever come out of the wait loop?

## Circular Buffer – 2<sup>nd</sup> cut



mutex buf\_lock = <initially unlocked>

```
Producer(item) {  
    acquire(&buf_lock);  
    while (buffer full) {release(&buf_lock);  
acquire(&buf_lock);}   
    enqueue(item);  
    release(&buf_lock);  
}
```

```
Consumer() {  
    acquire(&buf_lock);  
    while (buffer empty) {release(&buf_lock);  
acquire(&buf_lock);}   
    item = dequeue();  
    release(&buf_lock);  
    return item  
}
```

What happens when one is waiting for the other?

- Multiple cores ?
- Single core ?

# Semaphores

---

Semaphores are a type of generalized lock

First defined by Dijkstra in late 60s

Main synchronization primitive used in original UNIX



# Semaphores

---

A Semaphore has a **non-negative integer value** and supports the following operations:

- Set value when you initialize
- **Down()** or **P()**: an atomic operation that waits for semaphore to become positive, then decrements it by 1
  - » Think of this as the wait() operation
- **Up()** or **V()**: an atomic operation that increments the semaphore by 1, waking up a waiting P, if any
  - » Think of this as the signal() operation

# Semaphores Like Integers Except...

---

Semaphores are like integers, except:

- No negative values
- Only operations allowed are P and V – can't read or write value, except initially
  - Operations must be atomic
    - » Two P's together can't decrement value below zero
    - » Thread going to sleep in P won't miss wakeup from V – even if both happen at same time

# Two Uses of Semaphores

---

Mutual Exclusion (initial value = 1)

Also called “Binary Semaphore” or “mutex”.

Can be used for mutual exclusion, just like a lock:

```
        semaP(&mysem);  
// Critical section goes here  
        semaV(&mysem);
```

# Two Uses of Semaphores

---

Scheduling Constraints (initial value = 0)

- Allow thread 1 to wait for a signal from thread 2
- thread 2 **schedules** thread 1 when a given **event** occurs

Suppose you had to implement ThreadJoin which must wait for thread to terminate:

**Initial value of semaphore = 0**

```
ThreadJoin {  
    semaP(&mysem);  
}  
ThreadFinish {  
    semaV(&mysem);  
}
```